

TDDC17: Introduction to Automated Planning

Jonas Kvarnström

Automated Planning Group

Department of Computer and Information Science

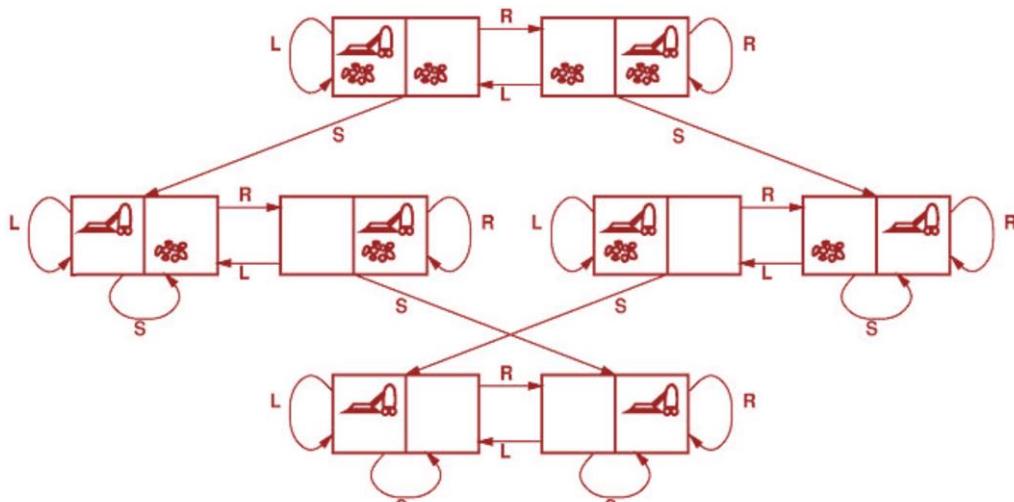
Linköping University

Introduction to Planning

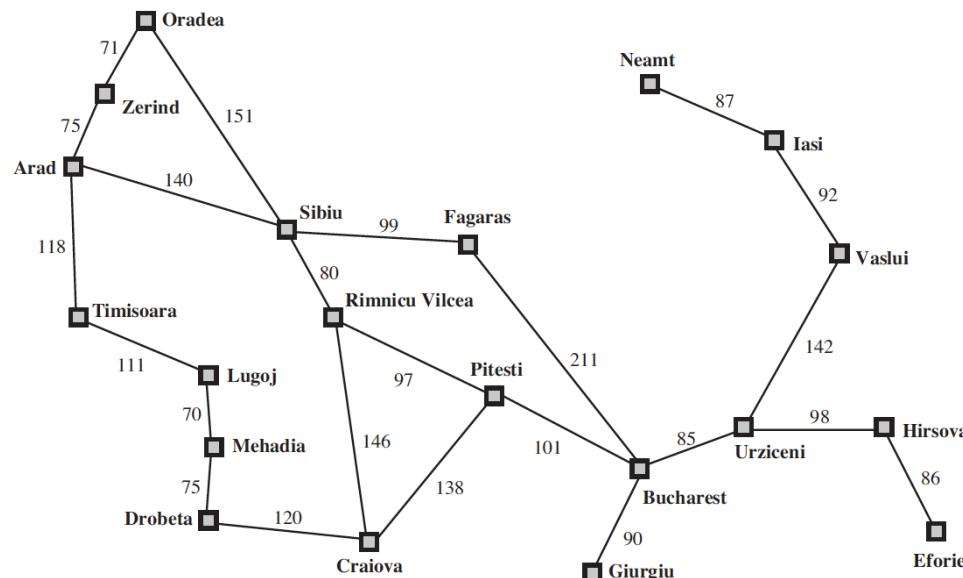
One way of defining planning:

*Using **knowledge** about the world,
including possible actions and their results,
to **decide** what to do and when
in order to achieve an **objective**,
before you actually start doing it*

You have done this before!

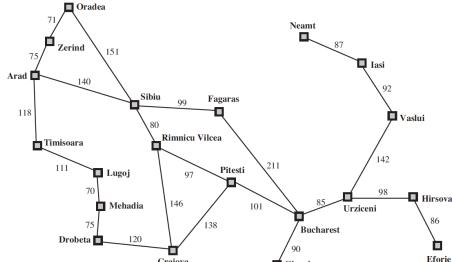


Using **knowledge** about the world, including possible actions and their results, to **decide** what to do and when in order to achieve an **objective**, **before** you actually start doing it

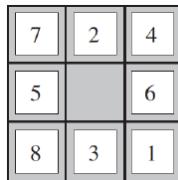


Are we done?

Domain-specific search guidance – too much (human) work!



Straight line distance from city n to goal city n'



$h_2(n)$: The sum of the manhattan distances for each tile that is out of place.

($3+1+2+2+2+3+3+2=18$) . The manhattan distance is an under-estimate because there are tiles in the way.

More complex



???

Planning:
General heuristics

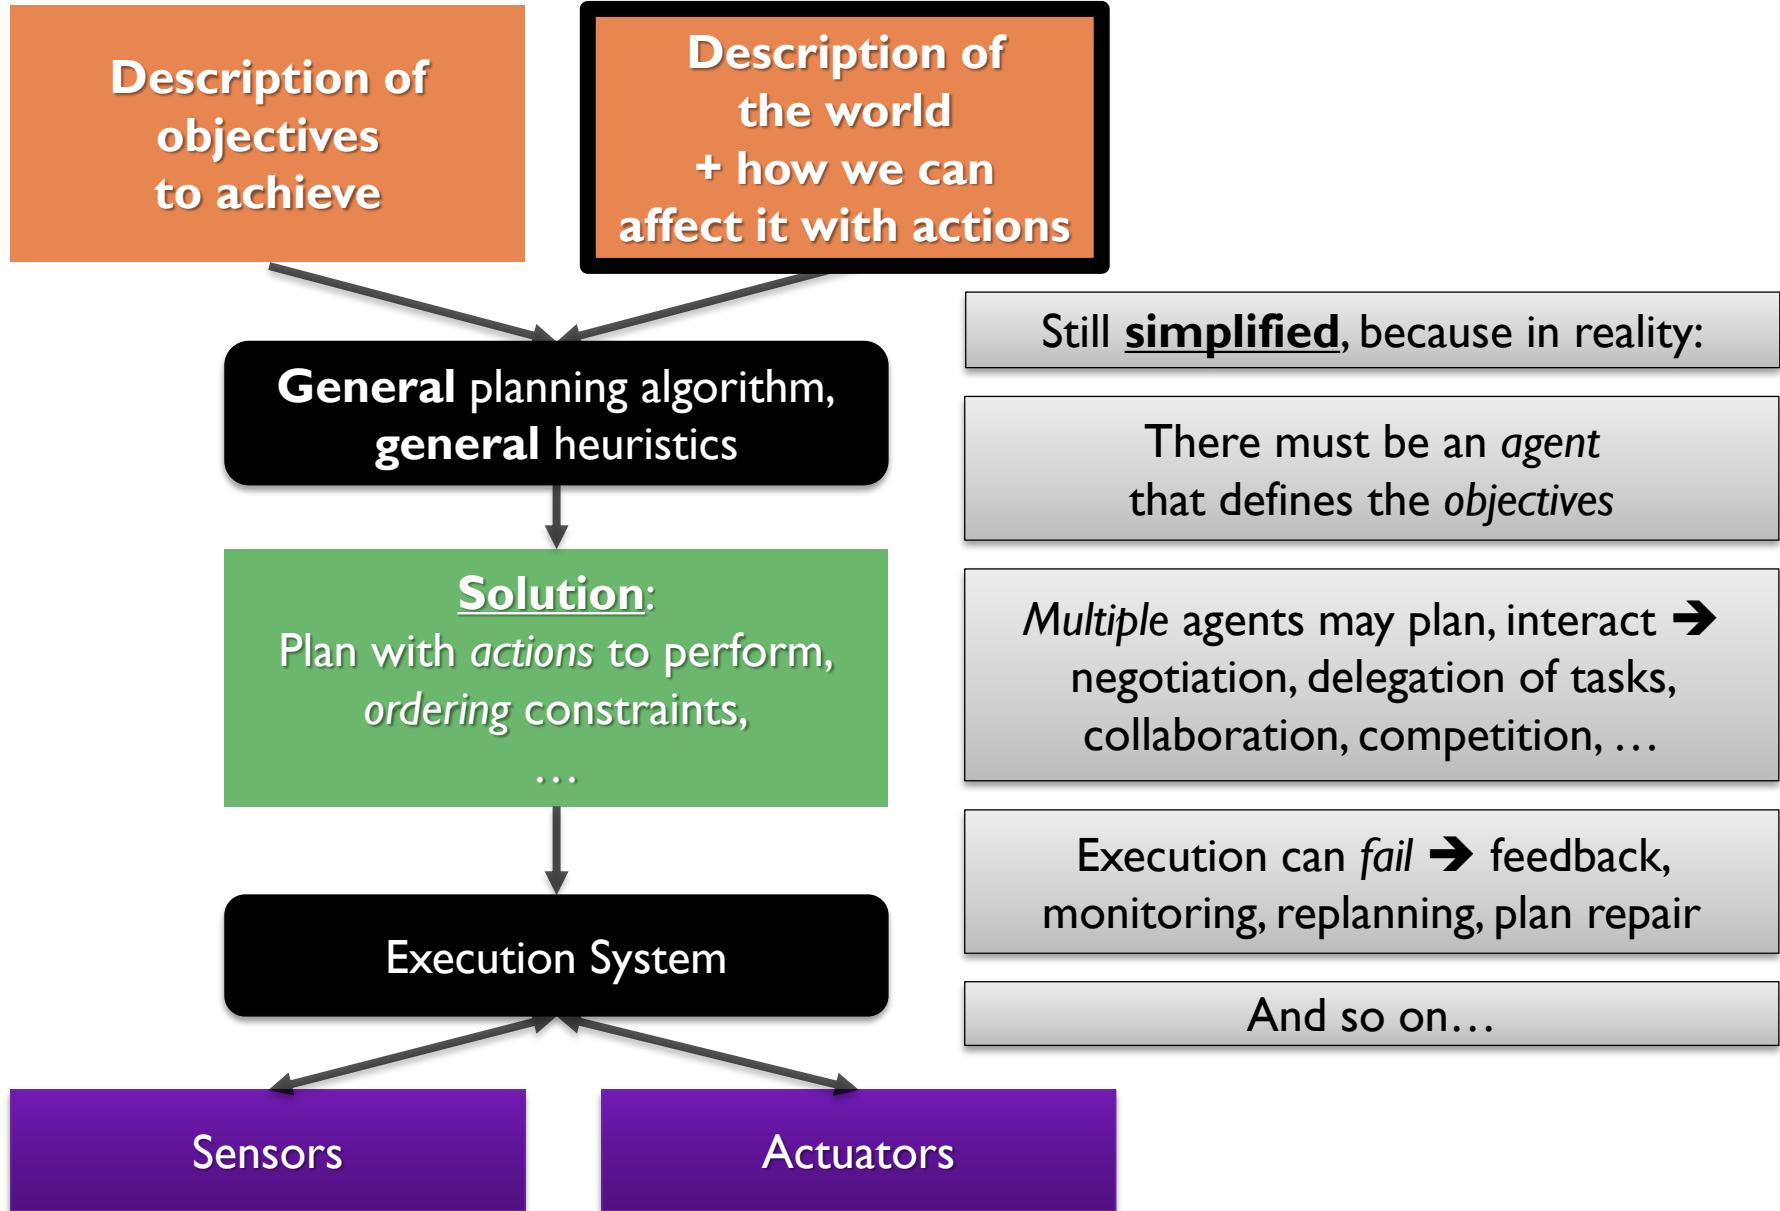
Pattern Databases
Landmarks
 FF , h^m , merge-and-shrink, ...

Planning:
Entirely different search spaces

Partial Order Causal Link
SAT planning
Planning Graphs, ...

Need a common, well-structured representation
for planners/heuristics to analyse!

AI Planning: A Simplified View



Domains 1: Blocks World



- Must support simple examples such as the **Blocks World**

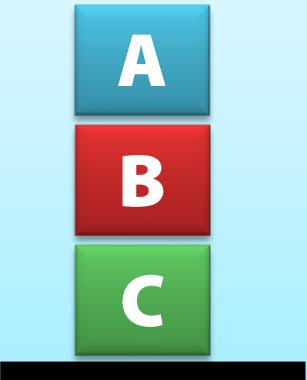
You



Current state of the world



Your greatest desire



Domains 2: Towers of Hanoi



Another "toy domain" – useful for examples!



General Knowledge: Problem Domain

- There are *pegs* and *disks*
- A disk can be *on* a peg
- One disk can be *above* another
- Actions:
Move topmost disk from *x* to *y*, without placing larger disks
on smaller disks

Specific Knowledge and Objective: Problem Instance

- 3 pegs, 7 disks
- All disks currently on peg 2, in order of increasing size
- We want: All disks on the *third* peg, in order of increasing size

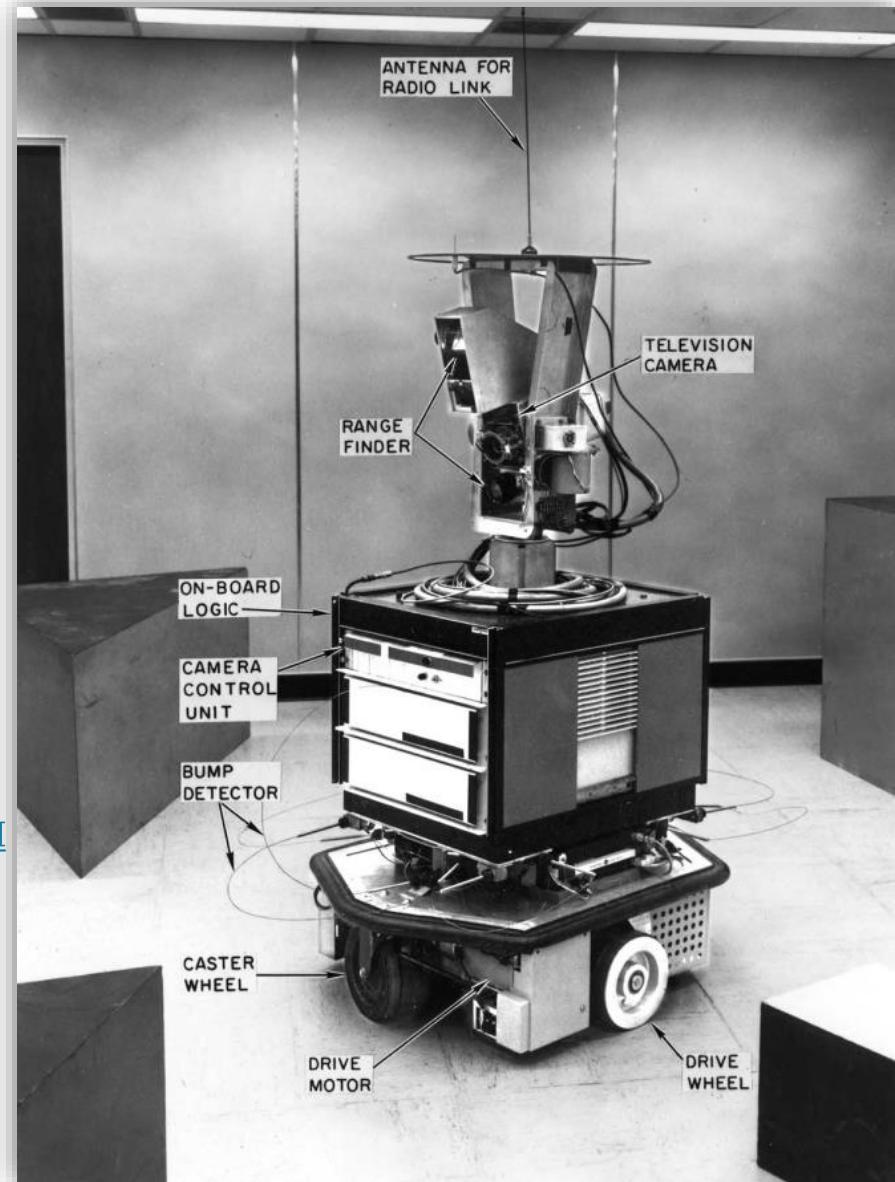
Domains 3: Shakey



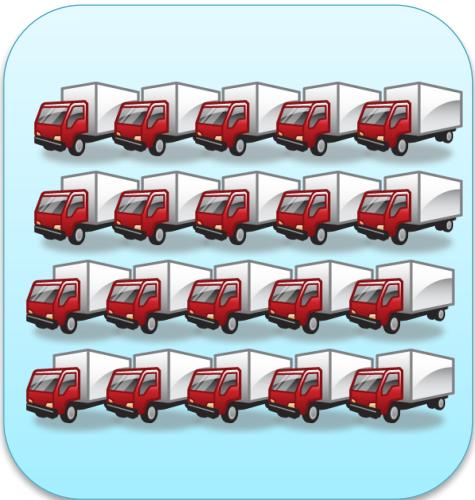
- Classical robot example:
Shakey (1969)

- Available **actions**:
 - *Moving* to another location
 - *Turning* light switches on and off
 - *Opening* and *closing* doors
 - *Pushing* movable objects around
 - ...

- **Goals:**
 - *Be in room 4 with objects A,B,C*
 - <http://www.youtube.com/watch?v=qXdn6ynwpI>

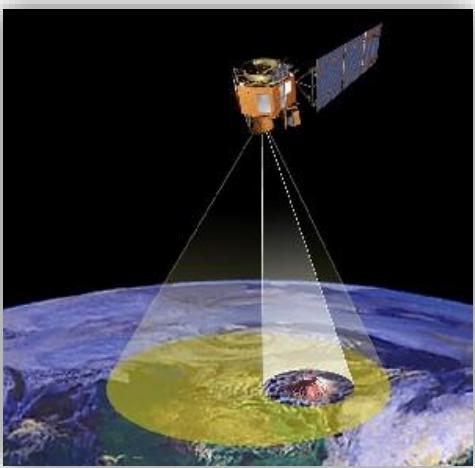


Domains 4



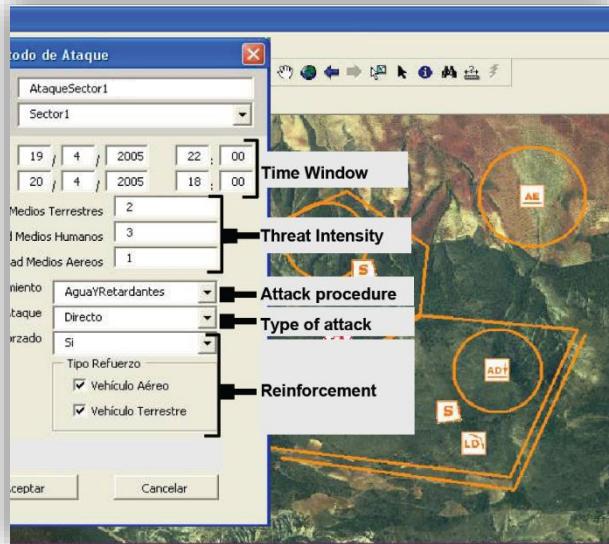
Logistics:

Use a fleet of trucks
to efficiently deliver
packages



On-board planning
to view interesting
natural events:

<http://ase.jpl.nasa.gov/>



SIADEX –
plan for firefighting
Limited resources
Plan execution is
dangerous!

Domains 5: Dock Worker Robots (DWR)



Containers shipped
in and out of a harbor



Cranes move containers
between "piles" and robotic trucks



Classical Planning



How to define a formal language for planning problems?

- We want a general language:
 - Allows a wide variety of domains to be modeled



- We want good algorithms
 - Generate plans quickly
 - Generate high quality plans
- Easier with more limited languages

Conflicting desires!

- Many early planners made similar tradeoffs
 - Later called "classical planning"
Restricted, but a good place to start

Modeling Classical Planning Problems

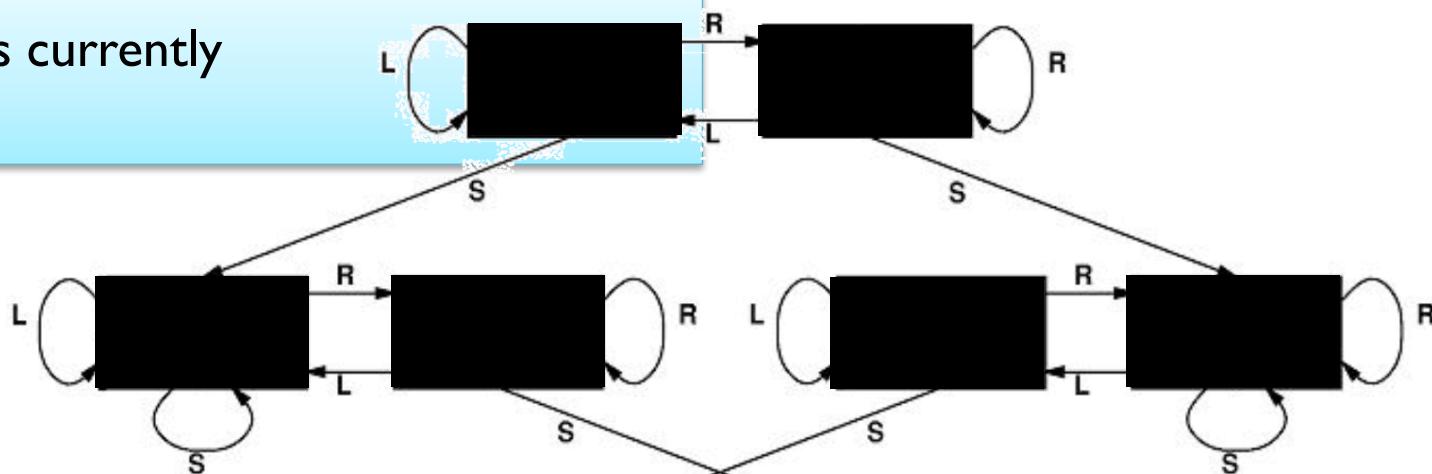
Facts and States: Introduction

14

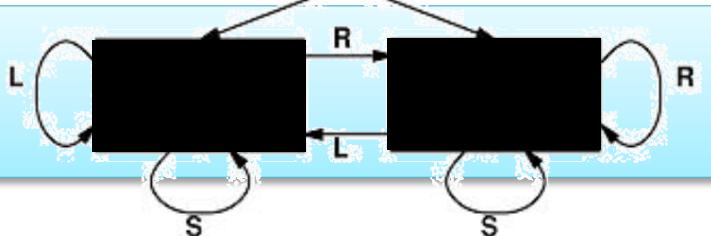
jonkv@ida

- As you saw before, we are interested in **states of the world**

The world is currently
in *this* state



We want to reach
one of *these* states



Is this information sufficient?

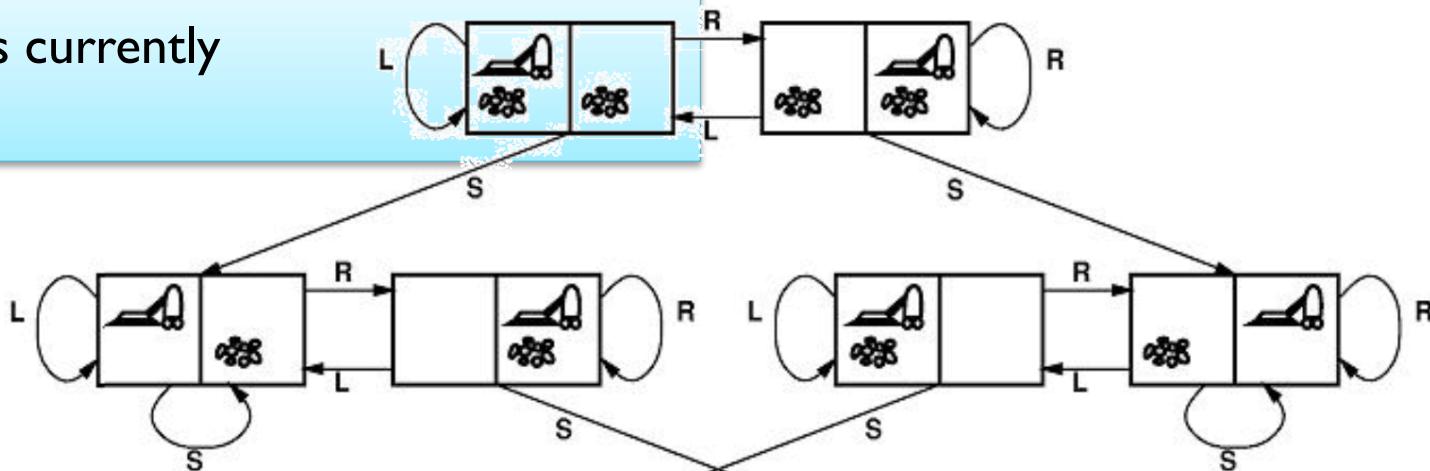
No – Need to **analyze** states,
find **differences** compared to goal states,
find **relevant** actions, ...!

Facts and States: Introduction

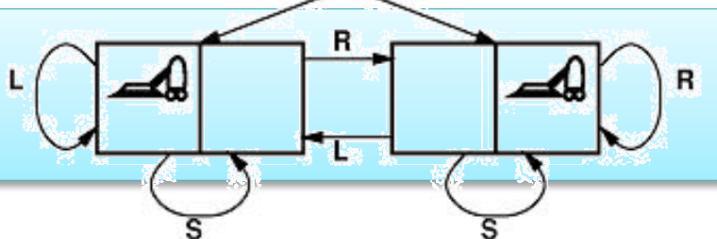


- We need **information about** every state:

The world is currently
in *this* state



We want to reach
one of *these* states



**First: Rooms and
vacuum cleaners
are objects**

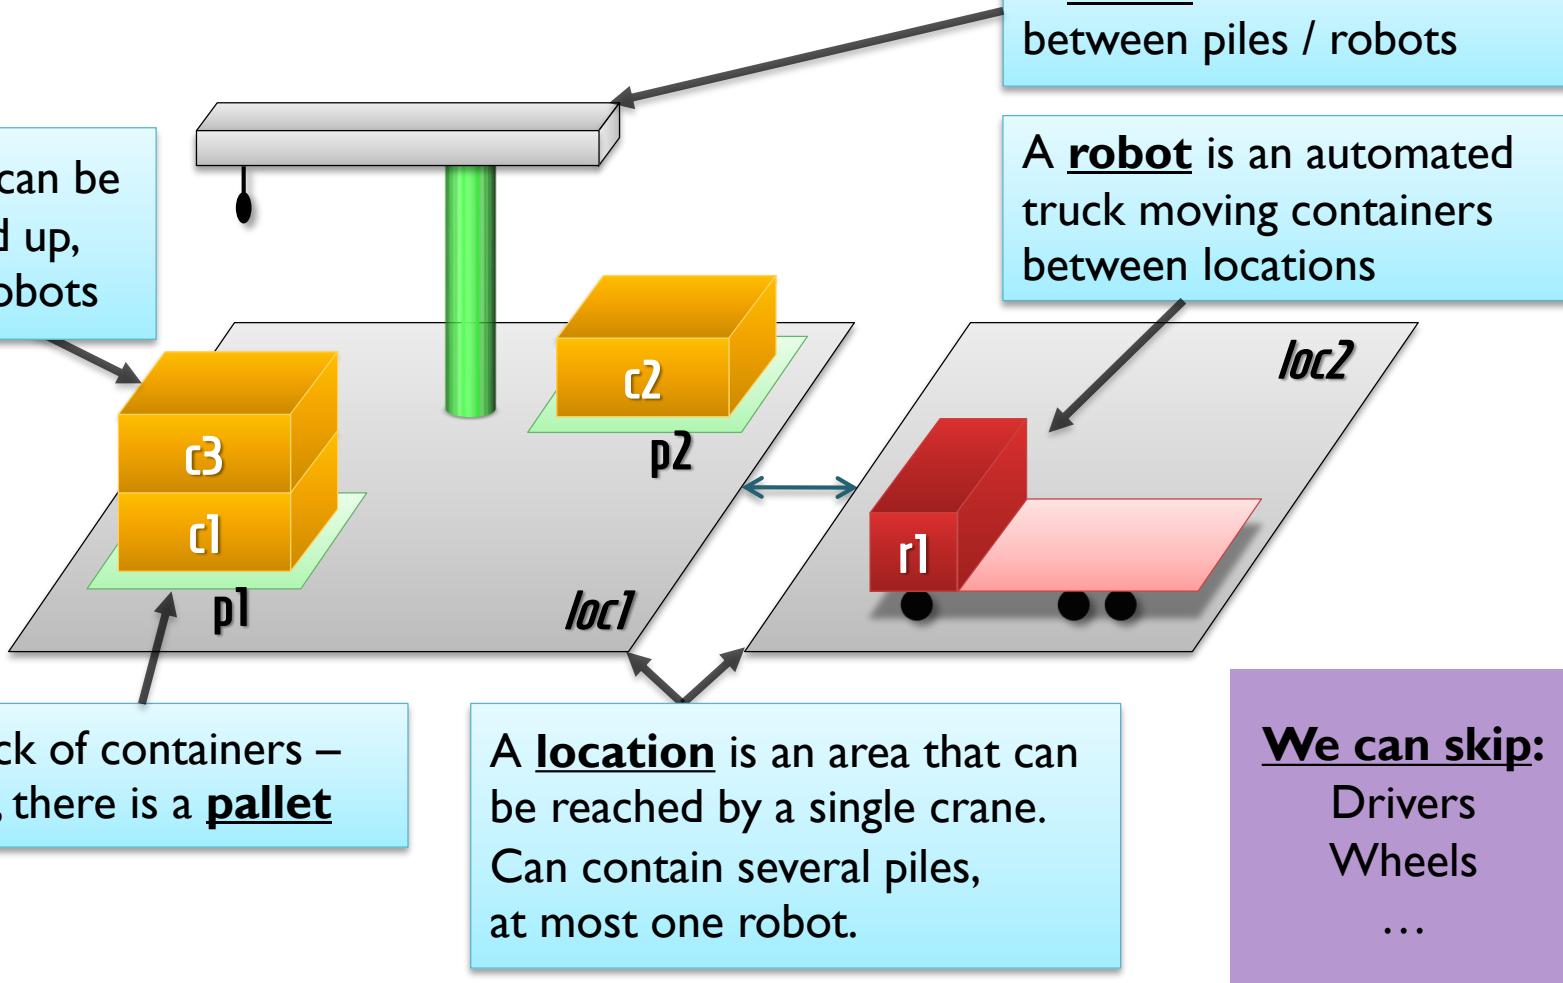
Efficient planning depends on describing states as collections of facts:

We are in a state where there is dirt in both rooms,
and the vacuum cleaner is in the leftmost room

Objects 1: Object Types



- We can often specify **object types**



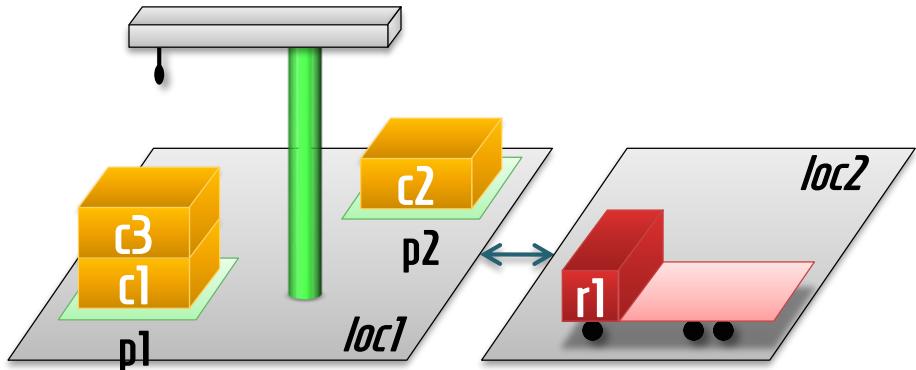
Essential: Determine which types are **relevant** for the **problem** and **objective**!

Objects 2: Actual objects

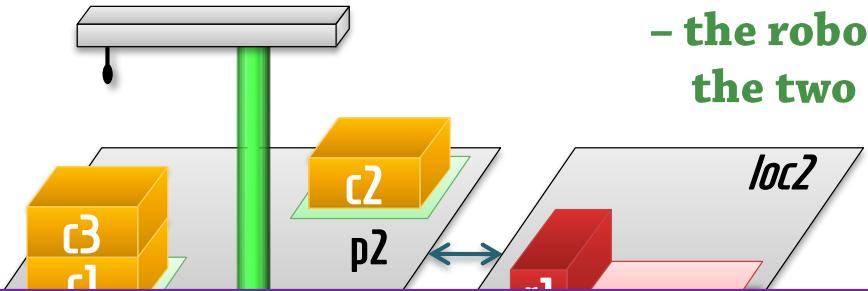


- We then specify sets of **actual objects**

- robot:** { r1 }
- location:** { loc1, loc2 }
- crane:** { k1 }
- pile:** { p1, p2 }
- container:** { c1, c2, c3, pallet }



- Most planners use a **first-order representation**:
 - Every **fact** is represented as a logical **atom**: Predicate symbol + arguments
 - **Properties of the world**
 - **raining** - it is raining [not in the standard DWR domain!]
 - **Properties of single objects...**
 - **empty**(crane) - the crane is not holding anything
 - **Relations between objects**
 - **attached**(pile, location) - the pile is in the given location
 - **Relations between >2 objects**
 - **can-move**(robot, location, location)
 - the robot can move between the two locations [Not in DWR]



Essential: Determine what is **relevant** for the **problem** and **objective**!

Facts / Predicates in DWR



- **Reference:** All predicates for DWR, and their intended meaning

"Fixed/Rigid"
(can't
change)

adjacent $(loc1, loc2)$

; can move from $loc1$ directly to $loc2$

attached (p, loc)

; pile p attached to loc

belong (k, loc)

; crane k belongs to loc

"Dynamic"
(modified by
actions)

at (r, loc)

; robot r is at loc

occupied (loc)

; there is a robot at loc

loaded (r, c)

; robot r is loaded with container c

unloaded (r)

; robot r is empty

holding (k, c)

; crane k is holding container c

empty (k)

; crane k is not holding anything

in (c, p)

; container c is somewhere in pile p

top (c, p)

; container c is on top of pile p

on $(c1, c2)$

; container $c1$ is on container $c2$

States 1: State of the World



- A **state (of the world)** should specify exactly which facts (**ground atoms**) are true/false in the world at a given time

Ground = without variables

We know all **predicates** that exist:
adjacent(location, location), ...

We know which **objects** exist
for each type

We can calculate all ground atoms

adjacent(loc1,loc1)
adjacent(loc1,loc2)
...
attached(pile1,loc1)
...

These are the **facts** to keep track of!

We can find all possible states!

Every assignment of true/false to the ground atoms is a distinct state

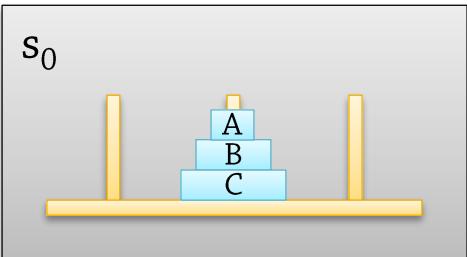
Number of states: $2^{\text{number of atoms}}$ – enormous, but finite (for classical planning!)

States 2: Efficient Representation



■ Efficient specification and storage for a single state:

- Specify which atoms are true
 - All other atoms have to be false – what else would they be?
- → A state of the world is specified as a set containing all variable-free atoms that [are, were, will be] true in the world
 - $s_0 = \{ \text{on}(A,B), \text{on}(B,C), \text{in}(A,2), \text{in}(B,2), \text{in}(C,2), \text{top}(A), \text{bot}(C) \}$



s_0

$\text{top}(A)$	
$\text{on}(A,B)$	$\text{in}(B,2)$
$\text{on}(B,C)$	$\text{in}(C,2)$
$\text{bot}(C)$	

$\text{top}(A) \in s_0 \rightarrow \text{top}(A) \text{ is true in } s_0$
 $\text{top}(B) \notin s_0 \rightarrow \text{top}(B) \text{ is false in } s_0$

States 3: Initial State

22

jonkv@ida

Initial states in classical STRIPS planning:

- We assume *complete information* about the **initial state** s_0 (before any action)

Complete **relative to the model**:

We must know everything
about those predicates and objects
we have specified...

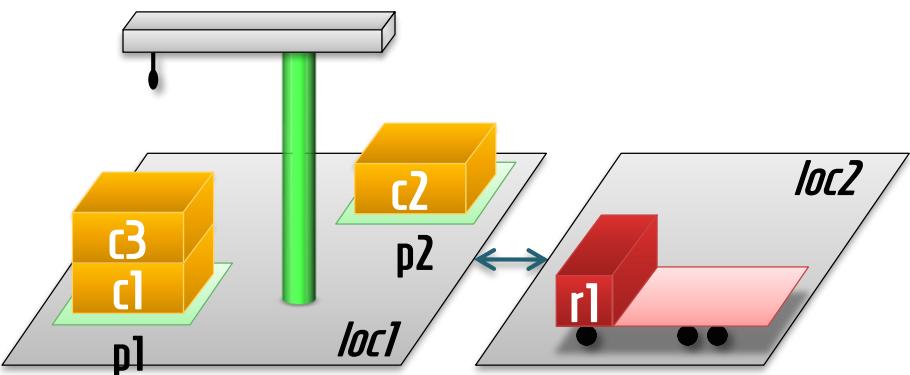
But not whether it's raining!

- So we can use a set of true atoms

{

attached(p1, loc1), in(c1, p1), on(c1, pallet), in(c3, p1), on(c3, c1), top(c3, p1),
attached(p2, loc1), in(c2, p2), on(c2, pallet), top(c2, p2),
belong(crane1, loc1), empty(crane1),
at(r1, loc2), unloaded(r1), occupied(loc2) ,
adjacent(loc1, loc2), adjacent(loc2, loc1),

}

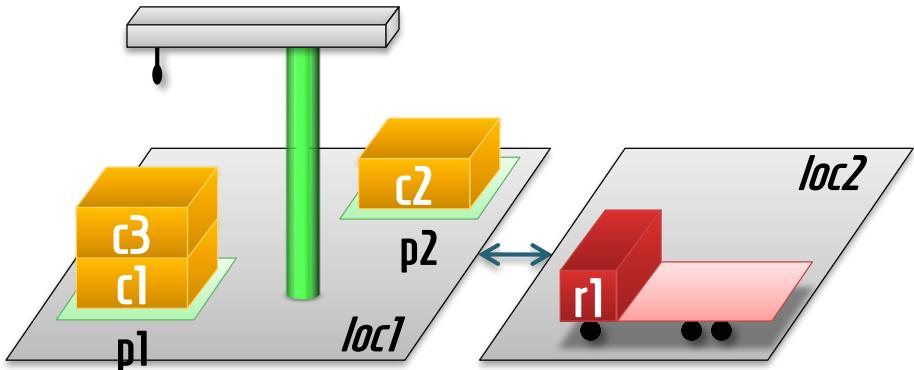


States 4: Goal States

23

jonkv@ida

- Classical STRIPS planning: Reach one of possibly many **goal states**
 - Can be specified as a set of literals that must hold
 - Example: Containers 1 and 3 should be in pile 2
 - We don't care about their order, or any other fact
 - { **in(c1,p2),
in(c3,p2) }** }



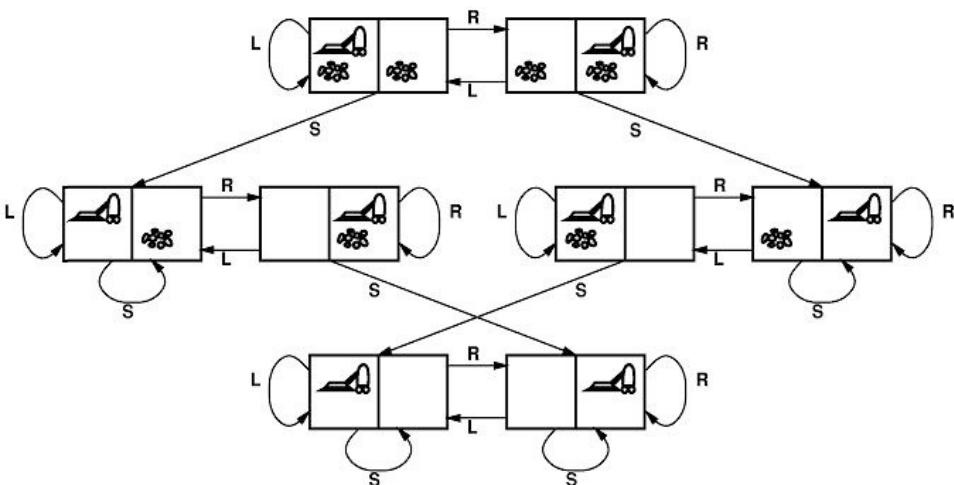
Actions 1: Intro

24

jonkv@ida

- Actions in plain search (lectures 2-3):
 - Assumed a *transition / successor function*

Result(State,Action) - A description of what each action does (Transition function)



- But how to specify it succinctly?
 - States consist of facts
 - Actions add or remove facts

Actions 2: Operators

25

jonkv@ida

- Define operators or action schemas:

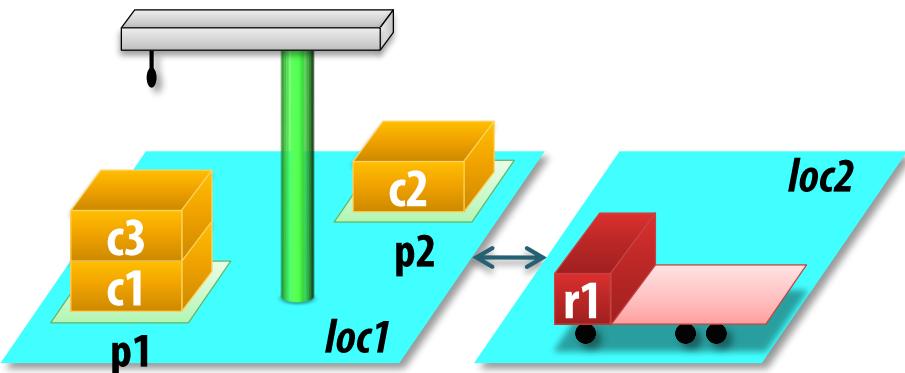
- **move**(*robot*, *location1*, *location2*)

- Precondition: **at**(*robot*, *location1*) \wedge
adjacent(*location1*, *location2*) \wedge
 \neg **occupied**(*location2*)

The action is **applicable** in a state *s* if its precond is true in *s*

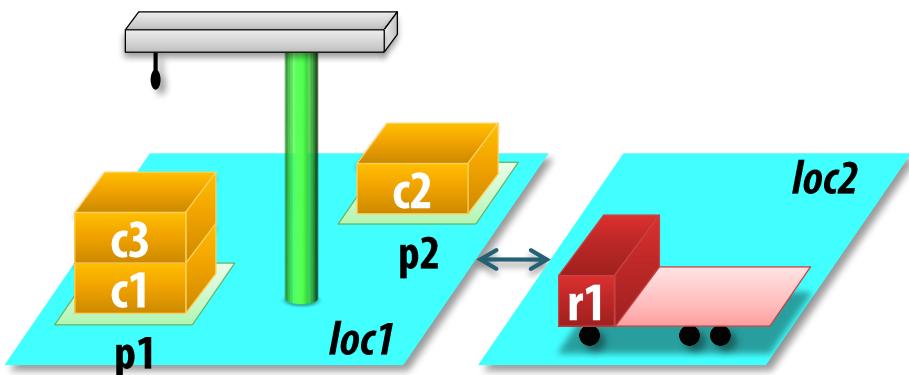
- Effects:
 \neg **at**(*robot*, *location1*),
at(*robot*, *location2*),
 \neg **occupied**(*location1*),
occupied(*location2*)

The result of applying the action in state *s*:
s – negated effects
+ positive effects



Actions 3: Instances

- The planner **instantiates** these schemas
 - Applies them to parameters of the correct type
 - **Example: move(r1, loc1, loc2)**
 - Precondition: **at(r1, loc1) \wedge adjacent(loc1, loc2) \wedge \neg occupied(loc2)**
 - Effects: **\neg at(r1, loc1),
at(r1, loc2),
 \neg occupied(loc1),
occupied(loc2)**



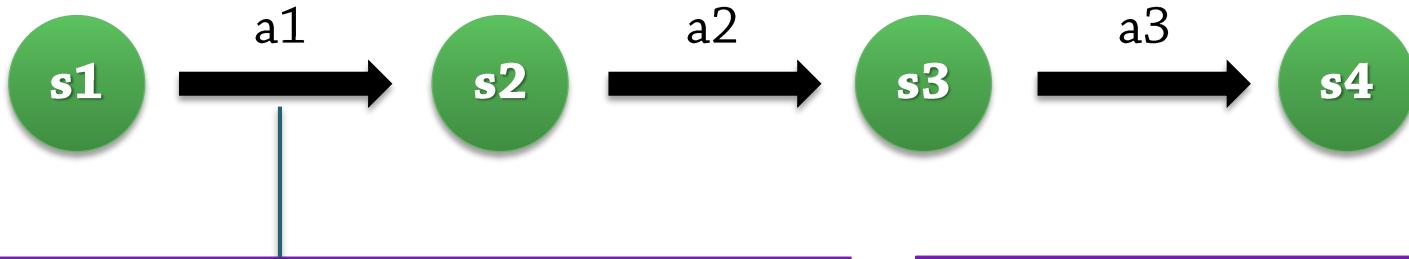
Actions 4: Step by Step

- In classical planning (the basic, limited form):

We know the initial state

Each action corresponds to one state update

Time is not modeled, and never multiple state updates for one action



We know how states are changed by actions

→ Deterministic, can completely predict the state of the world after a sequence of actions!

The solution to the problem will be a sequence of actions

Planning Domain, Problem Instance

28

jonkv@ida



Split knowledge into two parts

Planning Domain

- **General** properties of DWR problems
 - There are *containers*, *cranes*, ...
 - Each object has a *location*
 - Possible actions:
Pick up container, put down container, drive to location, ...

Problem Instance

- **Specific** problem to solve
 - Which containers and cranes exist?
 - Where is everything?
 - Where *should* everything be?
(More general:
What should we achieve?)

The State Space

State Spaces 1: Introduction

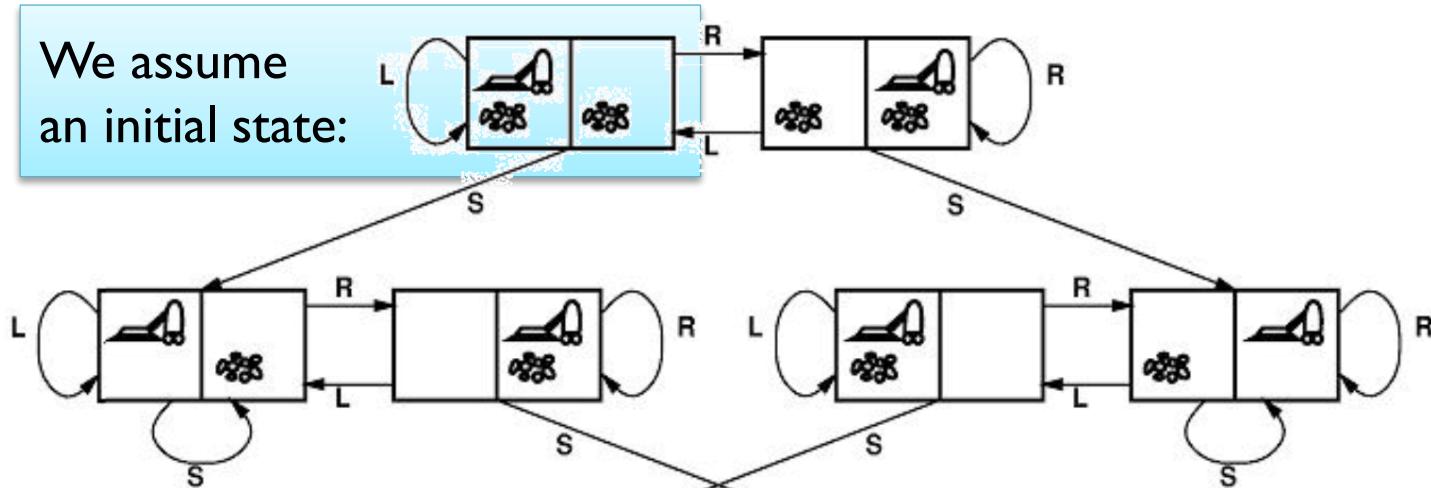
30

jonkv@ida

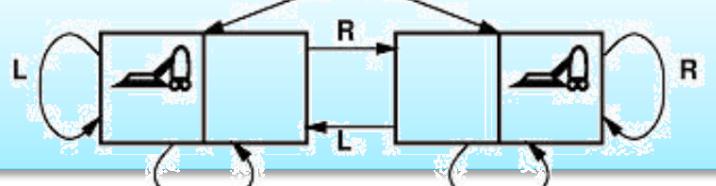
- Every classical planning problem has a state space – a **graph**
 - A node for every world state
 - An edge for every executable action

The planning problem: Find a path (not necessarily shortest)

We assume
an initial state:



And a number of
goal states ("no dirt"):



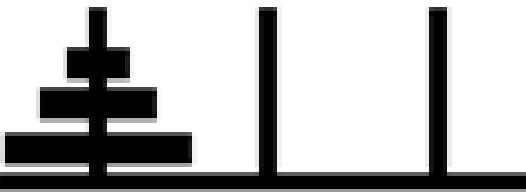
Example solutions: SRS, RSLS, LRLRLSSSRLRS, ...

State Spaces 2: Our Intuitions



- Our intuitions often identify states that we think are:

- "Normal"
- "Expected"
- "Physically possible"



- Usually:

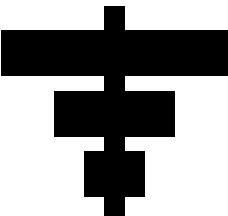
- The initial state is one of those states
- Mainly need to care about all states reachable from there (using the defined actions) – discussed later

State Spaces 3: Against our Intuitions



- But now that we introduced **facts**:

- **Every** combination of **facts** is a state
 - Some are "forbidden"
 - Towers of Hanoi:

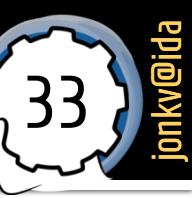


- Some are "counter-intuitive" – in Blocks World, there are states where:
 - $\text{on}(A,A)$ is true
 - $\text{holding}(A)$ and $\text{ontable}(A)$ are true at the same time

These are like
"variables" that can
independently be
true or false!

State Spaces 4: ToH, Actions

33

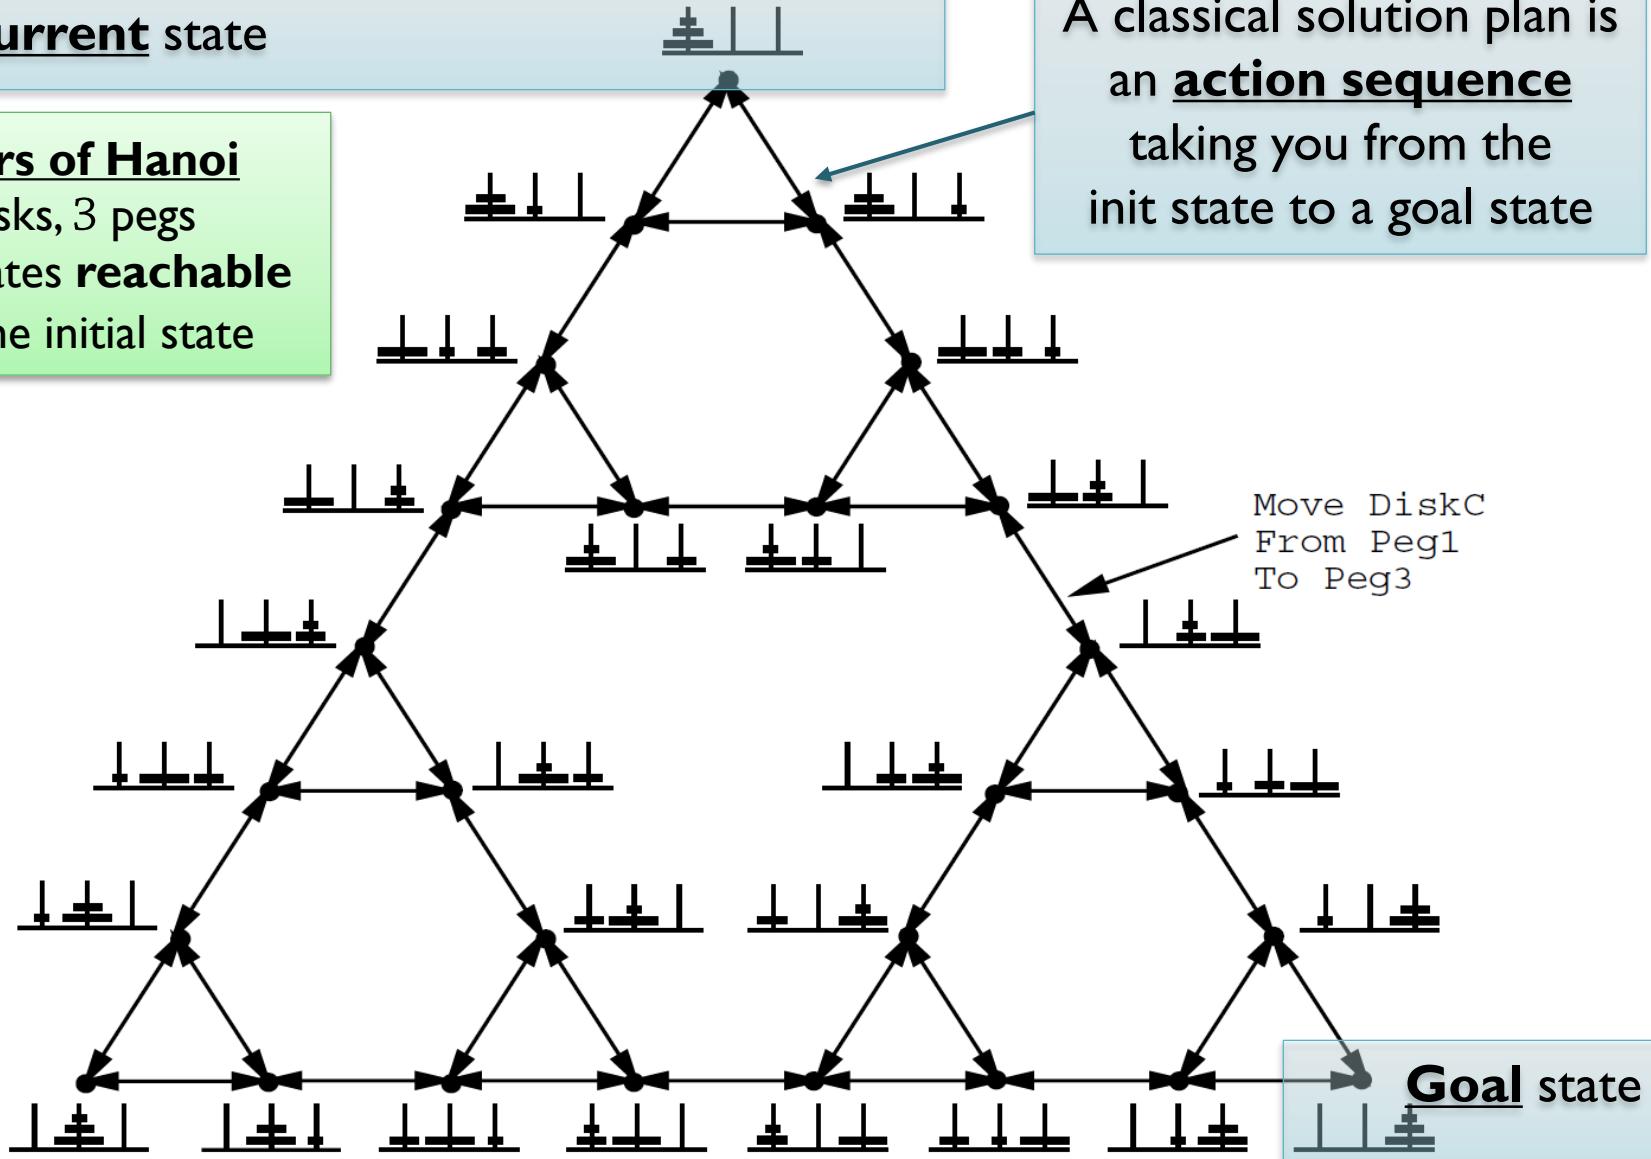


Initial/current state

Towers of Hanoi

3 disks, 3 pegs

→ 27 states **reachable**
from the initial state



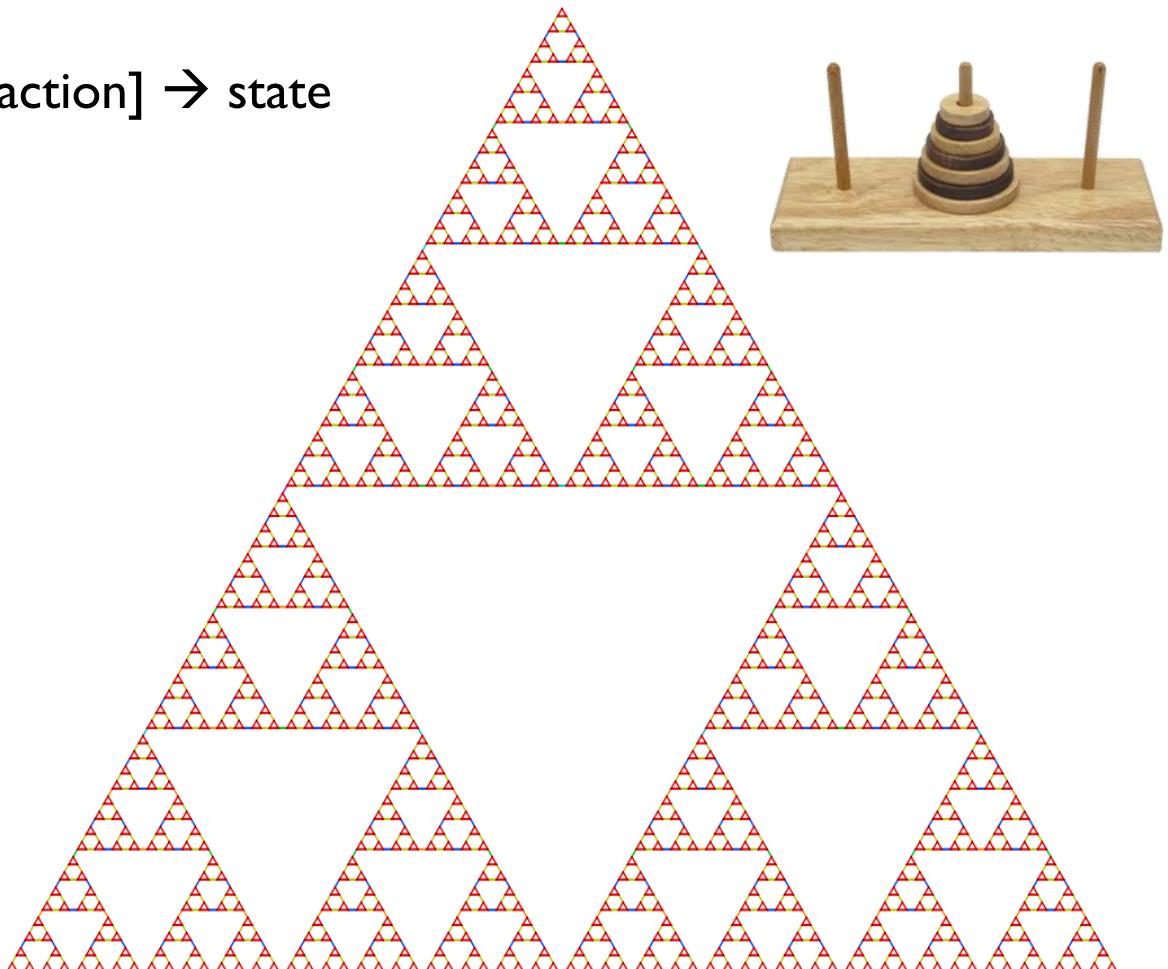
State Spaces 4: Larger Example

34

jonkv@ida

- Larger state space – interesting symmetry

- 7 disks
- 2187 “possible” states
- 6558 transitions, [state, action] → state



State Spaces 5: Blocks World



- A common blocks world version, with 4 operators

- pickup(x) – takes x from the table
- putdown(x) – puts x on the table
- unstack(x, y) – takes x from on top of ?y
- stack(x, y) – puts x on top of y

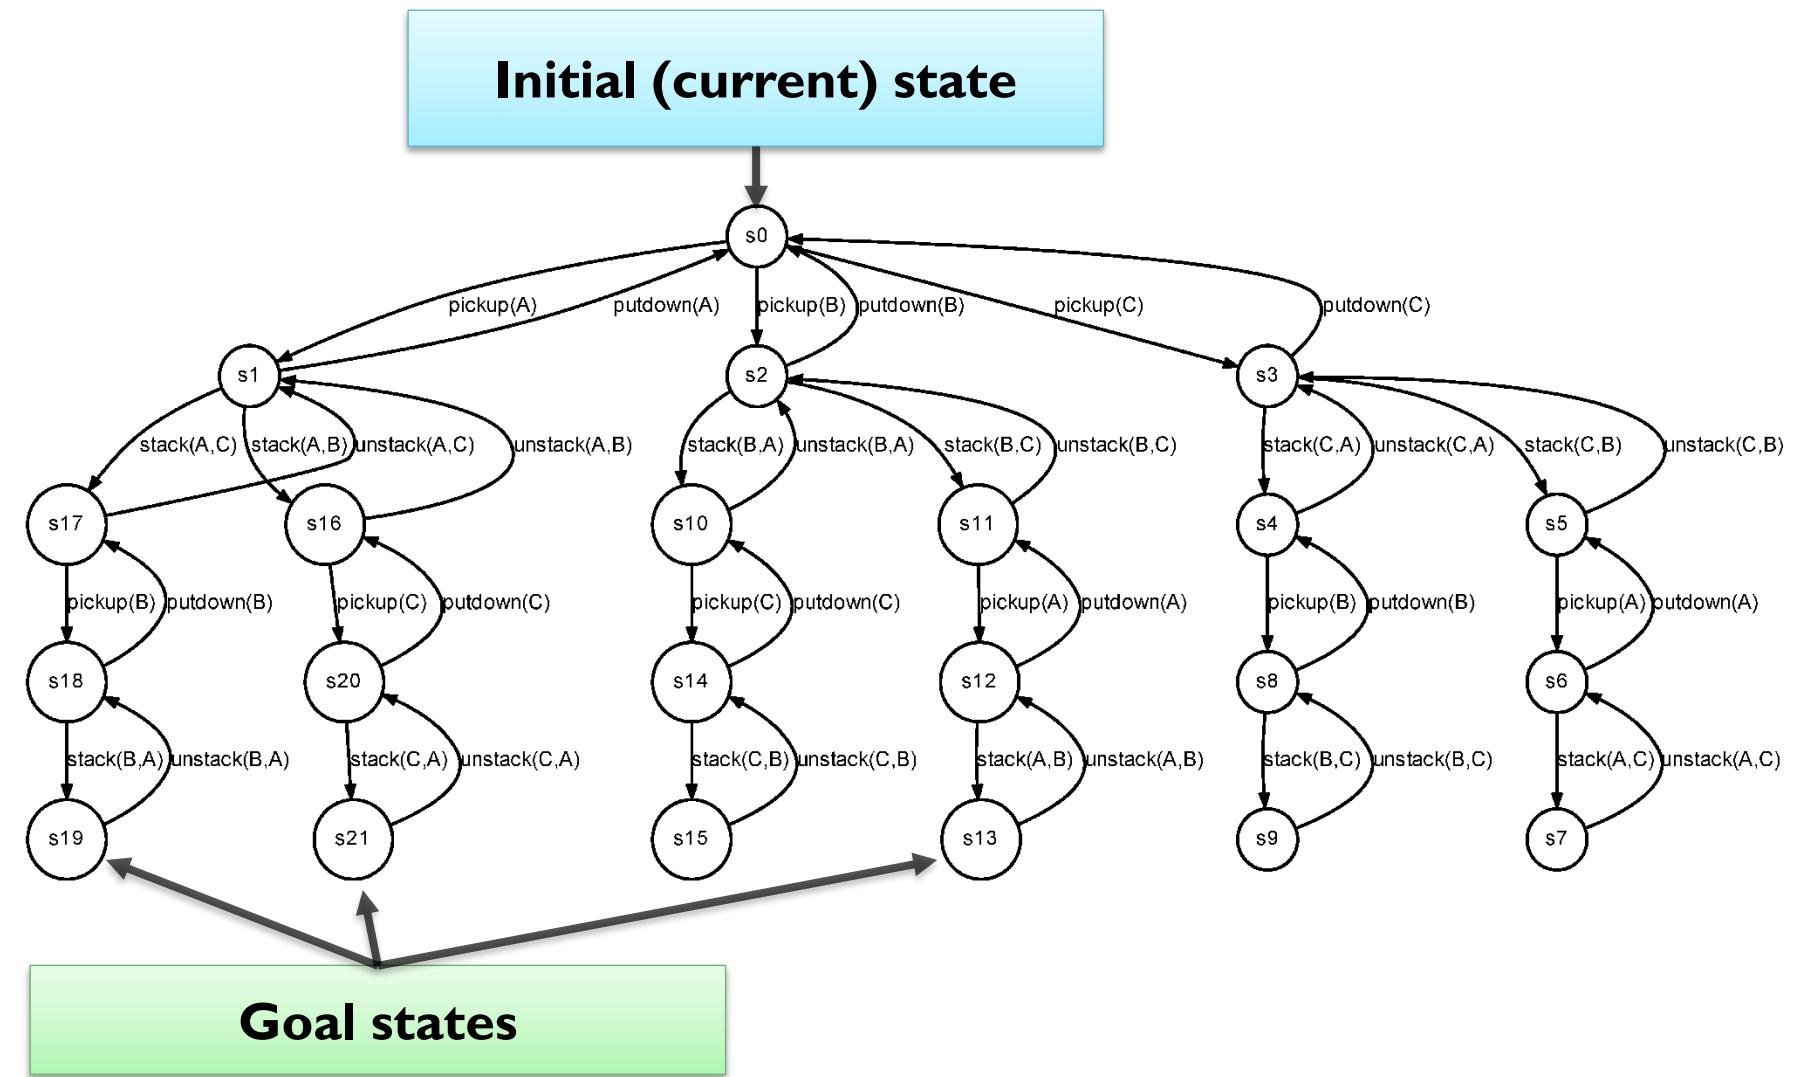


- Predicates (relations) used:

- on(x, y) – block x is on block y
- ontable(x) – x is on the table
- clear(x) – we can place a block on top of x
- holding(x) – the robot is holding block x
- handempty – the robot is not holding any block

clear(A)
on(A, C)
ontable(C)
clear(B) ontable(B)
clear(D) ontable(D)
handempty

State Spaces 6: Blocks World, 3 blocks



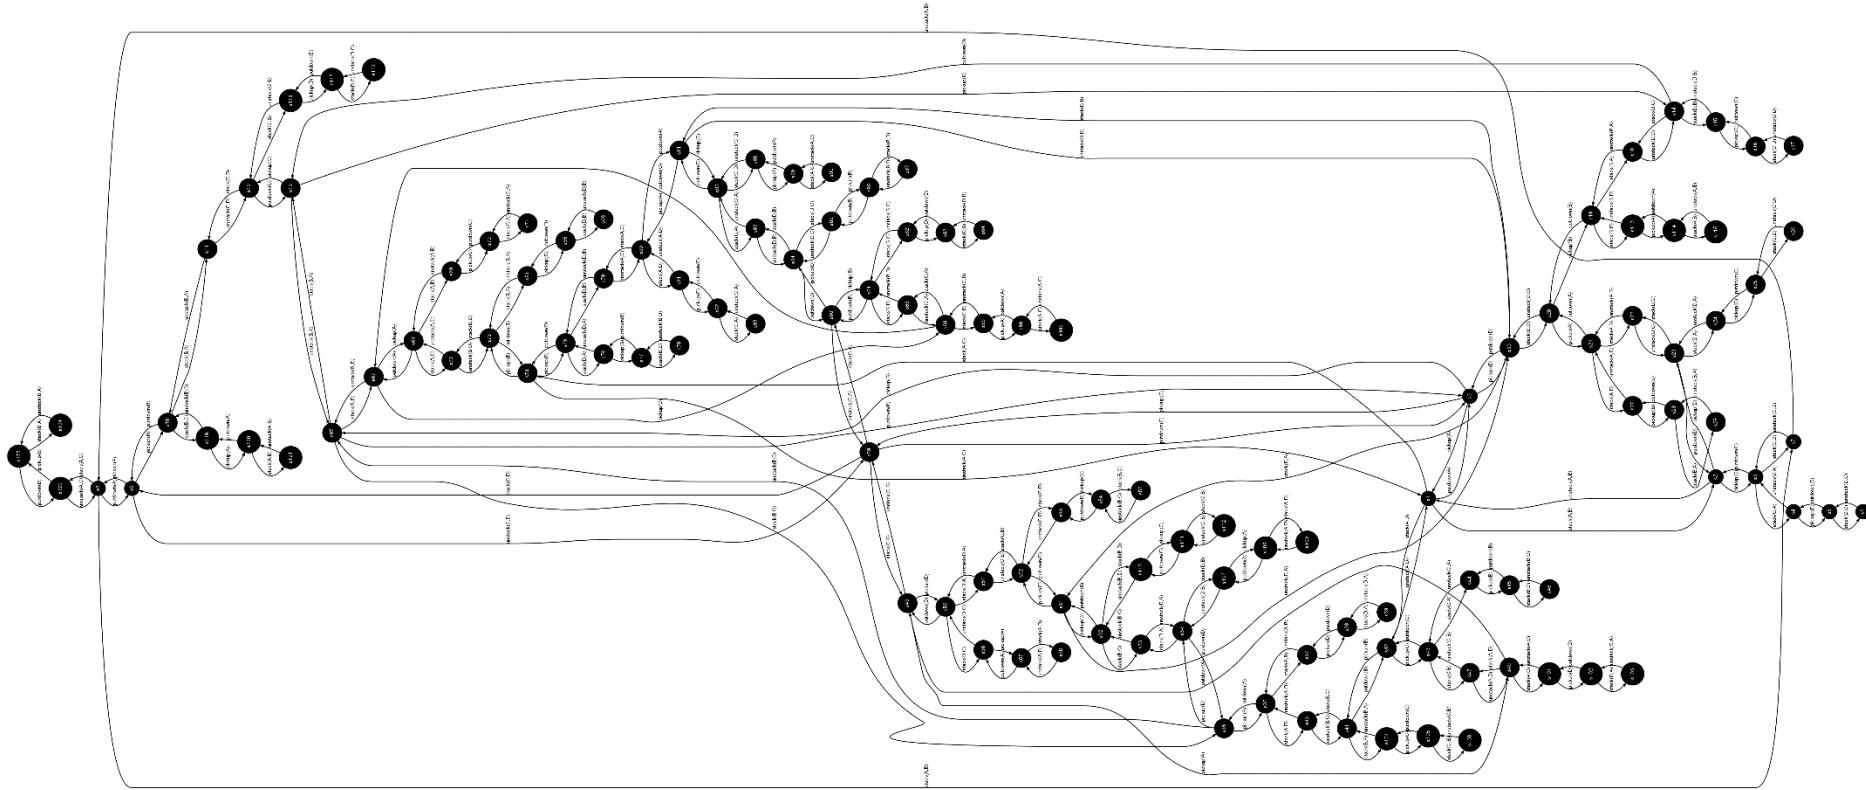
State Spaces 7: Blocks World, 4 blocks



37

jonkv@ida

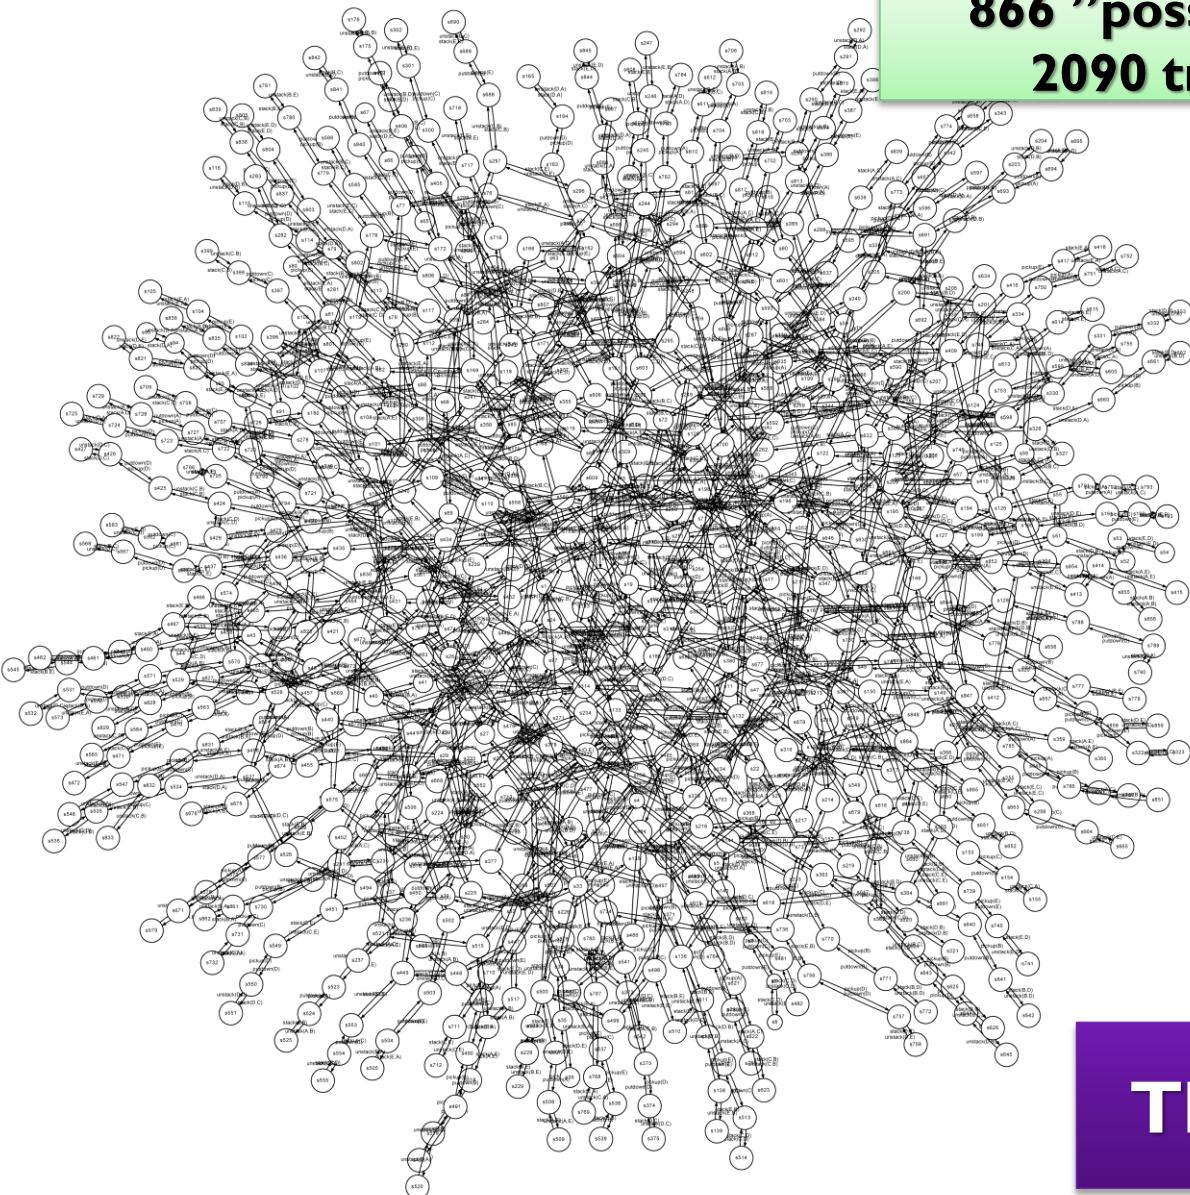
125 "possible" states
272 transitions



State Spaces 8: Blocks World, 5 blocks

38

jonkv@ida



**866 "possible" states
2090 transitions**

This is tiny!

State Spaces 9: Reachable States



Blocks	States reachable from "all on table"	Transitions (edges) in reachable part
0	1	0
1	2	2
2	5	8
3	22	42
4	125	272
5	866	2090
6	7057	18552
7	65990	186578
8	695417	2094752
9	8145730	25951122
10
...30	>197987401295571718915006598239796851	

Plan Generation Method 1: Forward State Space Search

Forward Search 1

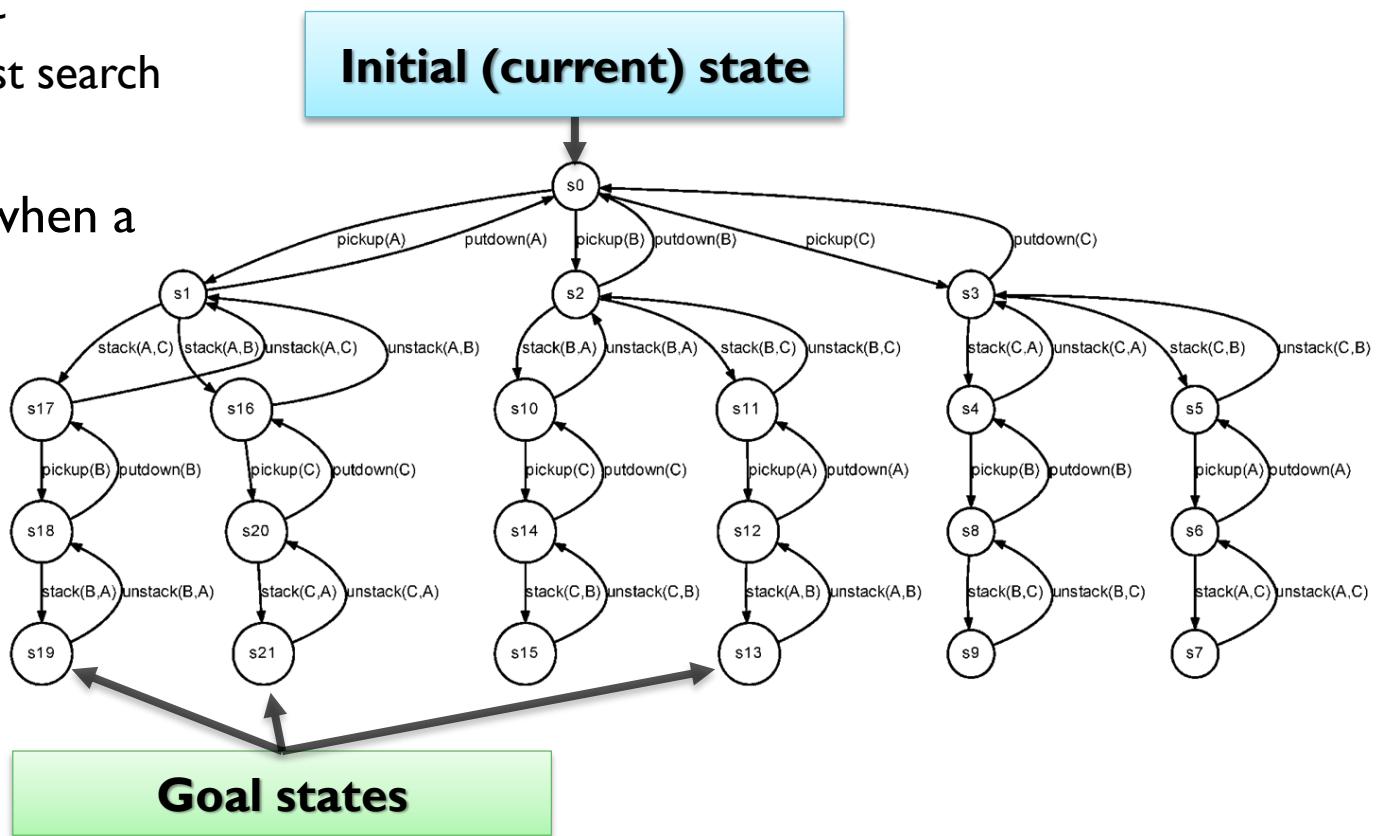
41

jonkv@ida

- Straight-forward planning: **Forward search in the state space**

- Start in the initial state
- Apply a **search** algorithm
 - Depth first
 - Breadth first
 - Uniform-cost search
 - ...

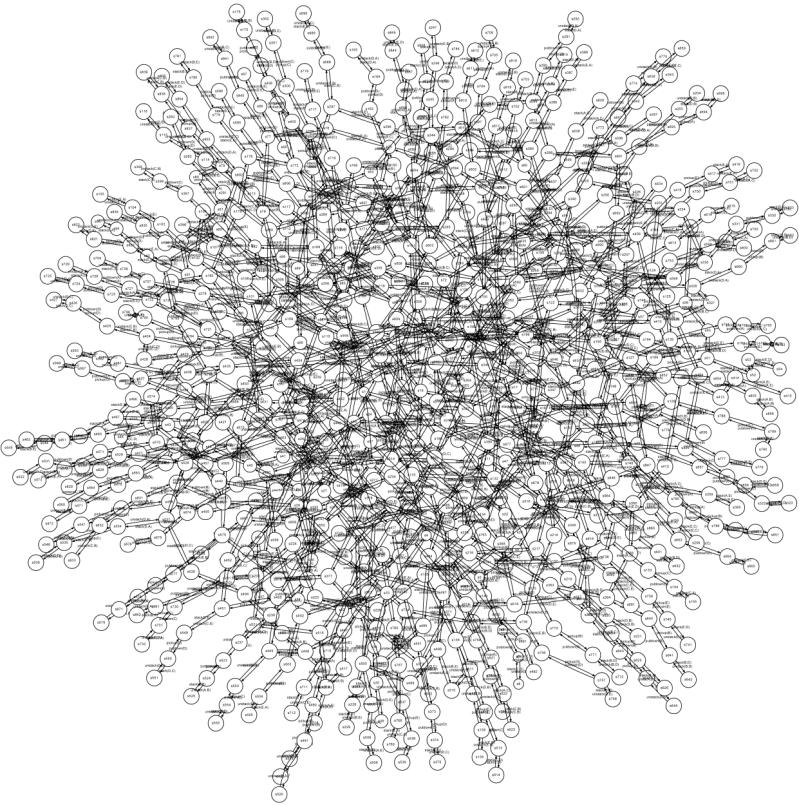
- Terminate when a goal state is found



Forward Search 2: Don't Precompute



- The planner is not given a complete precomputed search graph!



Usually too large!
→ Generate as we go,
hope we don't actually need the *entire* graph

Forward Search 3: Initial state

43

jonkv@ida

- The user (robot?) observes the current state of the world
 - The *initial* state



- Must describe this using the specified formal state syntax...
 - $s_0 = \{ \text{clear}(A), \text{on}(A,C), \text{ontable}(C), \text{clear}(B), \text{ontable}(B), \text{clear}(D), \text{ontable}(D), \text{handempty} \}$
- ...and give it to the planner, which creates one search node

{ **clear(A), on(A,C), ontable(C),**
clear(B), ontable(B), clear(D), ontable(D), handempty } }

Forward Search 4: Successors

44

jonkv@ida

- Given any search node...

```
{ clear(A), on(A,C), ontable(C),
  clear(B), ontable(B), clear(D), ontable(D), handempty }
```

- ...we can find successors – by applying actions!

- action** pickup(D)

- Precondition: **ontable(D) \wedge clear(D) \wedge handempty**

- Effects: $\neg\text{ontable}(D) \wedge \neg\text{clear}(D) \wedge \neg\text{handempty} \wedge \text{holding}(D)$

- This generates new reachable states...

...which can also
be illustrated

```
{ clear(A), on(A,C), ontable(C),
  clear(B), ontable(B), clear(D), ontable(D), handempty }
```

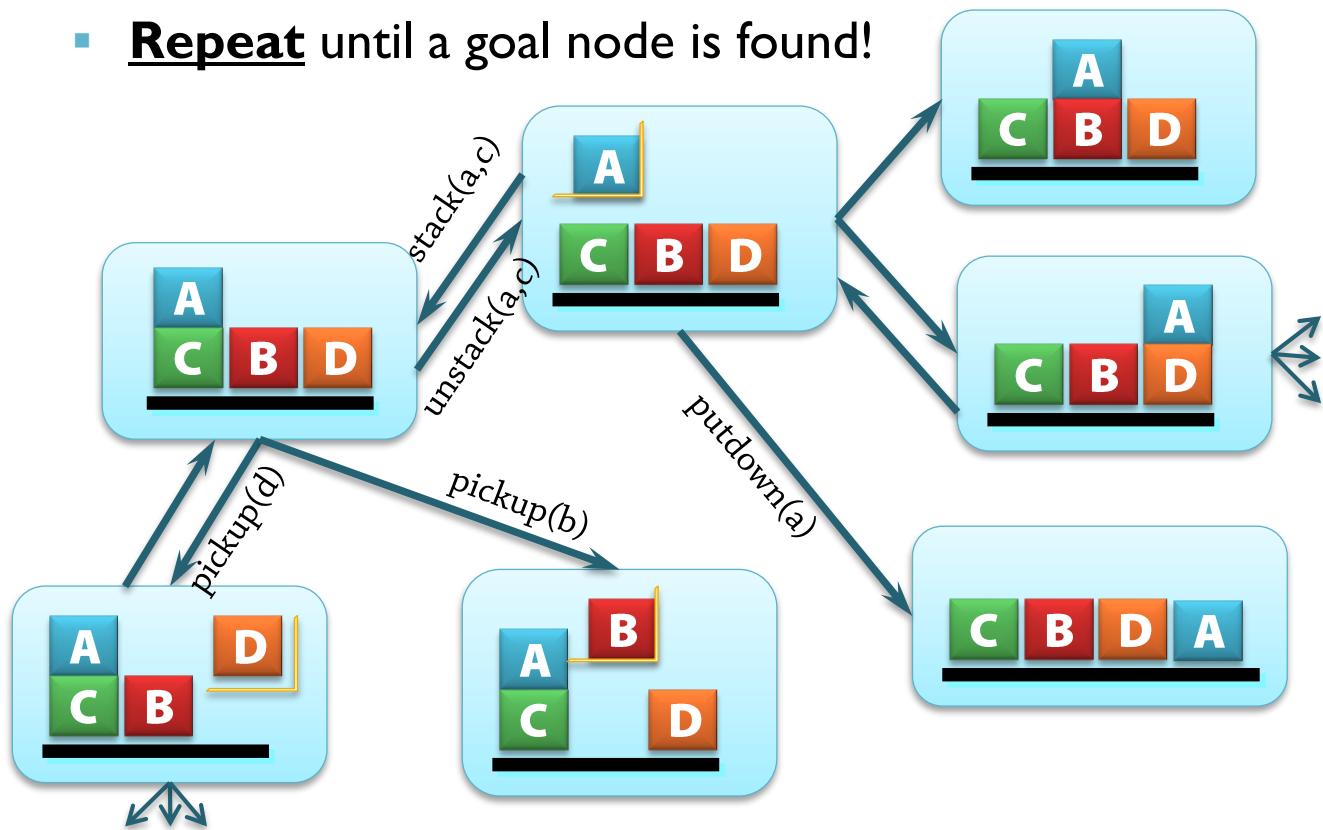
```
{ clear(A), on(A,C), ontable(C),
  clear(B), ontable(B), holding(D) }
```



Forward Search 5: Step by step

- A search strategy (depth first, A*, hill climbing, ...) will:

- Choose a node
- Expand the node, generating all possible successors
 - “What actions are applicable in the current state, and where will they take me?”
 - Generates new states by applying effects
- Repeat until a goal node is found!



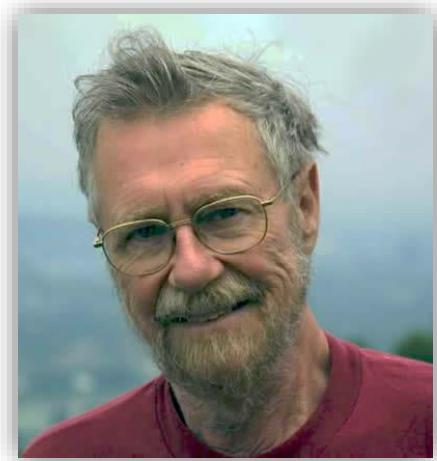
This is *illustrated* –
the planner works
with sets of facts

The blocks world is
symmetric: Can
always “return the
same way”
Not true for all
domains!

Uninformed Forward State Space Search

- One possible strategy: **Dijkstra's algorithm**

- Strategy can be described as *nearest-first*:
*“take an unexplored node
that can be reached at minimal cost from the initial node”*
 - **Efficient**: $O(|E| + |V| \log |V|)$
 - $|V|$ = the number of nodes
 - $|E|$ = the number of edges
 - And generates **optimal** (cheapest) paths/plans!



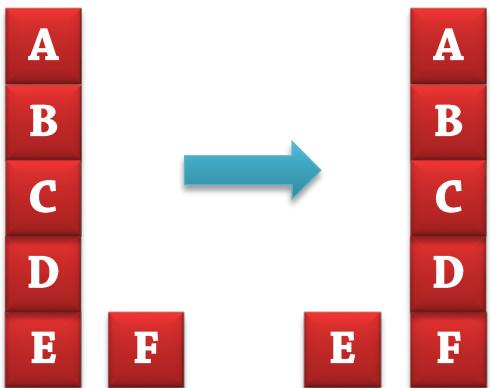
Would this algorithm be a solution?

Dijkstra's Algorithm: Example

48

jonkv@ida

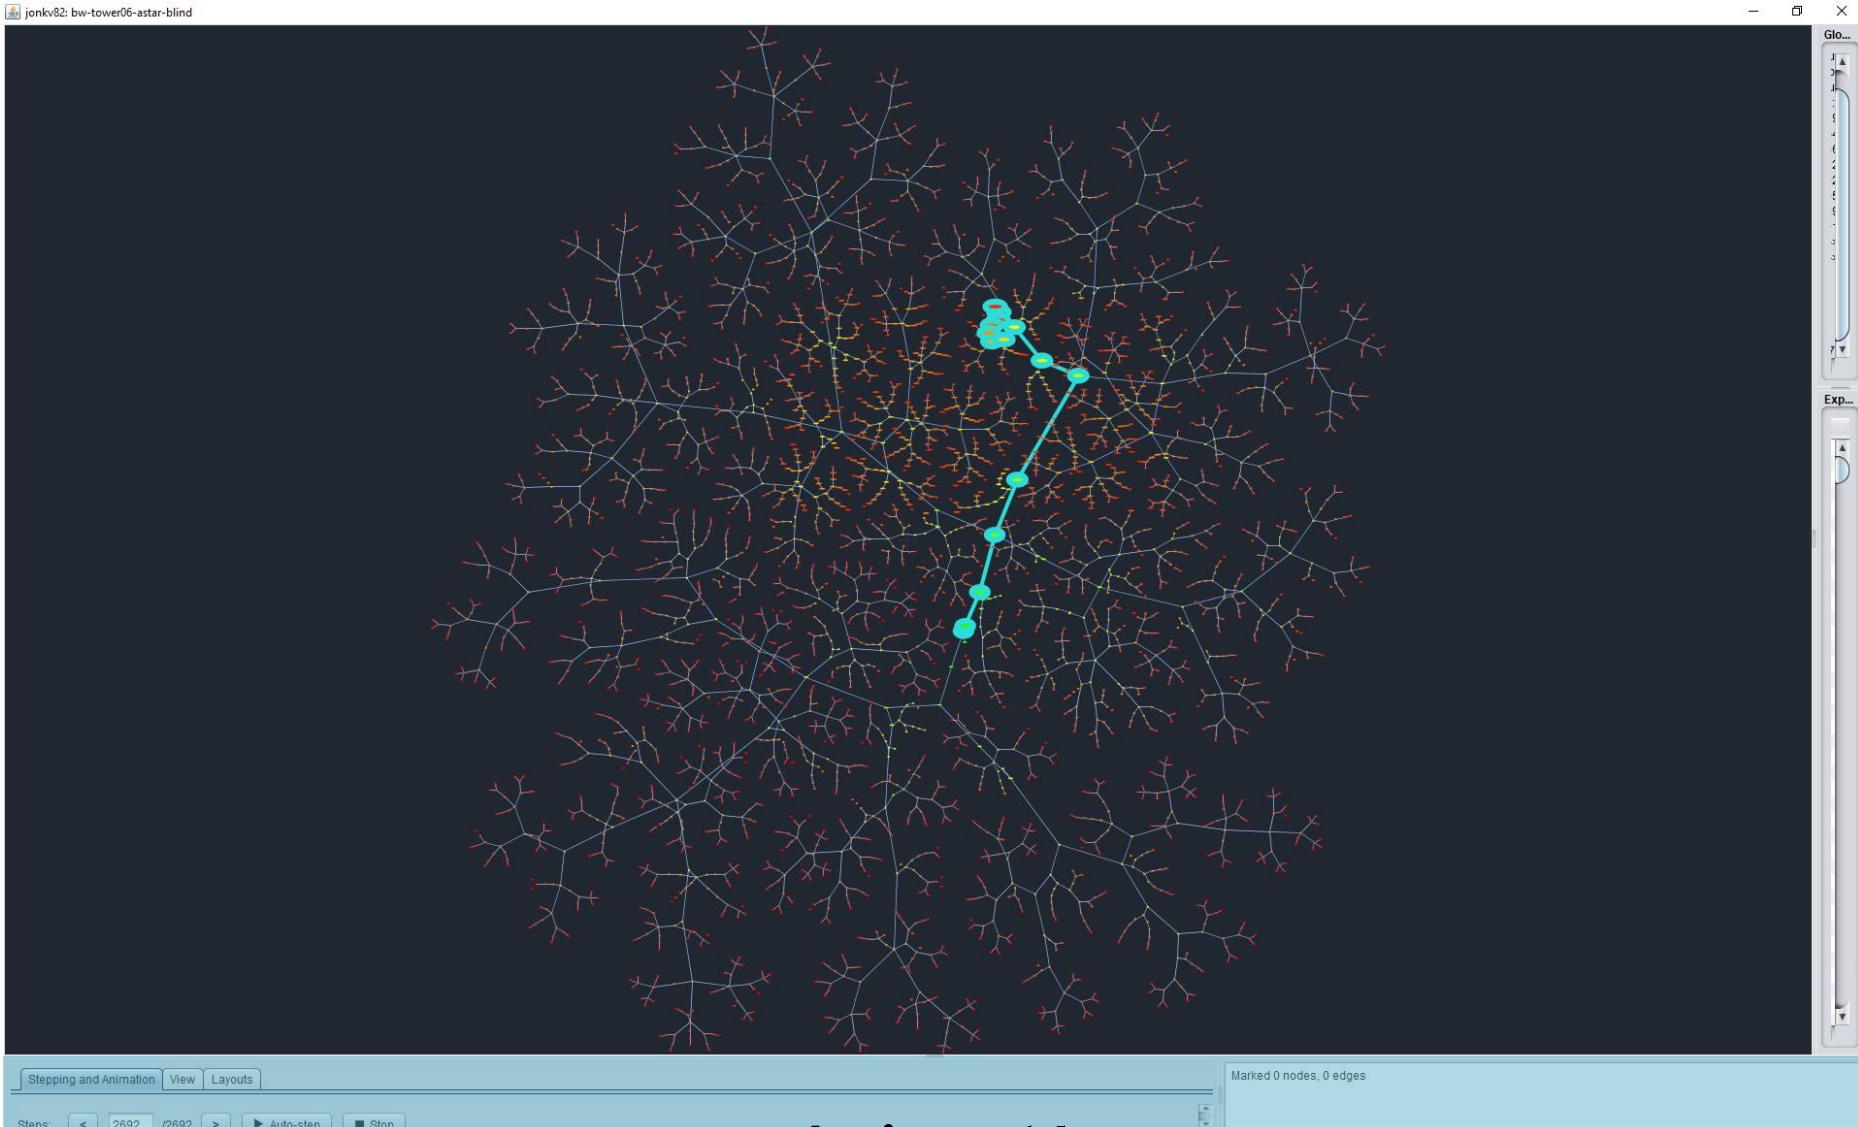
- A simple problem:



Goal
on(A,B)
on(B,C)
on(C,D)
on(D,F)
ontable(E)
ontable(F)

Optimal solution	
unstack(A,B)	pickup(D)
putdown(A)	stack(D,F)
unstack(B,C)	pickup(C)
putdown(B)	stack(C,D)
unstack(C,D)	pickup(B)
putdown(C)	stack(B,C)
unstack(D,E)	pickup(A)
stack(D,F)	stack(A,B)

bw-tower06-dijkstra: Only 6 blocks, Dijkstra search, no heuristic



Actions: 14

States: 8706 calculated, 2692 visited

Dijkstra's Algorithm: Analysis



- Blocks world, 400 blocks initially on the table, goal is a 400-block tower
 - Given that all actions have the same cost,
Dijkstra will first consider all plans that stack **less than 400 blocks!**
 - Stacking 1 block: = 400×399 plans, ...
 - Stacking 2 blocks: > $400 \times 399 \times 398$ plans, ...
 - More than

163056983907893105864579679373347287756459484163478267225862419762304263994207997664258213955766581163654137118
163119220488226383169161648320459490283410635798745232698971132939284479800304096674354974038722588873480963719
240642724363629154726632939764177236010315694148636819334217252836414001487277618002966608761037018087769490614
847887418744402606226134803936935233568418055950371185351837140548515949431309313875210827888943337113613660928
318086299617953892953722006734158933276576470475640607391701026030959040303548174221274052329579637773658722452
5497384594044525865036931693410912754853265795909113444084441755664211796
27432025699299231777374983037488265744484456318793090777961572990289194
810585217819146476629300233601372350568748665249021991849760646988031691
394386551194171193333144031541302649432305620215568850657684229678385177
7253589339861121273524529880313087201742432360729162527387508073225578630
777685901637435541458440833878709544174983977457450527557534417629122448835191721077333875230695681480990867109
051332104820413607822206465635272711073906611800376194410428900071013695438359094641682253856394743335678545824
320932106973317498515711006719985304982604755110167254854766188619128917053933547098435020659778689499606904157
077005797632287669764145095581565056589811721520434612770594950613701730879307727141093526534328671360002096924
483494302424649061451726645947585860104976845534507479605408903828320206131072217782156434204572434616042404375
21105232403822580540571315732915984635193126556273109603937188229504400

1.63 * 10¹⁷³⁵

Efficient in terms of the **search space size**: $O(|E| + |V| \log |V|)$

The search space is **exponential** in the size of the input description...

Fast Computers, Many Cores



- But computers are getting very fast!
 - Suppose we can check 10^{20} states per second
 - >10 billion states per clock cycle for today's computers, each state involving complex operations
 - Then it will only take $10^{1735} / 10^{20} = 10^{1715}$ seconds...

■ But we have multiple cores!

- The universe has at most 10^{87} particles, including electrons, ...
- Let's suppose every one is a CPU core
- → only 10^{1628} seconds
 $> 10^{1620}$ years
- The universe is around 10^{10} years old



Hopeless? No: We need informed search!

Informed Forward State Space Search

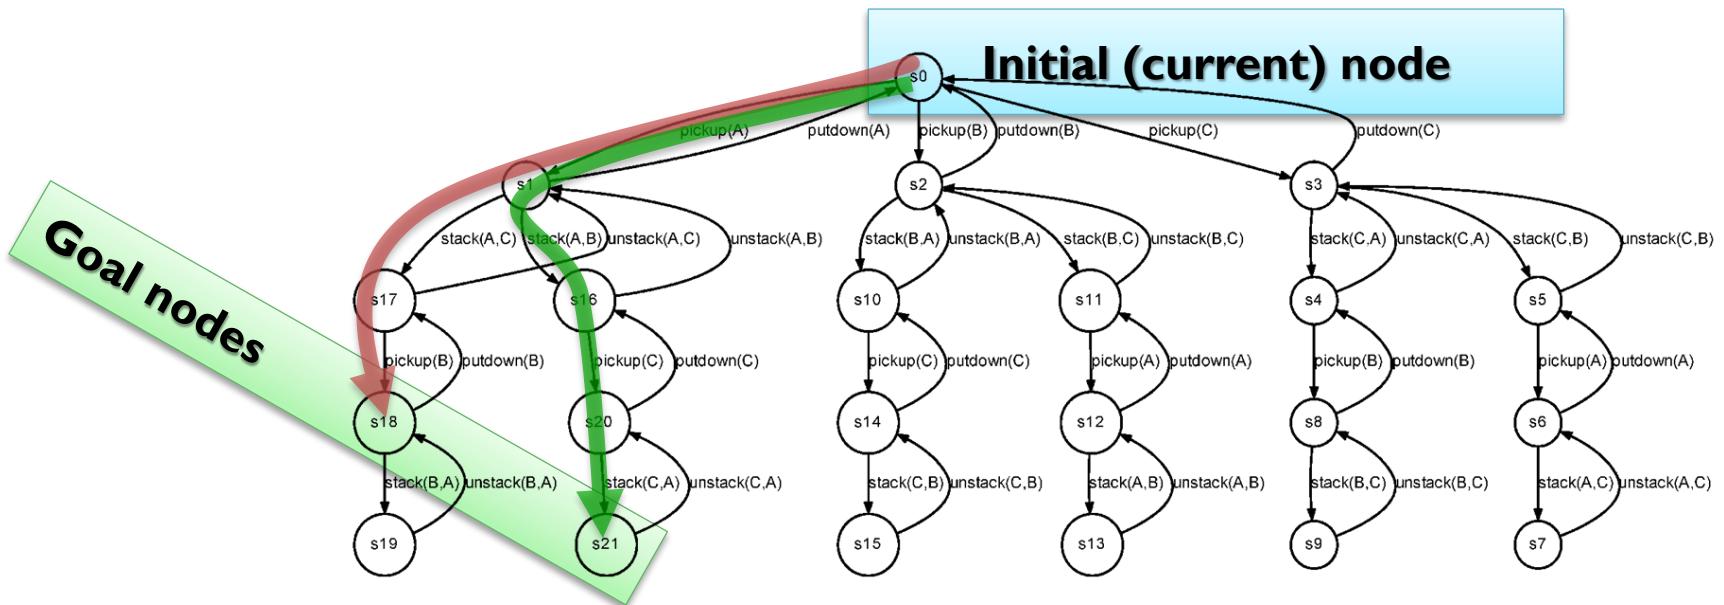
Intro (1)

53

jonkv@ida

We need:

- A heuristic function $h(n)$ estimating the cost of reaching a goal node from node n
 - Sometimes, we define **cost = number of actions**
 - More general: **each action has a cost $c(a)$ – longer plans may be cheaper!**



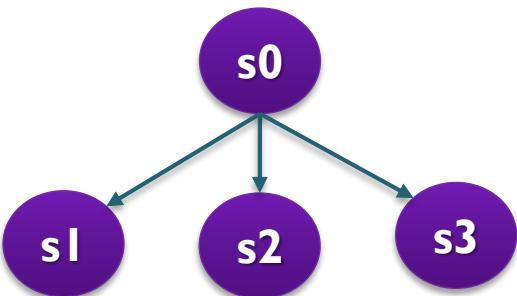
Intro (2)

- To use the heuristic, we will typically:

- Pick some state (let's say s_0):



- Expand it one step:



- Compute $h(n)$ for all the expanded nodes
 - $h(s_1), h(s_2), h(s_3)$
 - Apply some search strategy using $h(n)$ as part of its selection criteria, to:

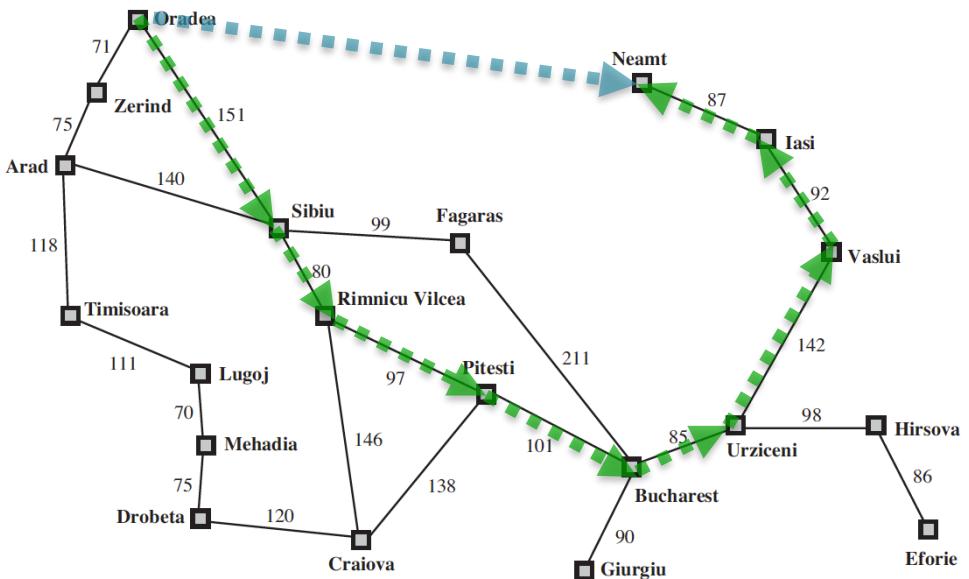


Intro (3)



- Previously we manually adapted heuristics to the problem
 - 8-puzzle → # pieces out of place, or sum of Manhattan distances
 - Romania Travel → straight line distance

7	2	4
5		6
8	3	1



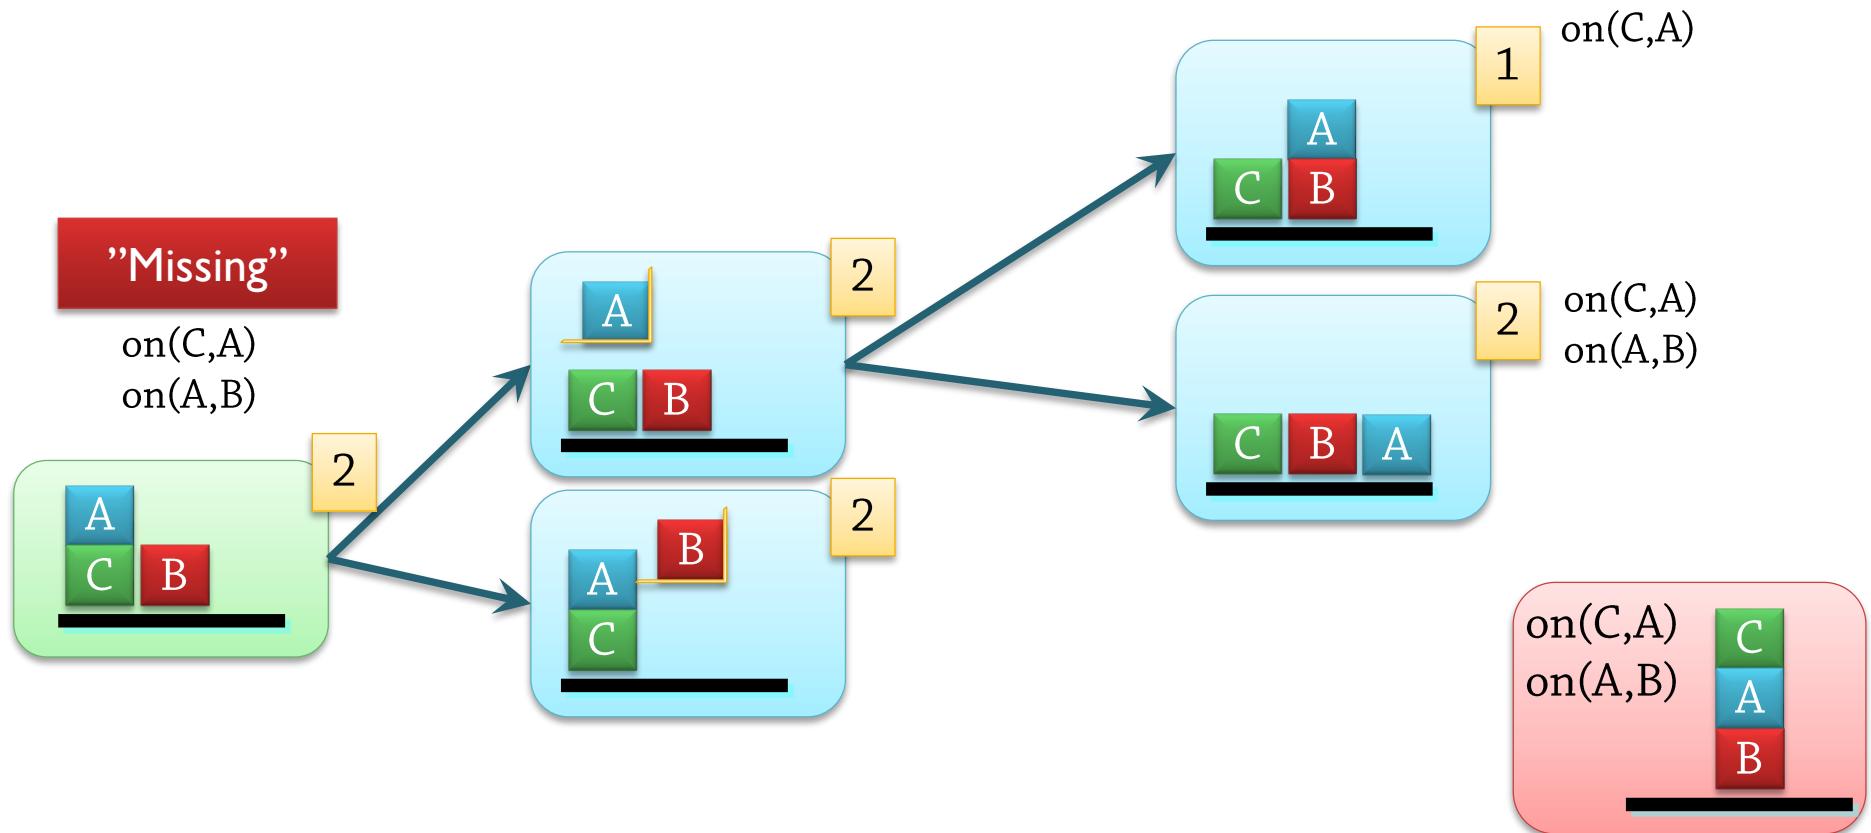
- Now: Want a general heuristic function
 - Without knowing what planning problem is going to be solved!

Domain-Independent Heuristics (1)



- A very simple domain-independent heuristic:
 - Count the number of unachieved goals in the current state
(similar to "number of pieces out of place")

Optimal:
unstack(A,C)
stack(A,B)
pickup(C)
stack(C,A)

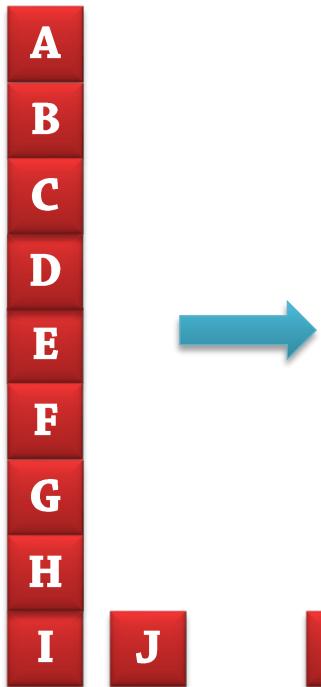


Domain-Independent Heuristics (2)

57

jonkv@ida

- Helpful, but not very...



Goal:

on(A,B)
on(B,C)
on(C,D)
on(D,E)
on(E,F)
on(F,G)
on(G,H)
on(H,J)

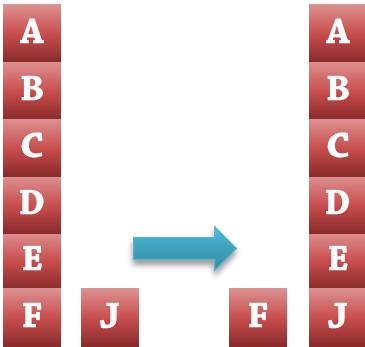
Optimal solution:

unstack(A,B)	pickup(G)
putdown(B)	stack(G,H)
unstack(B,C)	pickup(F)
putdown(C)	stack(F,G)
unstack(C,D)	pickup(E)
putdown(D)	stack(E,F)
...	...
unstack(H,I)	
stack(H,J)	

Only unachieved goal fact:
 $\text{on}(H,J)$

Must remove many
"good facts"
to reach the goal

bw-tower07-astar-gc: Only 7 blocks, A* search, based on goal count

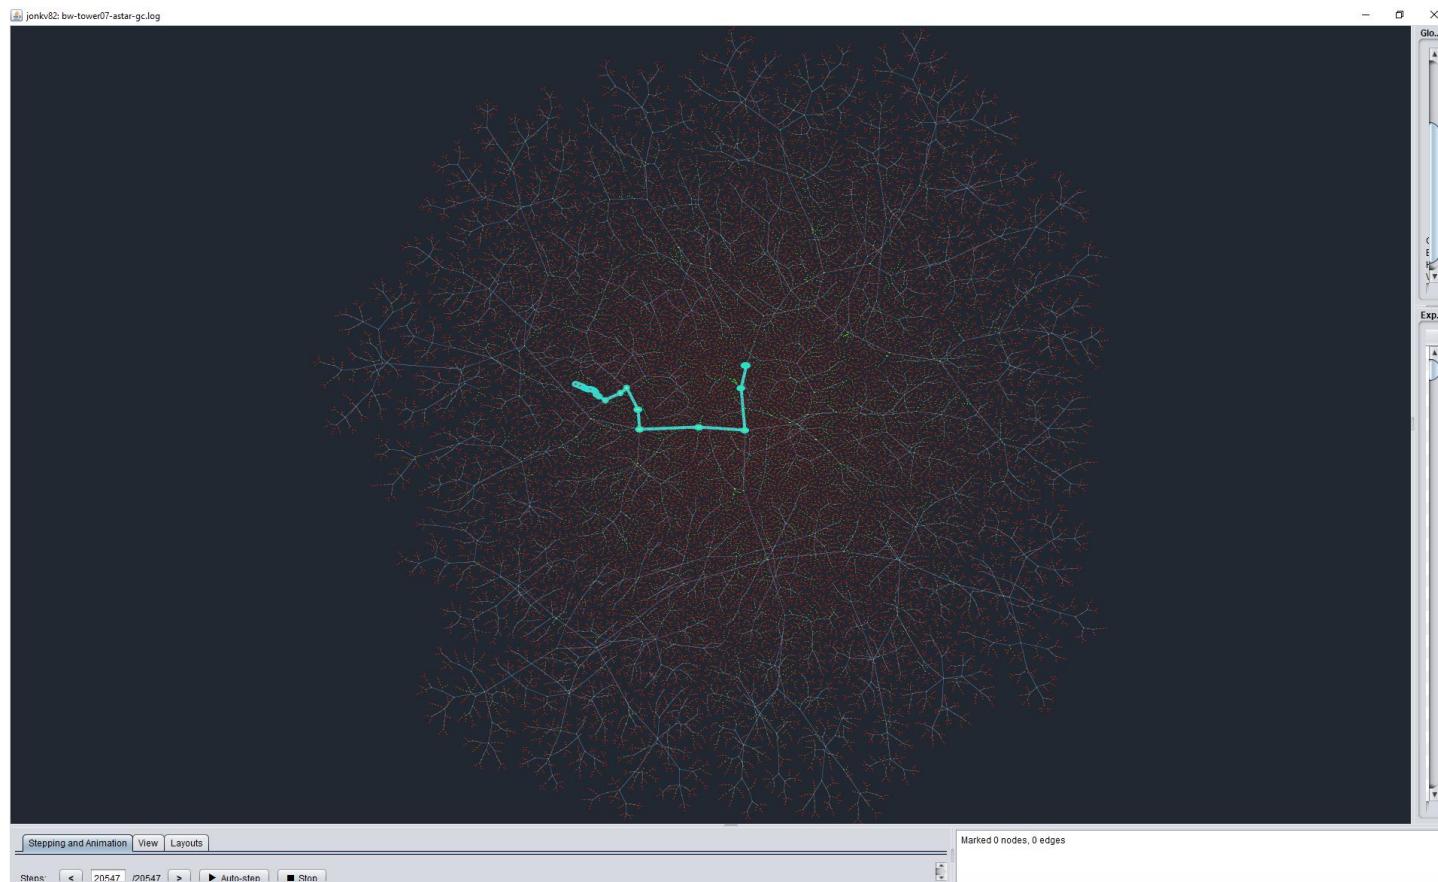


18 actions

States:

**6463 calculated,
3222 visited**

**(With Dijkstra,
43150 / 33436 –
improved, but we
can do better!)**



- $h(s_0) = 1$: Only one “missing” fact
- For a long time, all useful successors appear to increase remaining cost
 - Removing a block that must be moved
- And many useless successors appear to decrease remaining cost
 - Building towers that will need to be torn down

Optimal Classical Planning

Optimal 1: Introduction



■ Optimal plan generation:

- There is a quality measure for plans
 - Minimal number of actions
 - Minimal sum of action costs
 - ...
- We must find an optimal plan!



- Suboptimal plans
(0.5% more expensive):

Irrelevant

Optimal 2: A*



Optimal Plan Generation: Often uses A*

- A* focuses entirely on optimality
 - Find a **guaranteed optimal** plan as quickly as possible
 - But no point in trying to find a "reasonable" plan before the optimal one
 - Slowly expand from the initial state, systematically checking possibilities
- Requires admissible heuristics to guarantee optimality
 - Reason: Heuristic used for *pruning* (ignoring some search nodes)
 - Search queue ordered by $f = g$ (actual cost) + h (heuristic):

$$11 = 10 + 1$$

Pop – not a solution

$$12 = 10 + 2$$

Pop – not a solution

$$12 = 12 + 0$$

Pop – solution!

$$12 = 11 + 1$$

Ignore:
 g is known, h is an underestimate,
so solutions found by expanding
these nodes will cost $\geq g+h$
(and we have one of cost $\leq g+h$)

$$13 = 11 + 2$$

Creating Admissible Heuristic Functions: The General Relaxation Principle

How does relaxation apply to planning?

How can we make the definition more precise?

The Problem



■ We want:

- A general principle for developing different heuristic functions h(s)
 - Allowing us to explore different ideas for heuristics
- So that the heuristic functions we develop will:
 - Work for *all* classical planning problems, and *all* states s
 - Never overestimate the real cost of reaching the goal: h is admissible
 - Provide useful information – not simply $h(s) = 0$ for all s...

**What do we do?
Where do we start?
How do we think?**

An Obvious Method

64

jonkv@ida

■ One obvious(ish) method:

Every time we need $h(s)$ for some state s in a problem P :

1. Generate an optimal solution plan $\pi^*(s)$ starting in s

2. Let $h(s) = h^*(s) = \text{cost}(\pi^*(s))$
 - Admissible – why?
 - Informative – why?

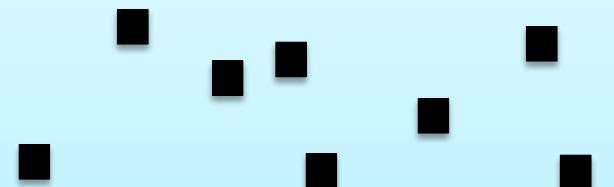
$h^*(s)$ = "the perfect heuristic" =
the actual cost of an optimal plan

■ Obvious, but stupid

- If we find $\pi^*(s)$, we're already done!

Also: These are hard to find
(or we wouldn't need
a heuristic)

Solutions π to P starting in s
(set of plans!)



Optimal solutions $\pi^*(s)$



Transforming a Problem



- Main problem: **performance** – we need something **faster**

Original problem P ,
current state s
(finding a solution: slow)

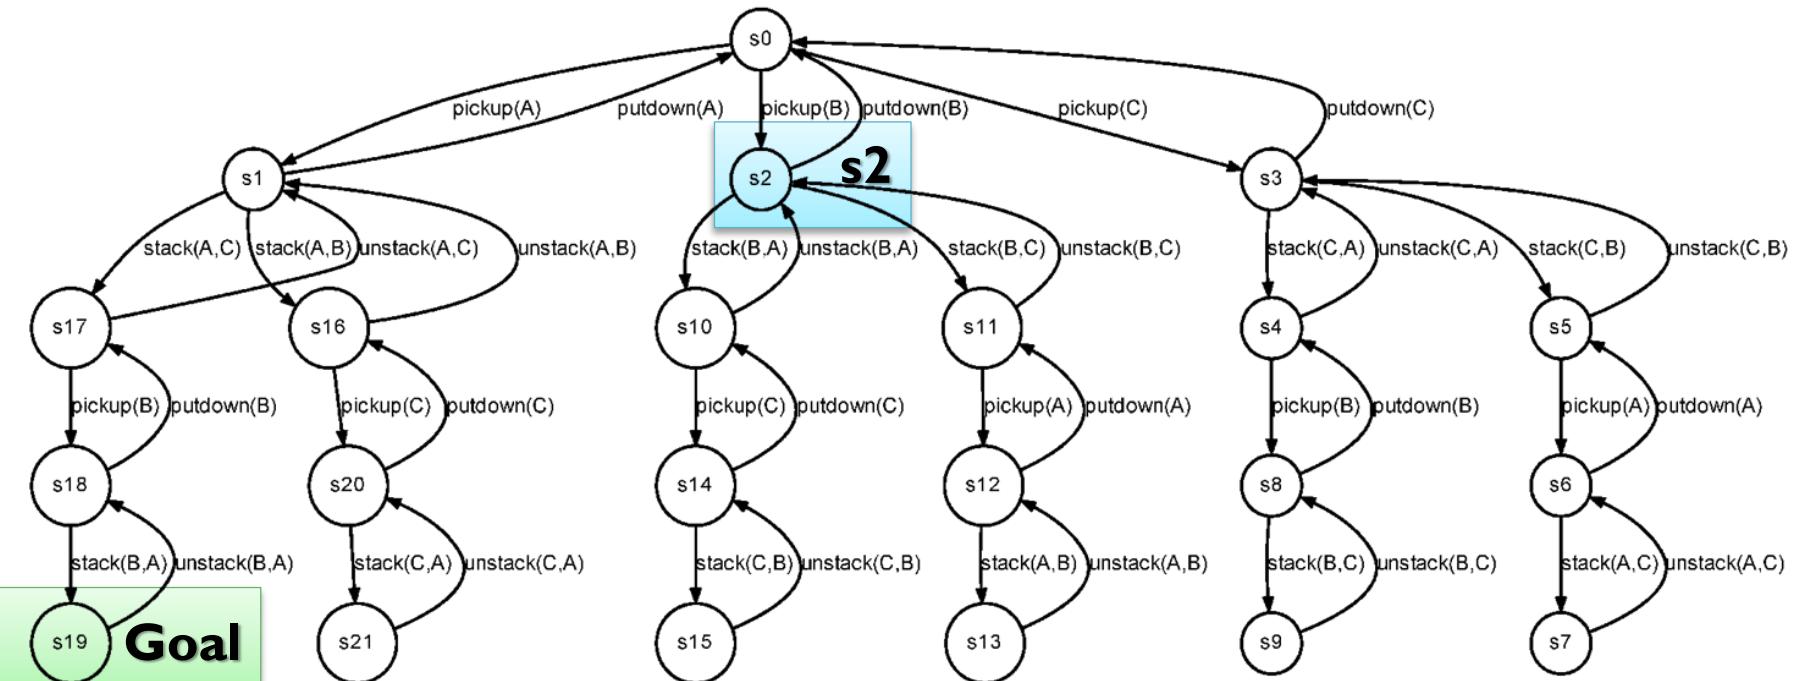


Transformed "simplified"
problem: P' , s'
(finding a solution: fast)

Simplification Example: Basis



- A simple planning problem (domain + instance)
 - Blocks world, 3 blocks
 - **Currently:** in state s2
 - **Goal:** (and (on B A) (on A C)) **(only satisfied in s19)**
 - Solutions from here:
All paths from s2 to goal (infinitely many – can have cycles)



Simplification Example 1

67

jonkv@ida

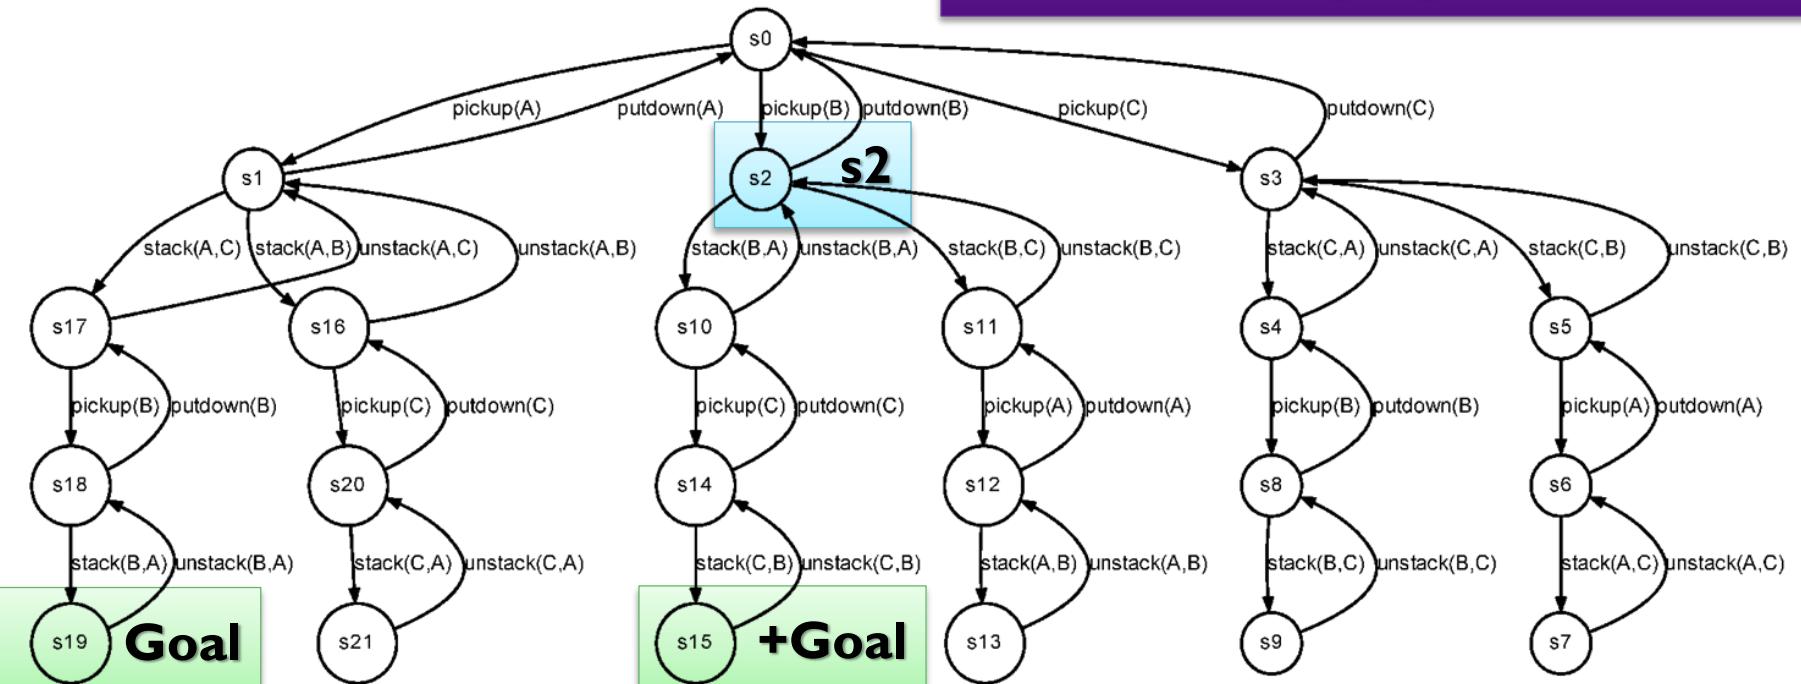
■ Adding goal states

- New goal formula: (and (on B A) **(or (on A C) (on C B))**)

- Informally:

More goal states → should be easier to find one from s_2 , on average

May need more transformations:
Just showing a general idea!



Computing a Heuristic Value



- **What do we do** with the transformed problem?

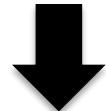
Original problem P ,
current state s
(finding a solution: slow)



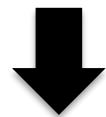
Transformed problem:
 P' , s'
(finding a solution: fast)



Still admissible? Yes, if
 $\text{cost}(\pi^*(P', s')) \leq \text{cost}(\pi^*(P, s))$
(The optimal plan we **found** for P'
was not more expensive than the
actual cost for P)



Quickly find $\pi^*(P', s')$ –
optimal plan for
transformed problem



Let $h(s) = c$



Compute
 $c = \text{cost}(\pi^*(P', s'))$

Simplification Example 1

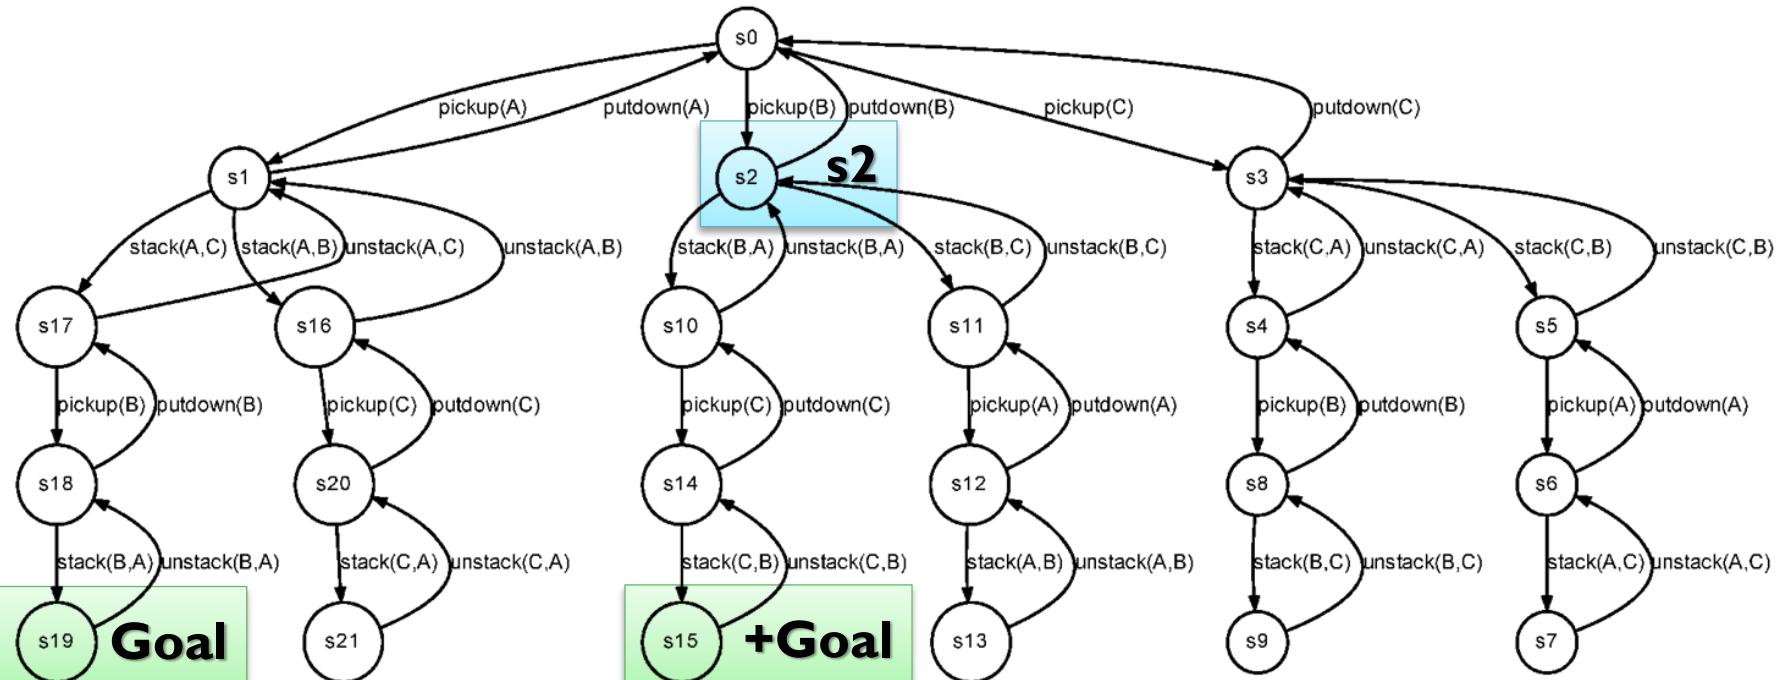
69

jonkv@ida

■ Adding goal states

- Can adding goal states make optimal plans more expensive?
- No, only cheaper from some states

How can we generalize this result?

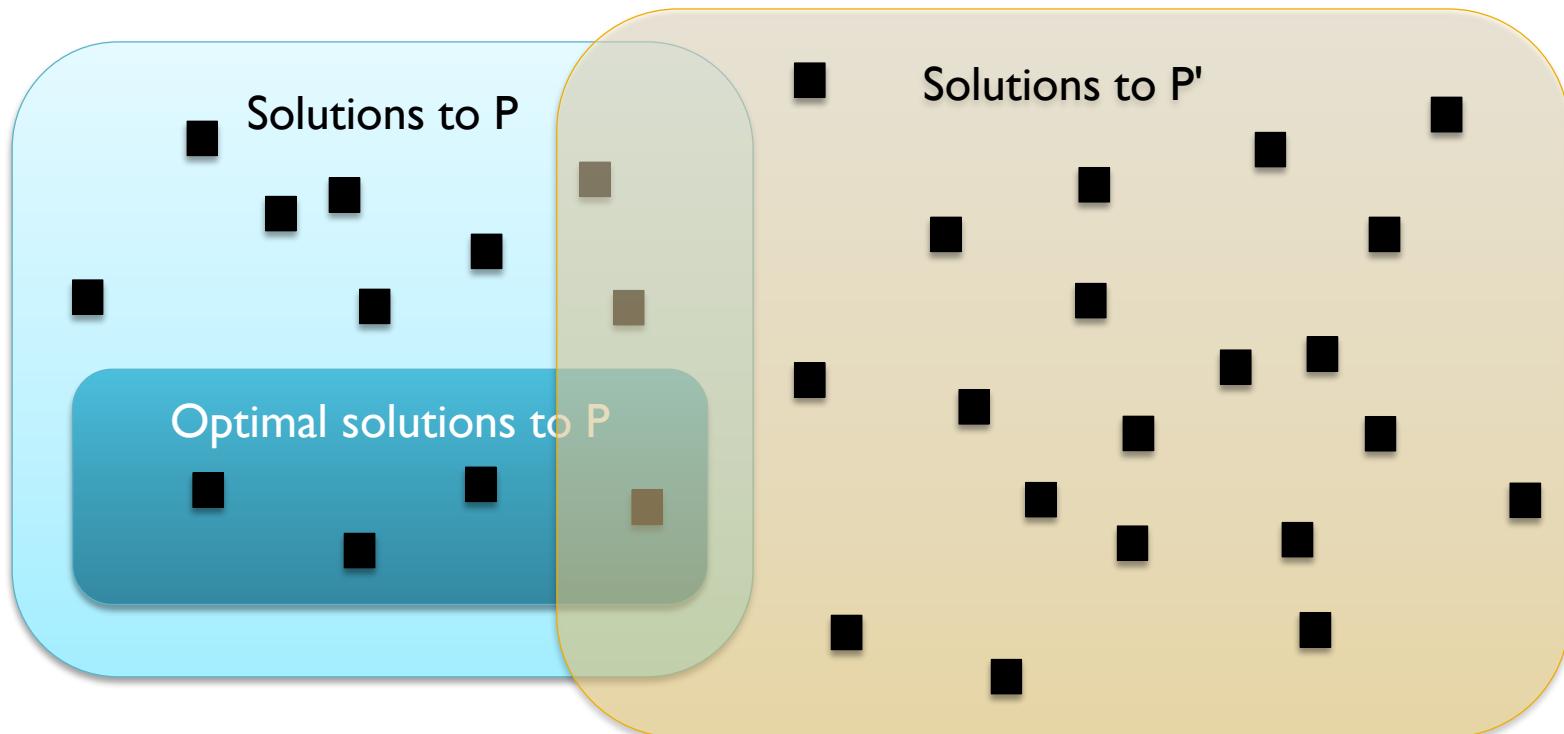


Solution Sets (1)



- How to prove $\text{cost}(\text{optimal-solution}(P')) \leq \text{cost}(\text{optimal-solution}(P))$?
 - Sufficient criterion: One optimal solution to P remains a solution for P'
 - $\text{cost}(\text{optimal-solution}(P')) = \min \{ \text{cost}(\pi) \mid \pi \text{ is any solution to } P' \} \leq \text{cost}(\text{optimal-solution}(P))$

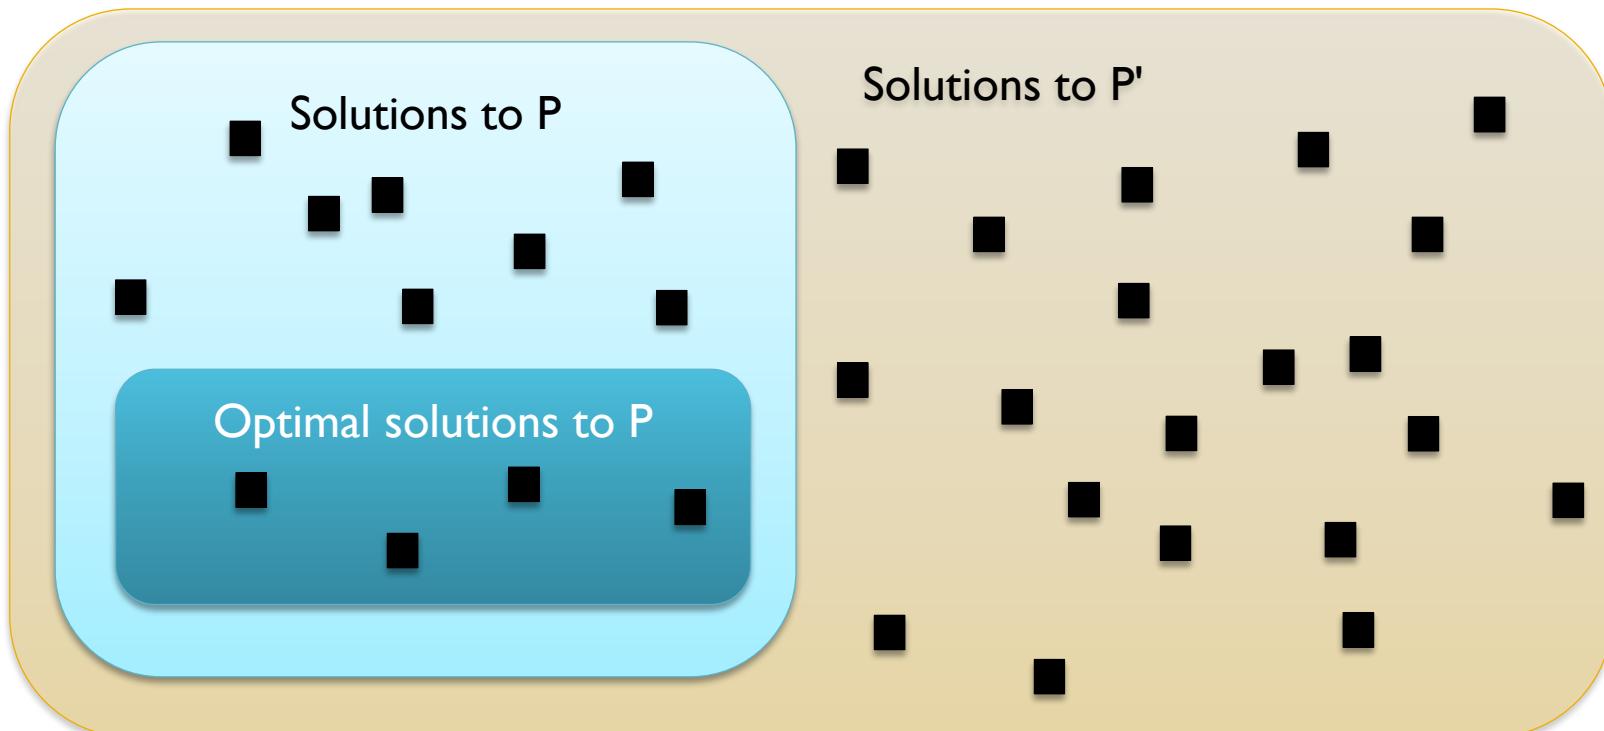
Includes the optimal solutions to P ,
so $\min \{ \dots \}$ cannot be greater



Solution Sets (2)



- Another sufficient criterion: All solutions to P remain solutions for P'
 - Stronger, but often easier to prove
 - This is called relaxation: P' is a relaxed version of P
 - Relaxes the constraint on what is accepted as a solution:
The **is-solution(plan)?** test is "expanded, relaxed" to cover additional plans

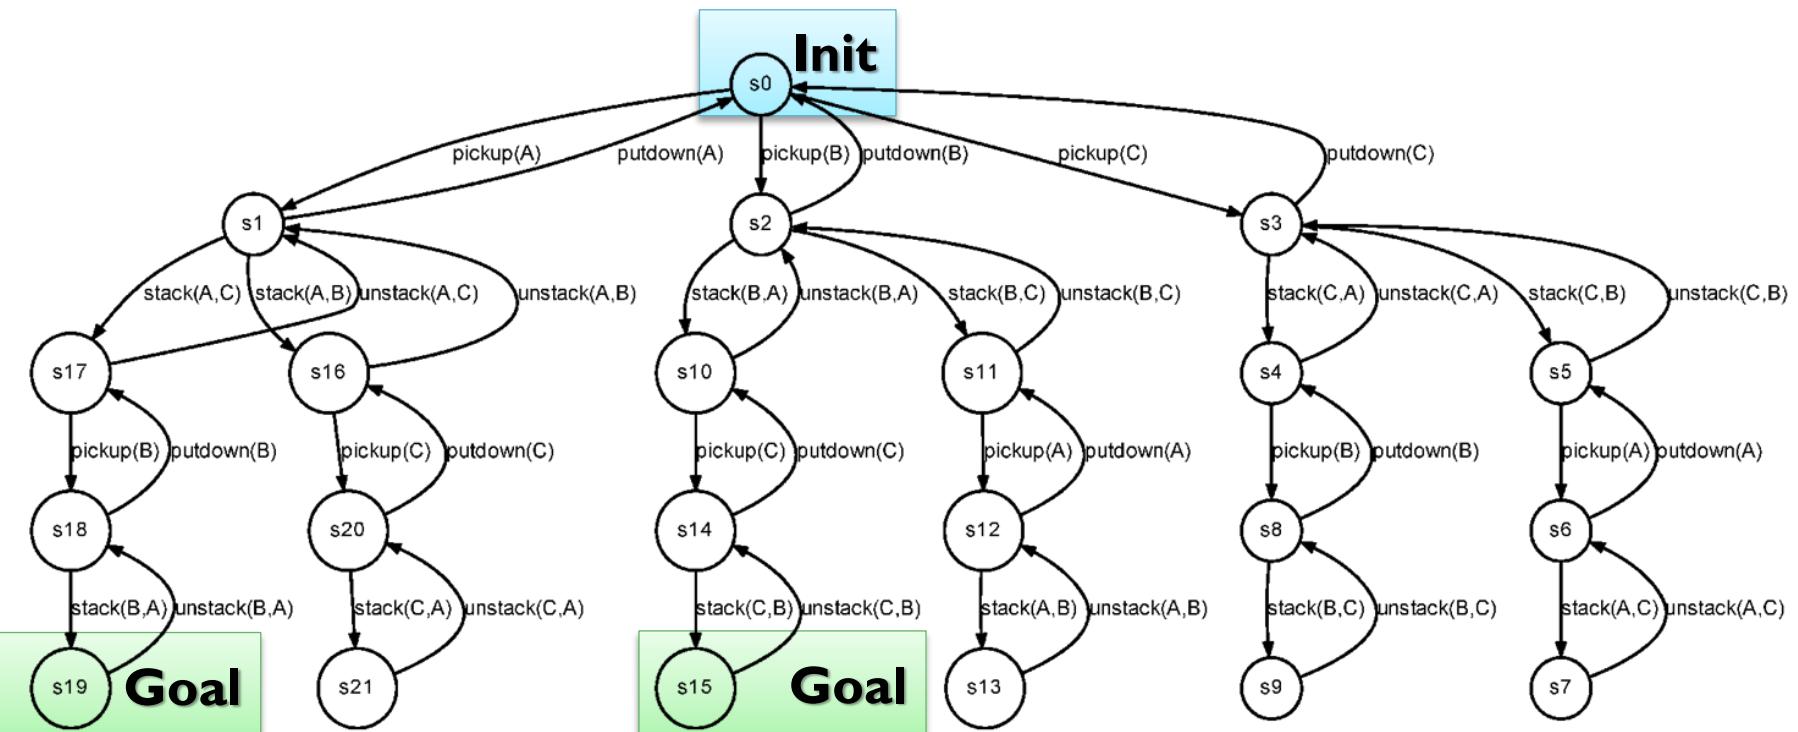


Relaxed Problem 1



■ Relaxation "method 1": Adding goal states

- New goal formula: (and (on B A) (**or** (on A C) (on C B)))
 - All old solutions still valid
 - Some non-solution paths become solutions

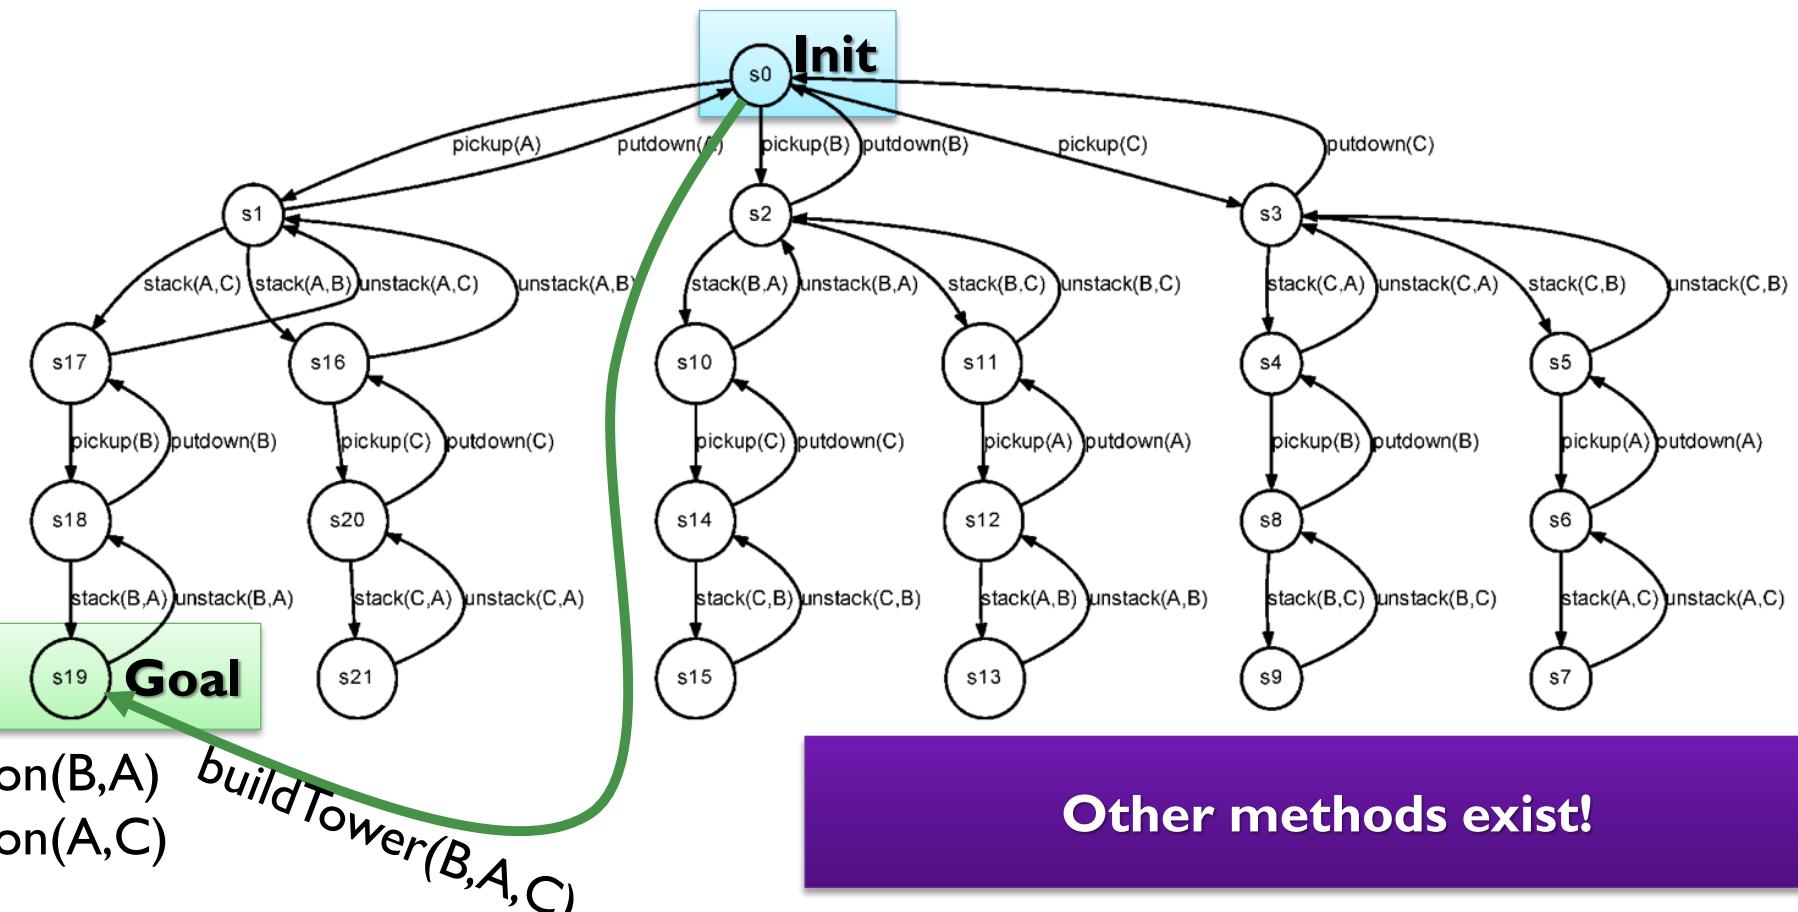


Relaxed Problem 2

73

jonkv@ida

- Relaxation "method 2": **Adding new actions to the domain**
 - **Modifies** the state transition system (states/actions)
 - This particular example: a *shorter* solution becomes possible



Understanding Relaxation

74

jonkv@ida

■ Important:

Relaxation is not simply "removing preconditions"

There are many other relaxations

Definition: Transformation that preserves all old solutions!

You cannot "use a relaxed problem as a heuristic".

What would that mean? The problem is not a number or function...

You use the cost of an optimal solution to the relaxed problem as a heuristic.

If you just take the cost of any solution to the relaxed problem,
it can be inadmissible!

You have to solve it optimally to get the admissibility guarantee.

You don't just solve the relaxed problem once.

Every time you reach a new state and want to calculate a heuristic value,
you have to solve the relaxed problem
of getting from that state to the goal.

Understanding Relaxation (2)



Relaxation is just one specific way of (1) finding a simplifying transformation, and (2) proving "not-more-expensive"!

You've seen relaxations adapted to the 8-puzzle

What can you do automatically, independently of the planning domain used?

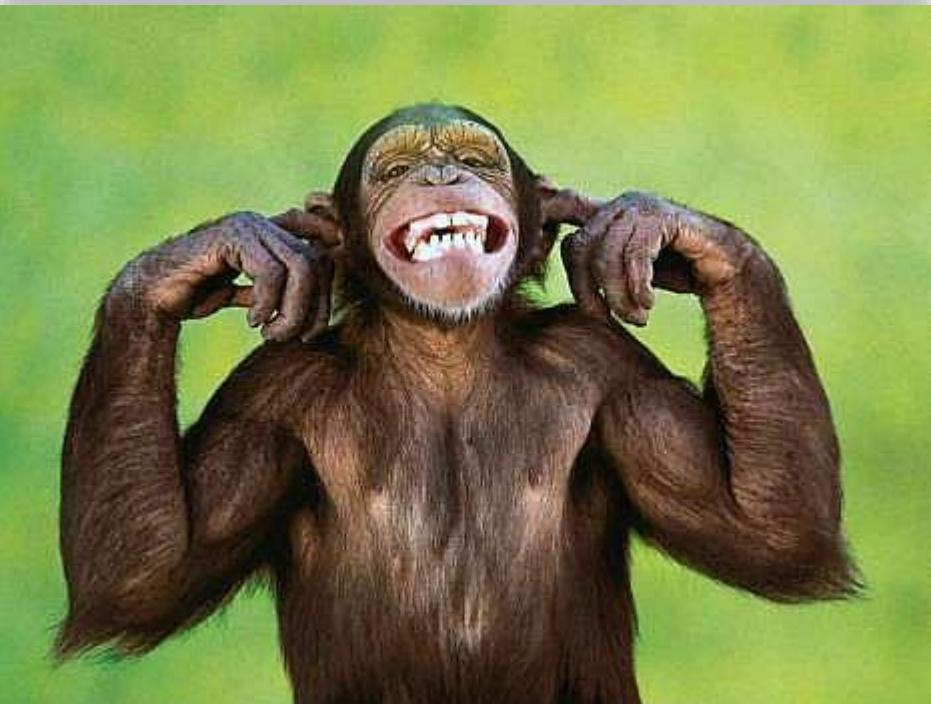
Example: Pattern Database Heuristics

PDB 1: Introduction

77

jonkv@ida

- Main idea behind pattern databases:
 - Let's ignore some facts – everywhere
 - Initial state and goals
 - Preconditions and effects
 - Compute optimal costs as if those facts didn't matter
 - Simpler problem, smaller state space



PDB 2: Dock Worker Robots

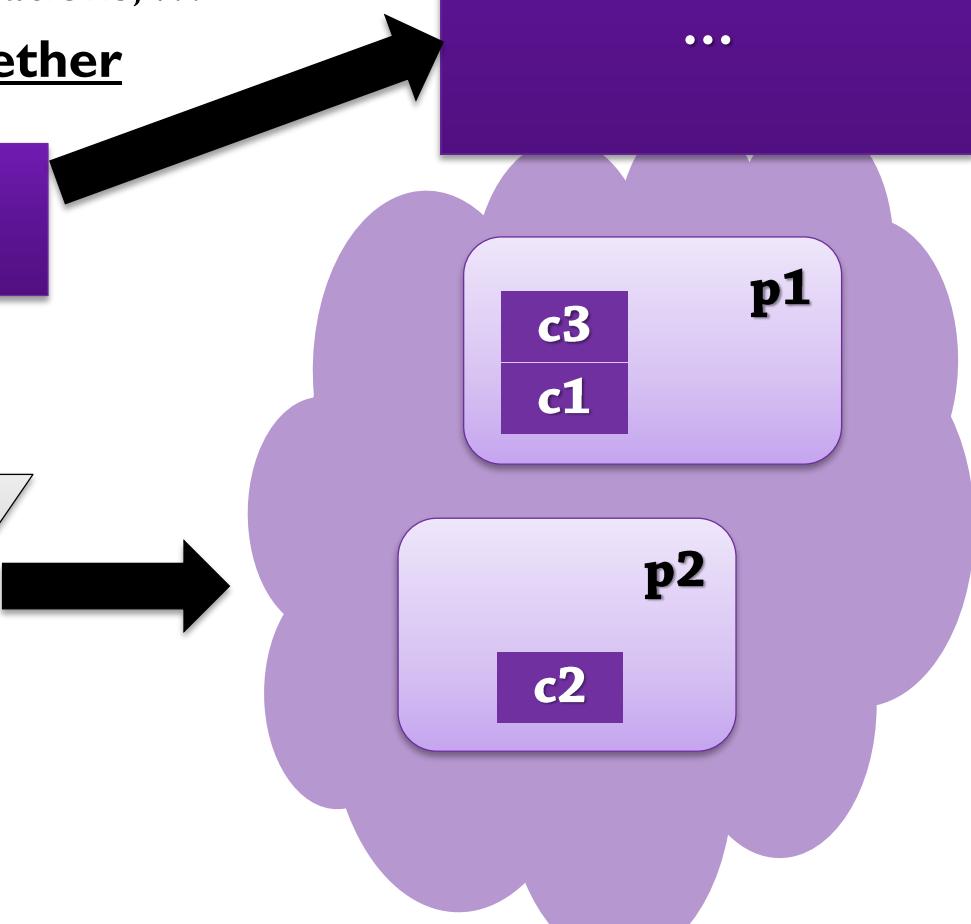
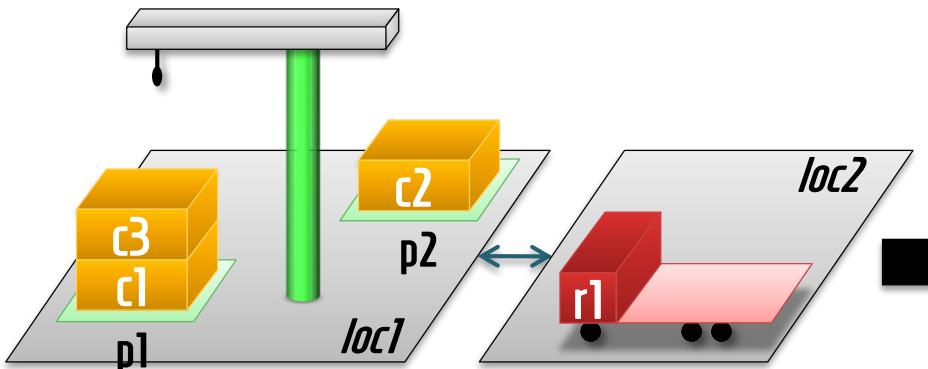
78

jonkv@ida

- Example: Dock Worker Robots
 - Care about facts related to **container locations**
 - $in(container, pile)$, $top(container, pile)$, $on(c1, c2)$, ...
 - Ignore robot locations, crane locations, ...
 - Original states are grouped together

Abstract state in P',
represents many
states in P where
c3 is on **c1** in **p1**,
...
...

Ordinary state in P,
all facts defined



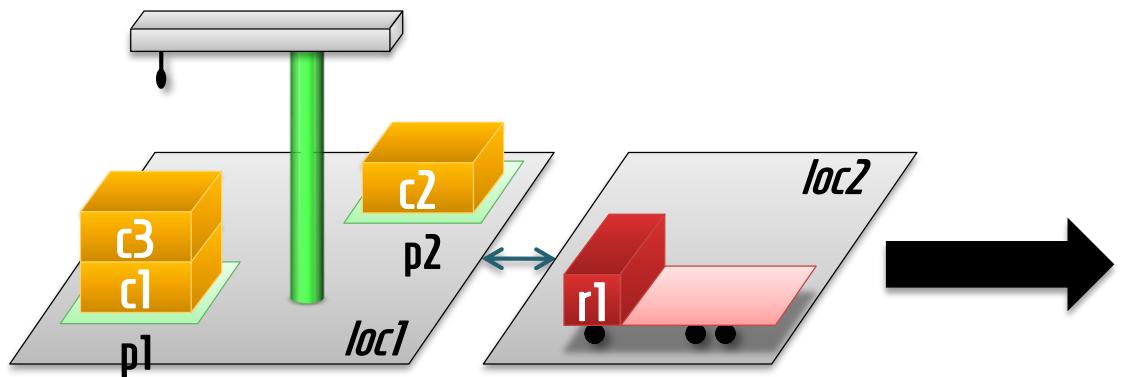
PDB 3: Planning in Patterns

79

jonkv@ida

- In P' we (pretend that we) can use the crane at p1 to:
 - pick up c3 (as we should)
 - place something on r1 (too far away, but we don't care)
 - place five containers on one truck
- But we can't:
 - pick up c1 (we do care about pile ordering)
 - immediately place c1 below c2, ...
 - → **Still a planning problem P' left to solve!**

New paths
to the goal!



Solve optimally, compute cost
→ admissible heuristic!

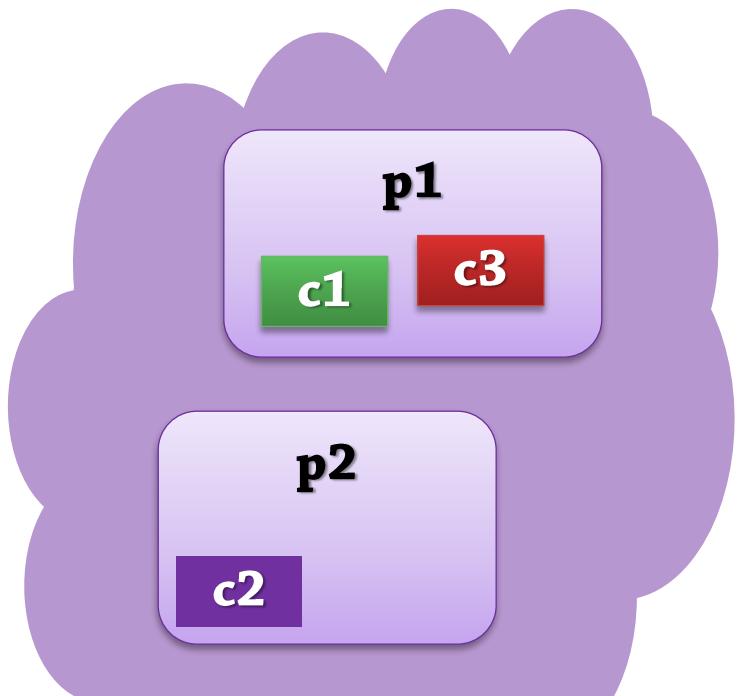
PDB 4: Computing a Heuristic Value



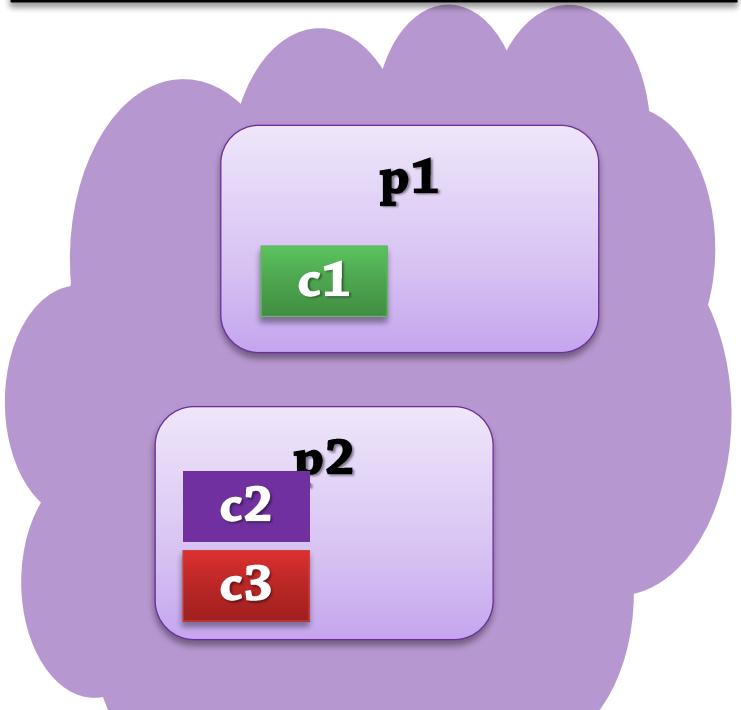
Solve P'(s) optimally:

- Take **c2** with the crane (it's in the way at the bottom)
- Take **c3** with the crane [relaxation – not checking if the crane is busy]
- Place **c3** at the bottom
- Place **c2** on the top

Abstract current state s



Abstract goal



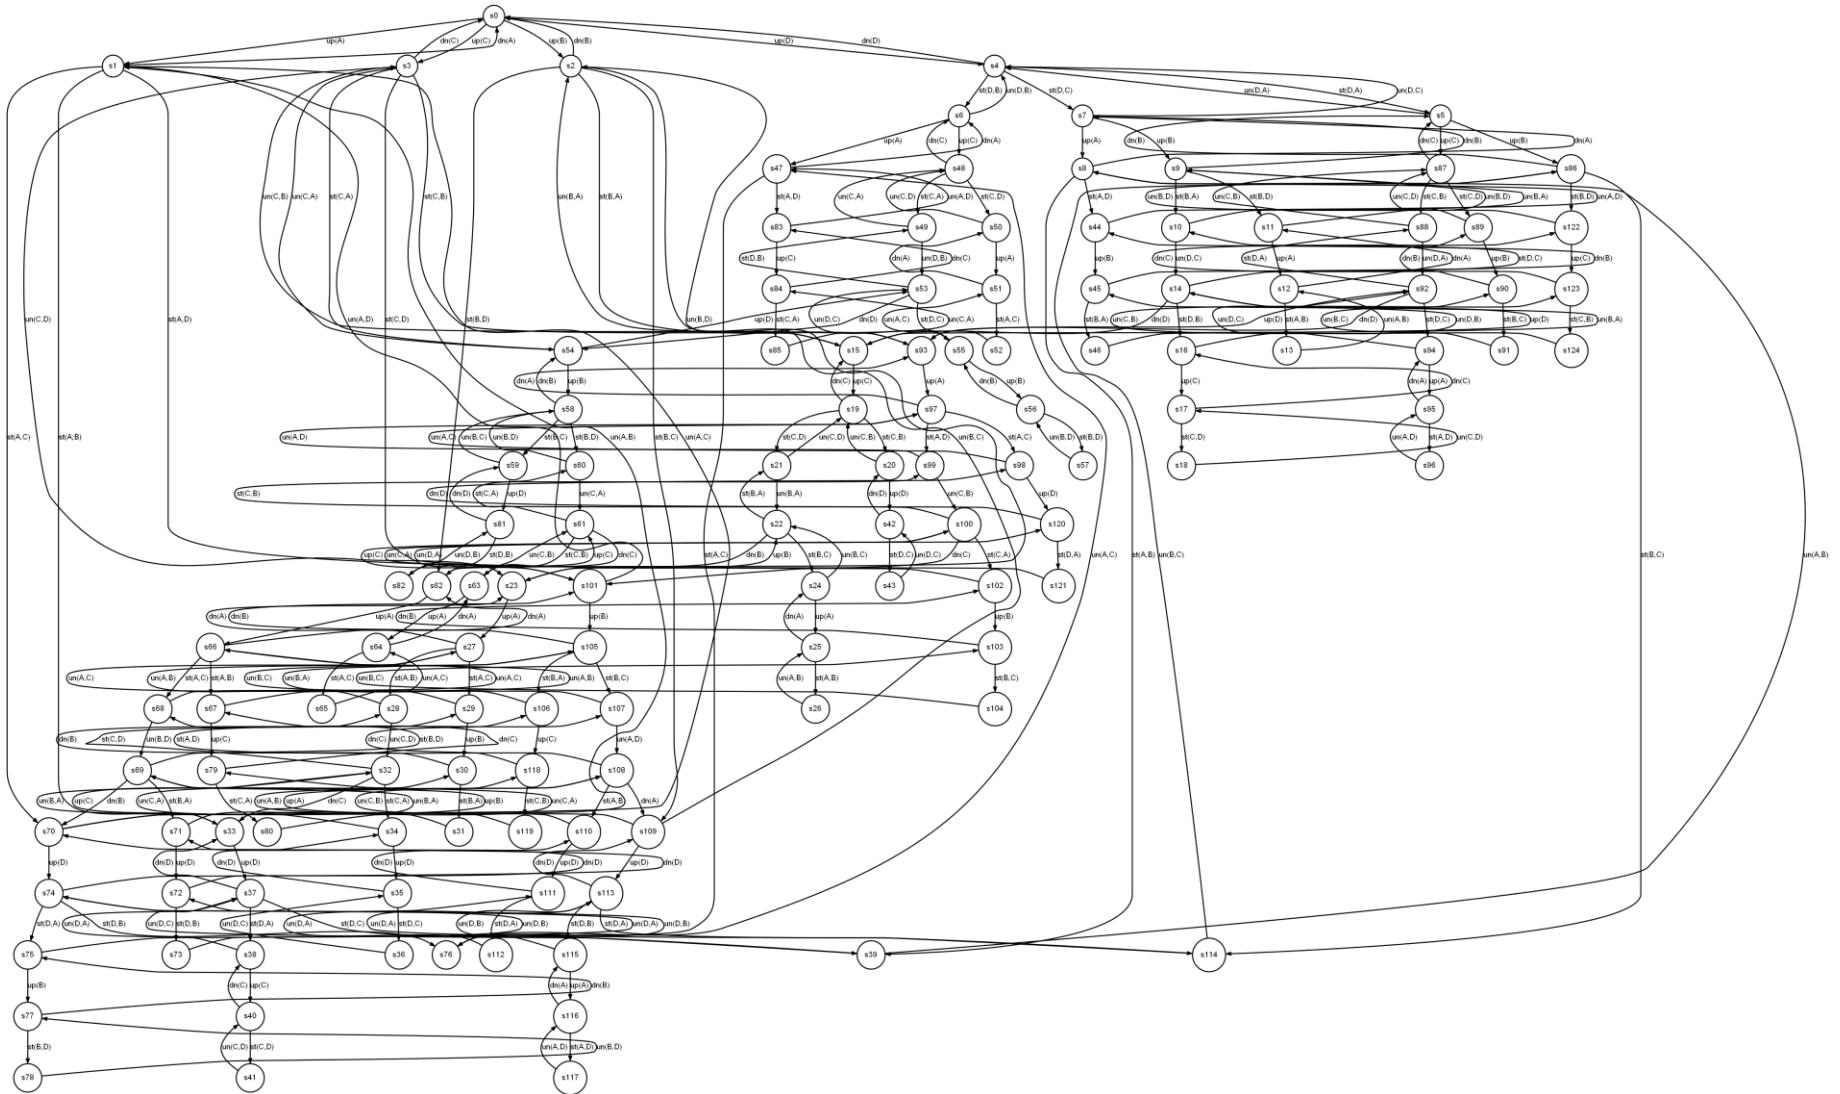
Let's formalize!

PDB: Blocks World size 4

82

jonkv@ida

- Consider physically achievable states in the blocks world, size 4:



PDB: Blocks World size 4, facts

83

jonkv@ida

- All **ground atoms (facts)** in this problem instance:

- **(on A A) (on A B) (on A C) (on A D)**
(on B A) (on B B) (on B C) (on B D)
(on C A) (on C B) (on C C) (on C D)
(on D A) (on D B) (on D C) (on D D)

(ontable A) (ontable B) (ontable C) (ontable D)

(clear A) (clear B) (clear C) (clear D)

(holding A) (holding B) (holding C) (holding D)

(handempty)

PDB: Ignoring Facts

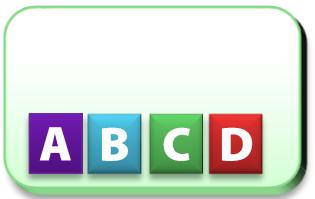
84

jonkv@ida

- Example: only consider 5 **ground facts** related to **block A**

- "Pattern": $p=\{(on\ A\ B), (on\ A\ C), (on\ A\ D), (clear\ A), (ontable\ A)\}$

- Initial state:**



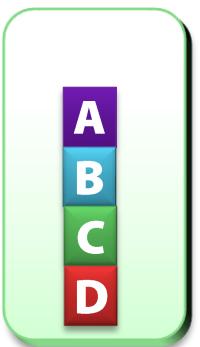
ontable(A)
ontable(B)
ontable(C)
ontable(D)
clear(A)
clear(B)
clear(C)
clear(D)
handempty



ontable(A)
clear(A)

An "abstract state"

- Goal:**



clear(A)
on(A,B)
on(B,C)
on(C,D)
ontable(D)
handempty



clear(A)
on(A,B)

An "abstract goal"

PDB: Transforming Actions



- Pattern $p = \{(on A B), (on A C), (on A D), (clear A), (ontable A)\}$

- Example action:** (unstack A B)

- Before transformation:**

```
:precondition (and (handempty) (clear A) (on A B))  
:effect (and (not (handempty)) (holding A) (not (clear A)) (clear B)  
         (not (on A B)))
```

- After transformation:**

```
:precondition (and (clear A) (on A B))  
:effect (and (not (clear A)) (not (on A B)))
```

Loses **some** preconditions
and effects

Let's call this action
 $transform(a, p)$

- Example action:** (unstack C D)

- Before transformation:**

```
:precondition (and (handempty) (clear C) (on C D))  
:effect (and (not (handempty)) (holding C) (not (clear C)) (clear D)  
          (not (on C D)))
```

- After transformation:**

```
:precondition (and)  
:effect (and)
```

Loses **all** preconditions and
effects → never used!

PDB: Patterns, Abstract States

86

jonkv@ida

- The **pattern** p is the **set of ground facts we care about**
 - A state s is represented by the **abstract state** $s \cap p$
 - If $s \cap p = s' \cap p$, the two states are considered **equivalent**

(**clear A**)
(on A B)
(on B C)
(on C D)
(ontable D)
(handempty)

(**clear A**)
(on A B)
(ontable B)
(clear C)
(on C D)
(ontable D)
(handempty)

(**clear A**)
(on A B)
(on B D)
(on D C)
(ontable C)
(handempty)

represented
by a single
abstract
state

(**clear A**)
(on A B)

(**clear A**)
(ontable A)
(clear B)
(on B C)
(on C D)
(ontable D)
(handempty)

(**clear A**)
(ontable A)
(holding B)
(clear C)
(on C D)
(ontable D)

(**clear A**)
(ontable A)
(clear B)
(on B D)
(on D C)
(ontable C)
(handempty)

represented
by a single
abstract
state

(**clear A**)
(ontable A)

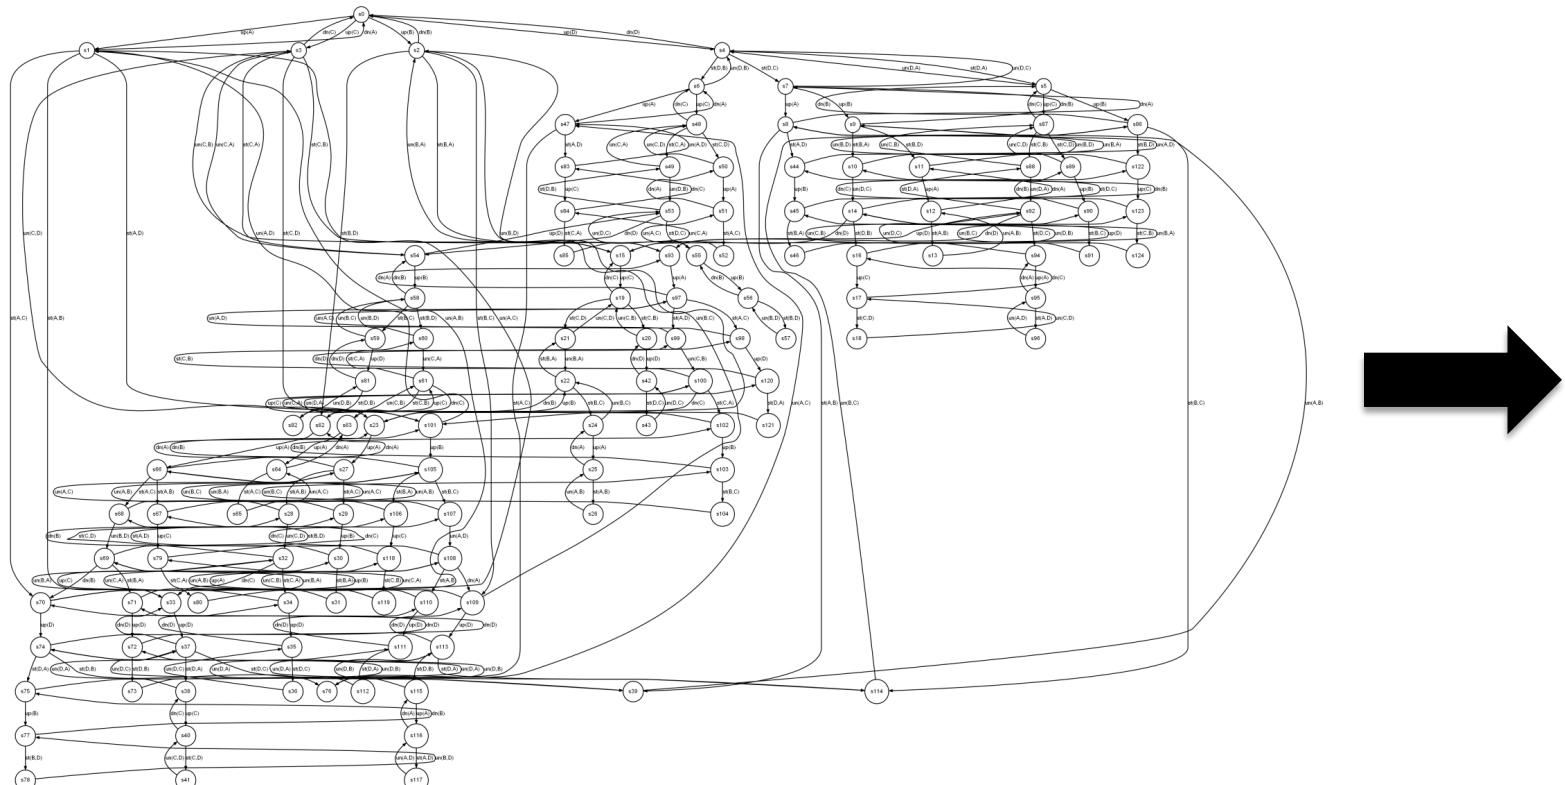
A pattern generally contains **few** facts – for performance!

PDB: New State Space

87



- Pattern $p = \{\text{(on A B)}, \text{(on A C)}, \text{(on A D)}, \text{(clear A)}, \text{(ontable A)}\}$



PDB: Relaxation?



■ Is this a relaxation?

- Yes
- Facts disappear from states...
 - $S' = \{s \cap p \mid s \in S\}$
- But also from precond/goal requirements!
 - If a_i could be executed in s ,
 $\text{transform}(a_i)$ can be executed in $s \cap p$
 - If $g \subseteq s$, then $g \cap p \subseteq s \cap p$
 - ...

```
ontable(A)
ontable(B)
ontable(C)
ontable(D)
clear(A)
clear(B)
clear(C)
clear(D)
handempty
```

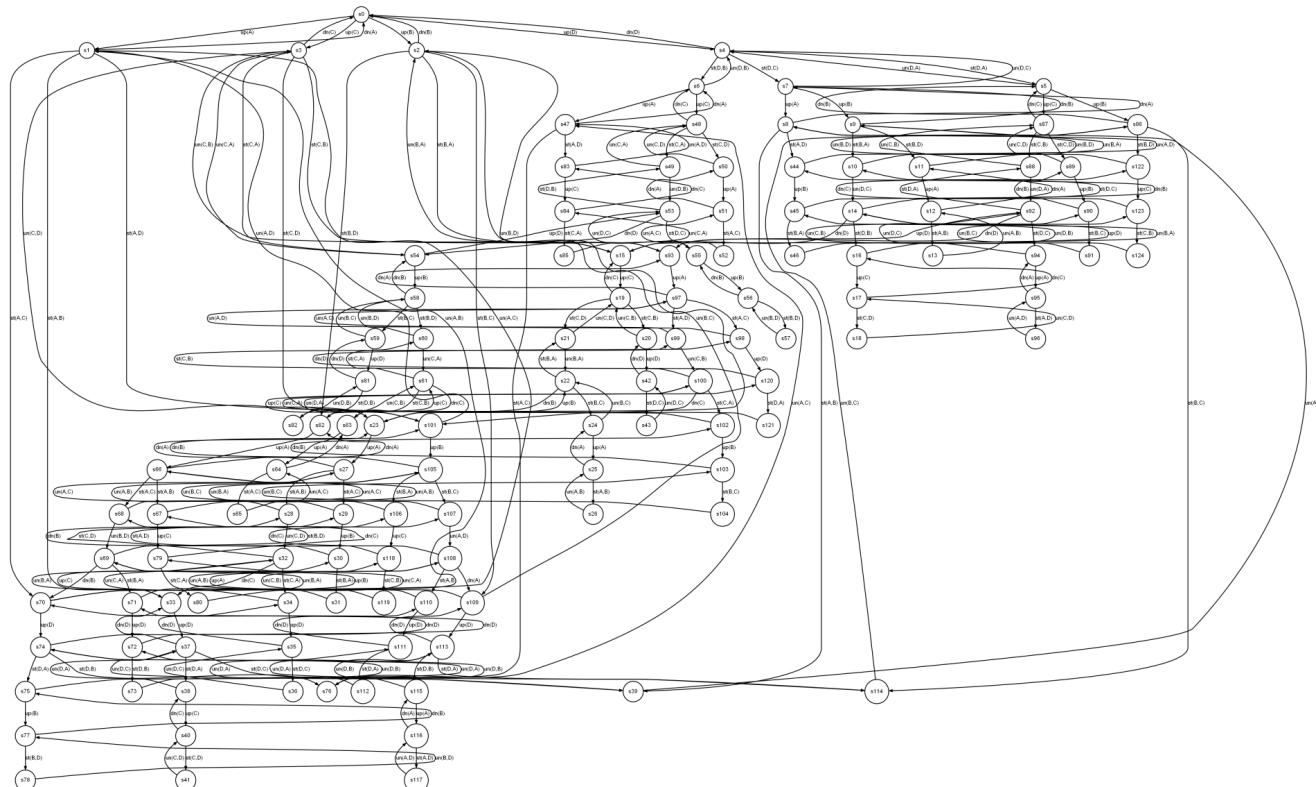


```
ontable(A)
ontable(B)
ontable(C)
ontable(D)
clear(A)
clear(B)
clear(C)
clear(D)
handempty
```

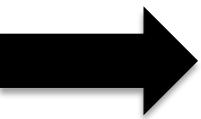
PDB: Preserving Solutions



Any path leading to a goal state here...



...also leads to an abstract goal state here!



`unstack(C,D)`



These paths can contain "loops", actions that do nothing in P':
`unstack(C,D)` may be part of the original solution,
but does not affect any of the *interesting facts*

PDB: State Transition Graph

90

jonkv@ida

■ New reachable state transition graph:

- **Current state:** Everything on the table, hand empty, all blocks clear
 - Abstract state: $s_0 = \{ (\text{ontable } A), (\text{clear } A) \}$

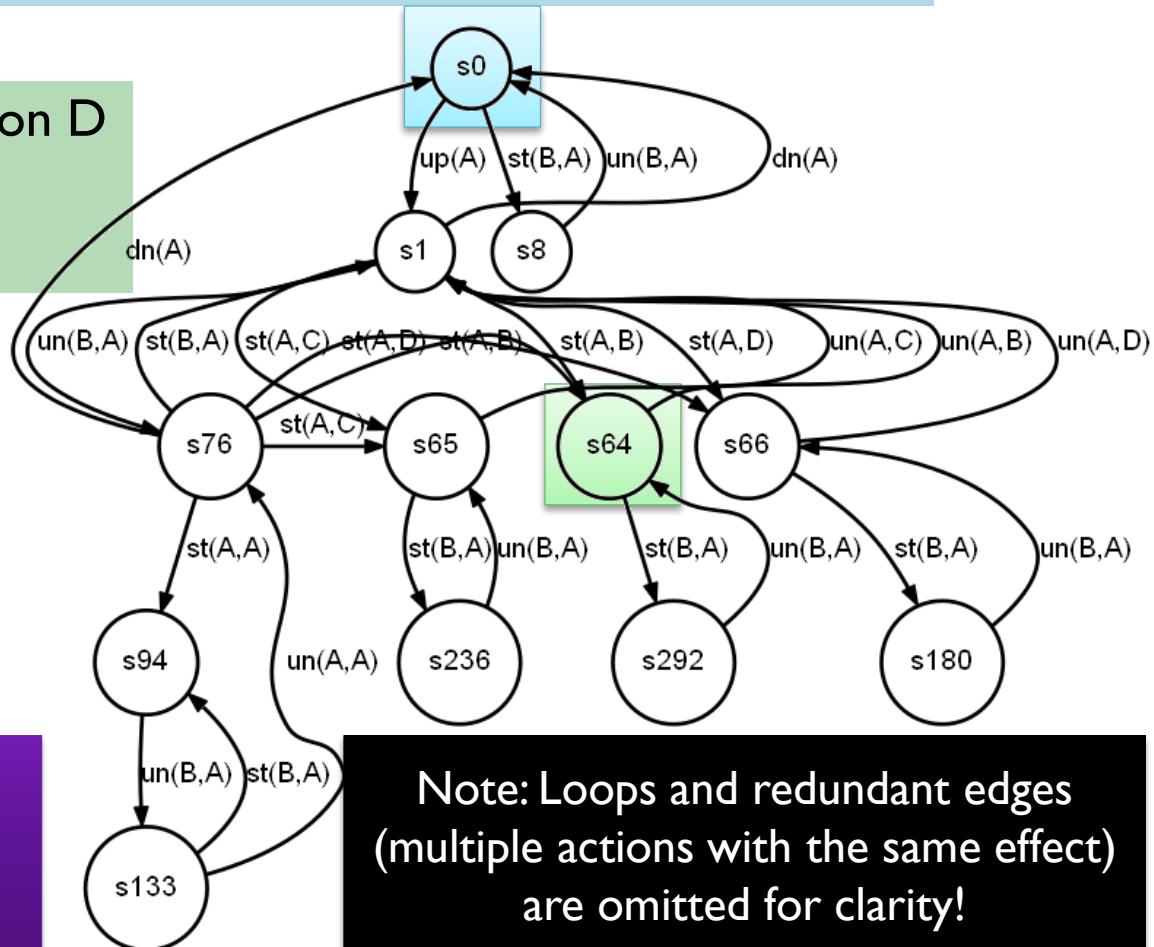
- **Goal state:** A on B on C on D

- Abstract goal: $s_{64} = \{ (\text{on } A B), (\text{clear } A) \}$

- Sufficiently few states to quickly compute optimal costs

- Cost is *at least* 2:
Shortest path $s_0 \rightarrow s_{64}$

Optimal cost of a relaxation
→
admissible heuristic



Note: Loops and redundant edges (multiple actions with the same effect) are omitted for clarity!

PDB: More information



- To make PDB heuristics **more informative**:
 - Calculate costs for **several** patterns
 - Suppose we only care about **{clear(A), ontable(A)}**
 - Suppose we only care about **{on(A,B), on(C,D)}**
 - Suppose we...
 - Take the **maximum** of the computed heuristic values
 - Cost is at least c1 and at least c2 → cost is at least $\max(c1, c2)$
- Real difficulty:
 - Choosing which patterns to use...

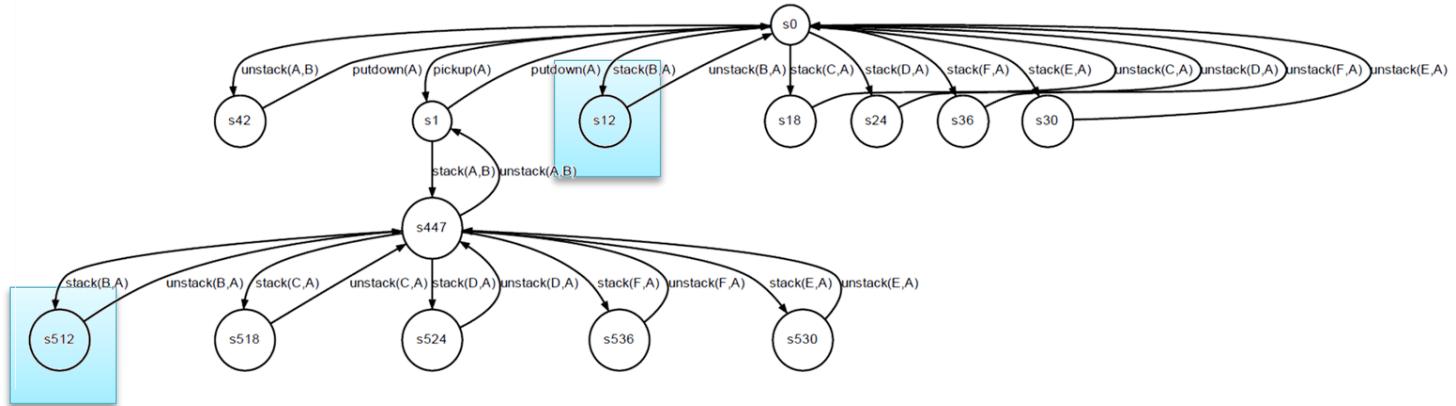
PDB: Databases

92

jonkv@ida

■ Where did the databases go?

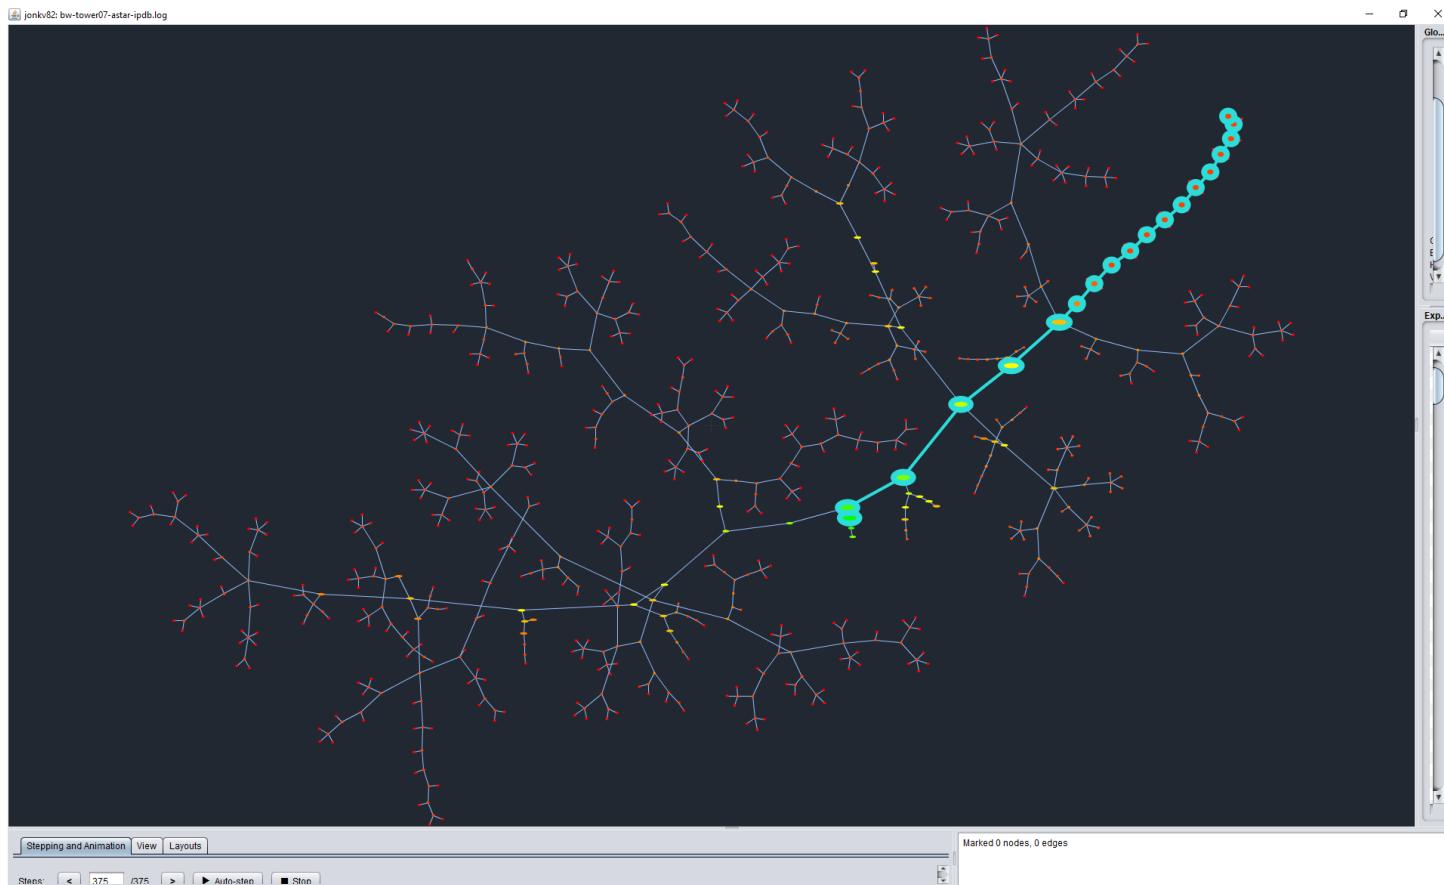
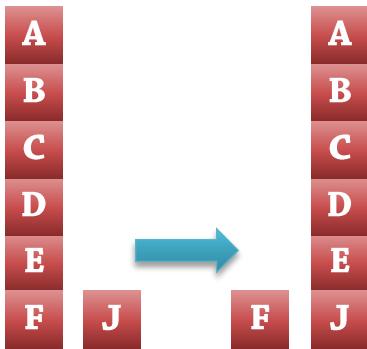
- Planning visits many **abstract** states over and over again
- For each pattern, we **precompute** all reachable abstract states



- Determine which ones satisfy the **abstract goal**
 - Abstract: { aboveA = B }
- Calculate shortest paths from **every** abstract state to **any** abstract goal state
 - Dijkstra's algorithm backwards – possible in a small search space
 - Store in a fast look-up table, a *database*

Only for efficiency!
The principle is more important!

bw-tower07-astar-ipdb: Only 7 blocks, A* search, based on PDB variation



- Blind A*: 43150 states calculated, 33436 visited
- A* + goal count: 6463 states calculated, 3222 visited
- A* + iPDB: 1321 states calculated, 375 visited

No heuristic is perfect – visiting some additional states is fine!

Satisficing Planning

Satisficing Planning



- Optimal plans are nice but often hard to find
 - Larger problem instances → too much time, memory (even with good heuristics)!
- Satisficing plan generation:
 - Find a plan that is sufficiently **good**, sufficiently **quickly**
- What's sufficient?
 - Usually not well-defined!
 - "*These strategies and heuristics seem to give pretty good results for the instances I tested...*"



Speed: Strategies

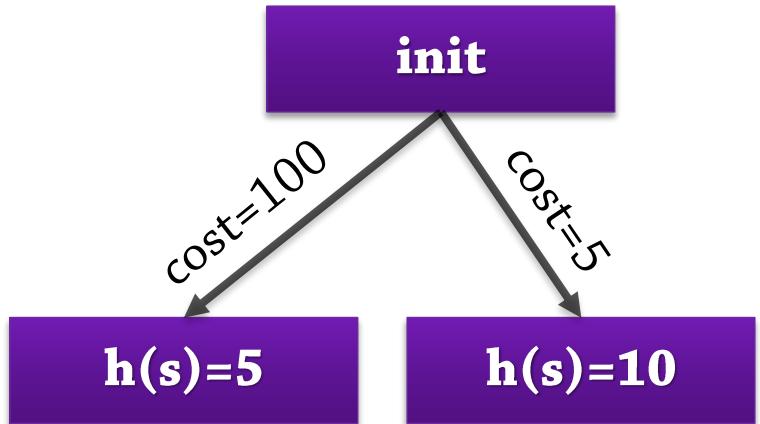
96

jonkv@ida

- One reason for speed: Other informed search strategies

- Simple: Greedy best first search

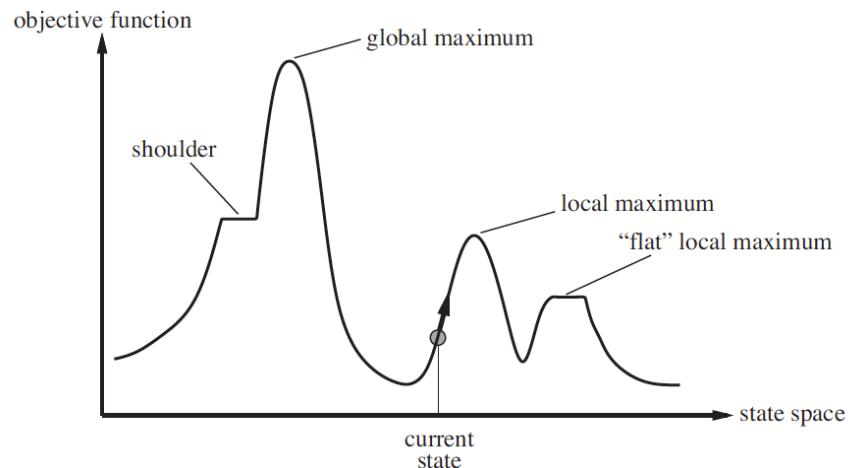
- Always expand the node that seems to be closest to the goal
- Who cares if getting there was expensive?
At least I might find a way to the goal!
- (Only care about $h(s)$, not about $g(s)$)



- Hill climbing

- Be stubborn:
If one direction seems promising,
continue in this direction

- Many others!



Speed: Heuristics

97

jonkv@ida

- One reason for speed: Often more informative heuristics

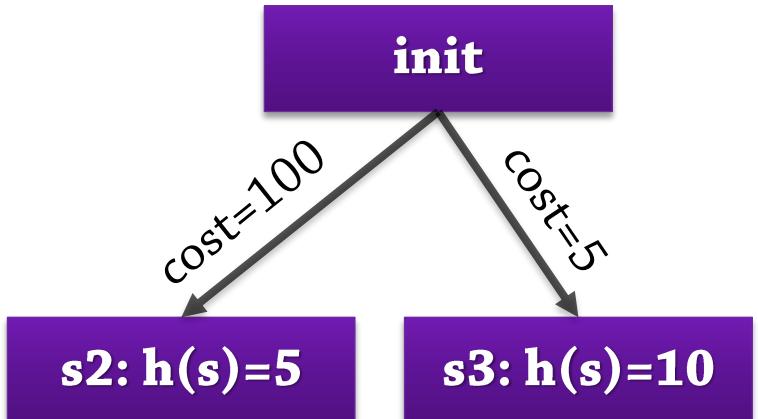
- Optimal planning:

- Often requires admissibility: "We must never overestimate, ever!"
 - Result: Usually *underestimates* (a lot!)

- Satisficing planning:

Extreme example – greedy best first

- Only important that the "best" successor has a low heuristic value
 - For GBF, what the value *is* doesn't matter!

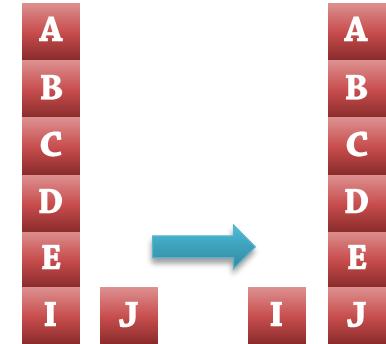
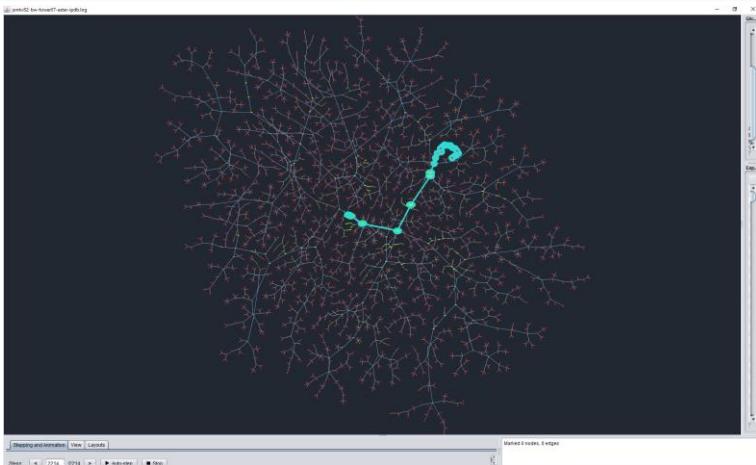


- In many cases:

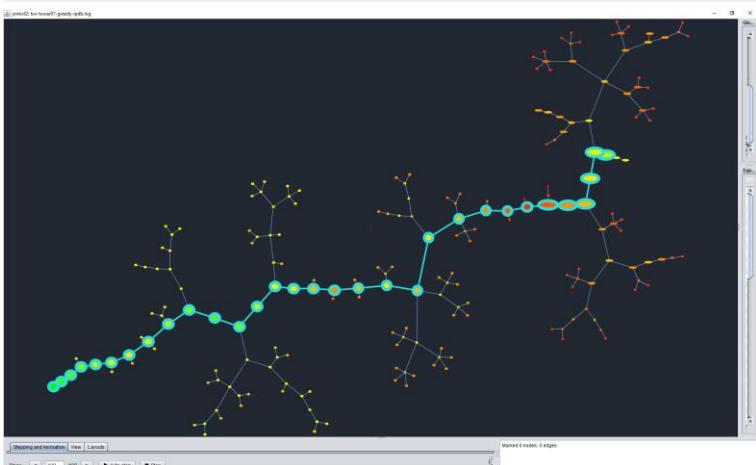
- Admissibility is not required
 - Lack of admissibility is not "only slightly harmful"
 - Lack of admissibility is *irrelevant*

7 blocks (tiny problem)

A*, IPDB: 1321 states, 18 actions



Greedy, IPDB: 171 states, 22 actions



Larger problem instances → the difference increases

Example: Landmark Heuristics

Landmark Heuristics (1)



Landmark:

**"a geographic feature used by explorers and others
to find their way back or through an area"**



Landmark Heuristics (2)



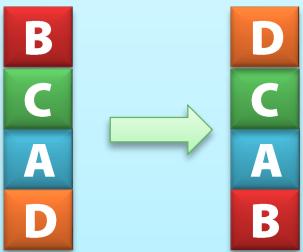
Landmarks in planning:

Something you must pass by/through in every solution to a specific planning problem

Assume we are currently in state s...

Fact Landmark for s:

A **fact** that is **not true** in s,
but must be true at some point
in every solution starting in s



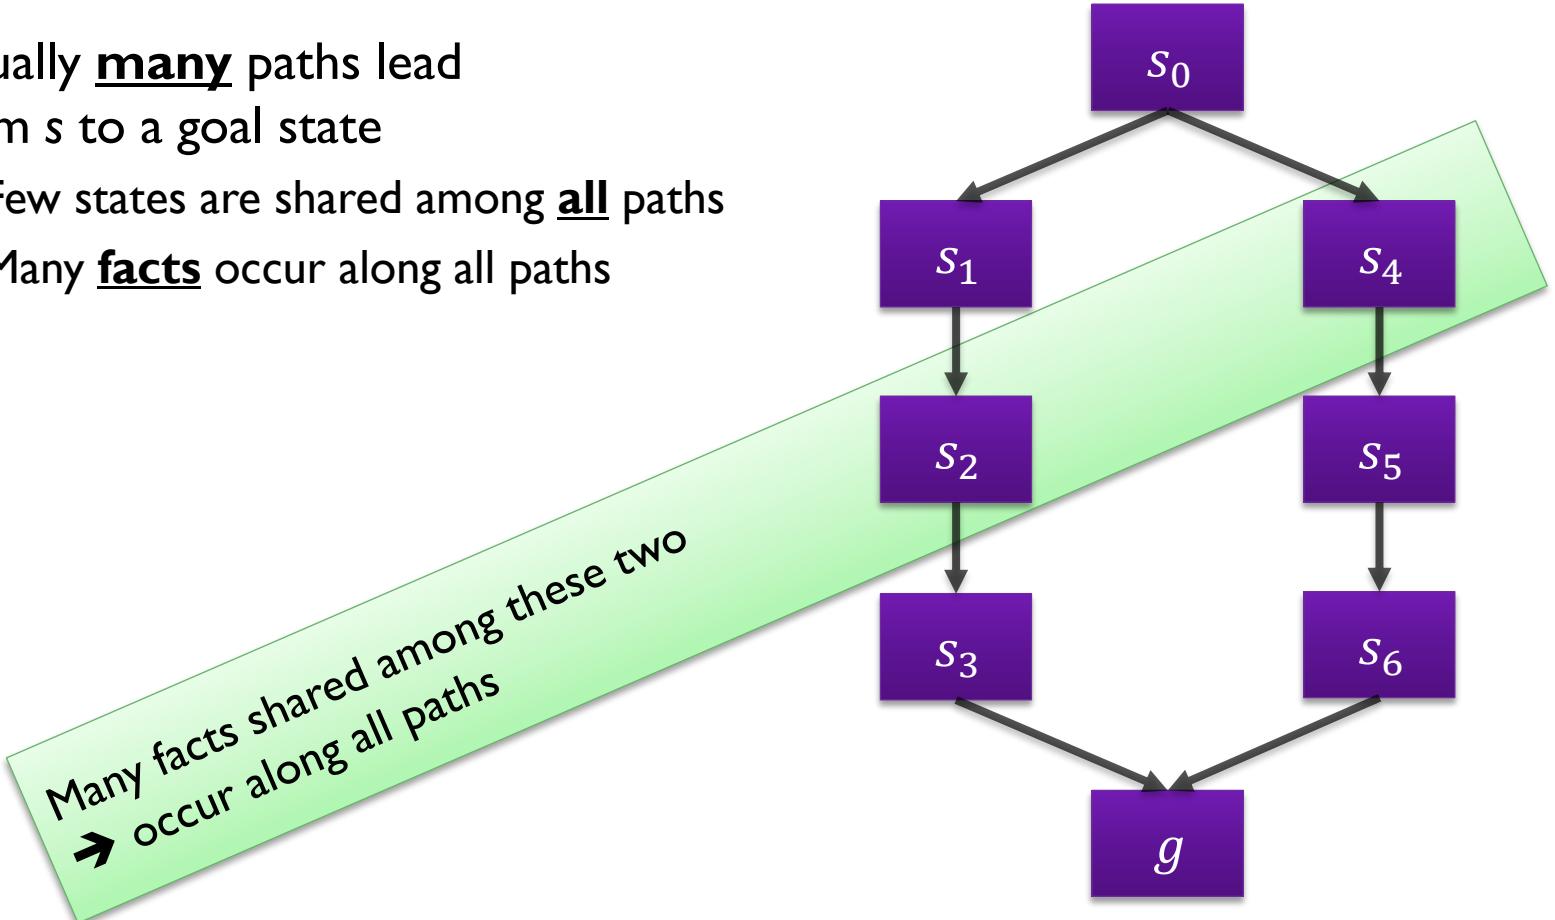
clear(A)
holding(C)
...

Landmark Heuristics (3)



Facts and formulas, not complete states!

- Usually many paths lead from s to a goal state
 - Few states are shared among all paths
 - Many facts occur along all paths



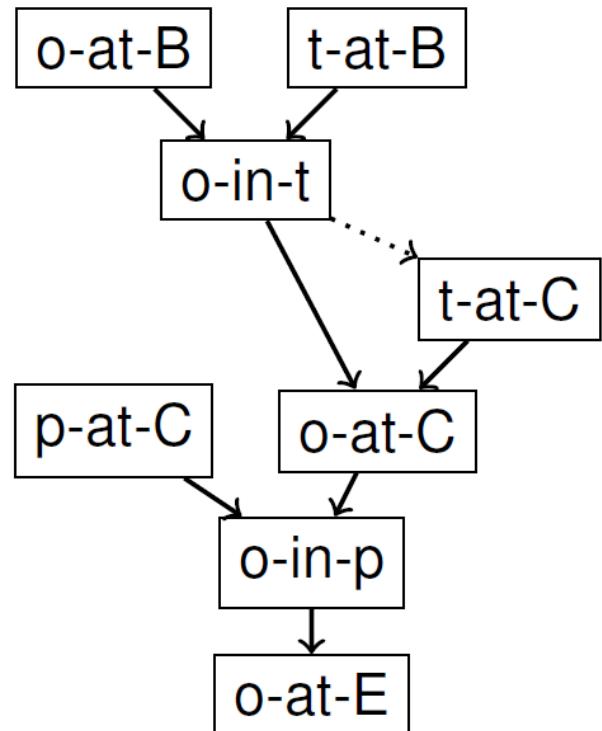
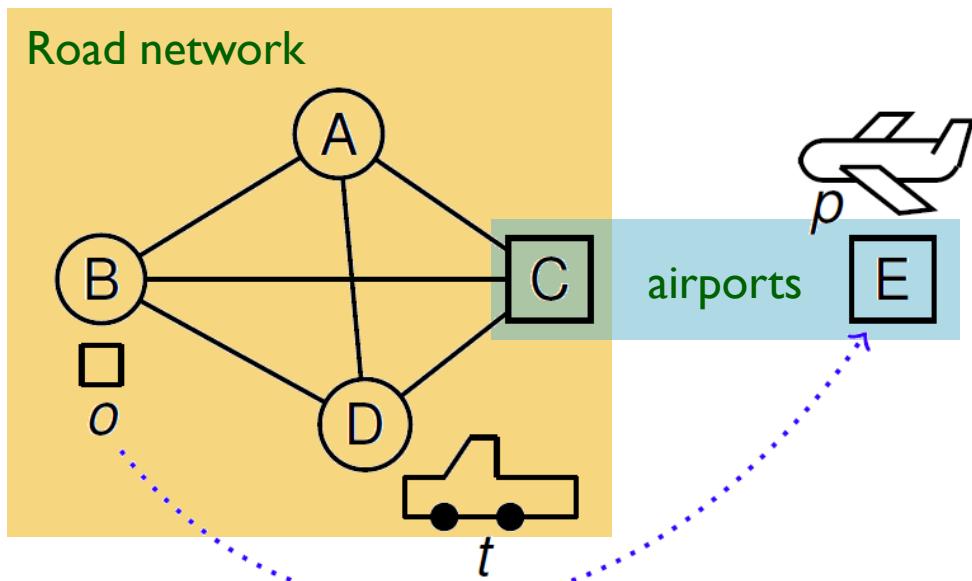
Finding landmarks: Complicated; not covered here

Landmarks as Subgoals



■ Use of landmarks:

- As subgoals: Infer necessary orderings between landmarks, then try to achieve each landmark in succession
 - Example from *Karpas & Richter: Landmarks – Definitions, Discovery Methods and Uses*



Landmark Counts and Costs

104

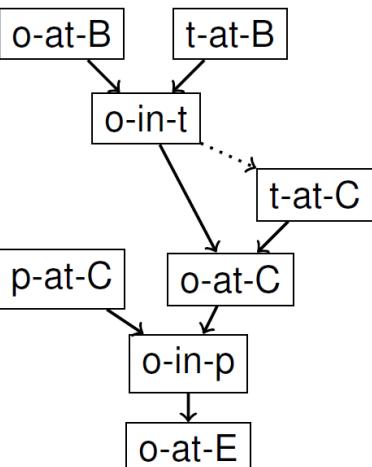
jonkv@ida

■ Another use of landmarks:

- As a basis for **heuristic estimates**
 - Used by LAMA, the winner of the *sequential satisficing* track of the International Planning Competition in 2008, 2011

■ LAMA **counts** landmarks:

- Identifies a set of landmarks that still need to be achieved after reaching state s through path (action sequence) π
- $L(s, \pi) =$



$$(L \setminus \text{Accepted}(s, \pi))$$

All discovered landmarks,
minus those that are
accepted as achieved
(has become true *after*
predecessors are achieved!)

$$\cup$$

$$\text{ReqAgain}(s, \pi)$$

Plus those we can show will
have to be re-achieved

One action can achieve multiple landmarks at once
→ not admissible (though this can be "adjusted")

Example: Hill Climbing

Search Algorithms for non-admissible heuristics



- Some planners use variations on **Hill Climbing**

function HILL-CLIMBING(*problem*) **returns** a state that is a local maximum

current \leftarrow MAKE-NODE(*problem*.INITIAL-STATE)

loop do

neighbor \leftarrow a highest-valued successor of *current*

if *neighbor*.VALUE $<$ *current*.VALUE **then return** *current*.STATE

current \leftarrow *neighbor*

Successors: States created by applying an action to the current state

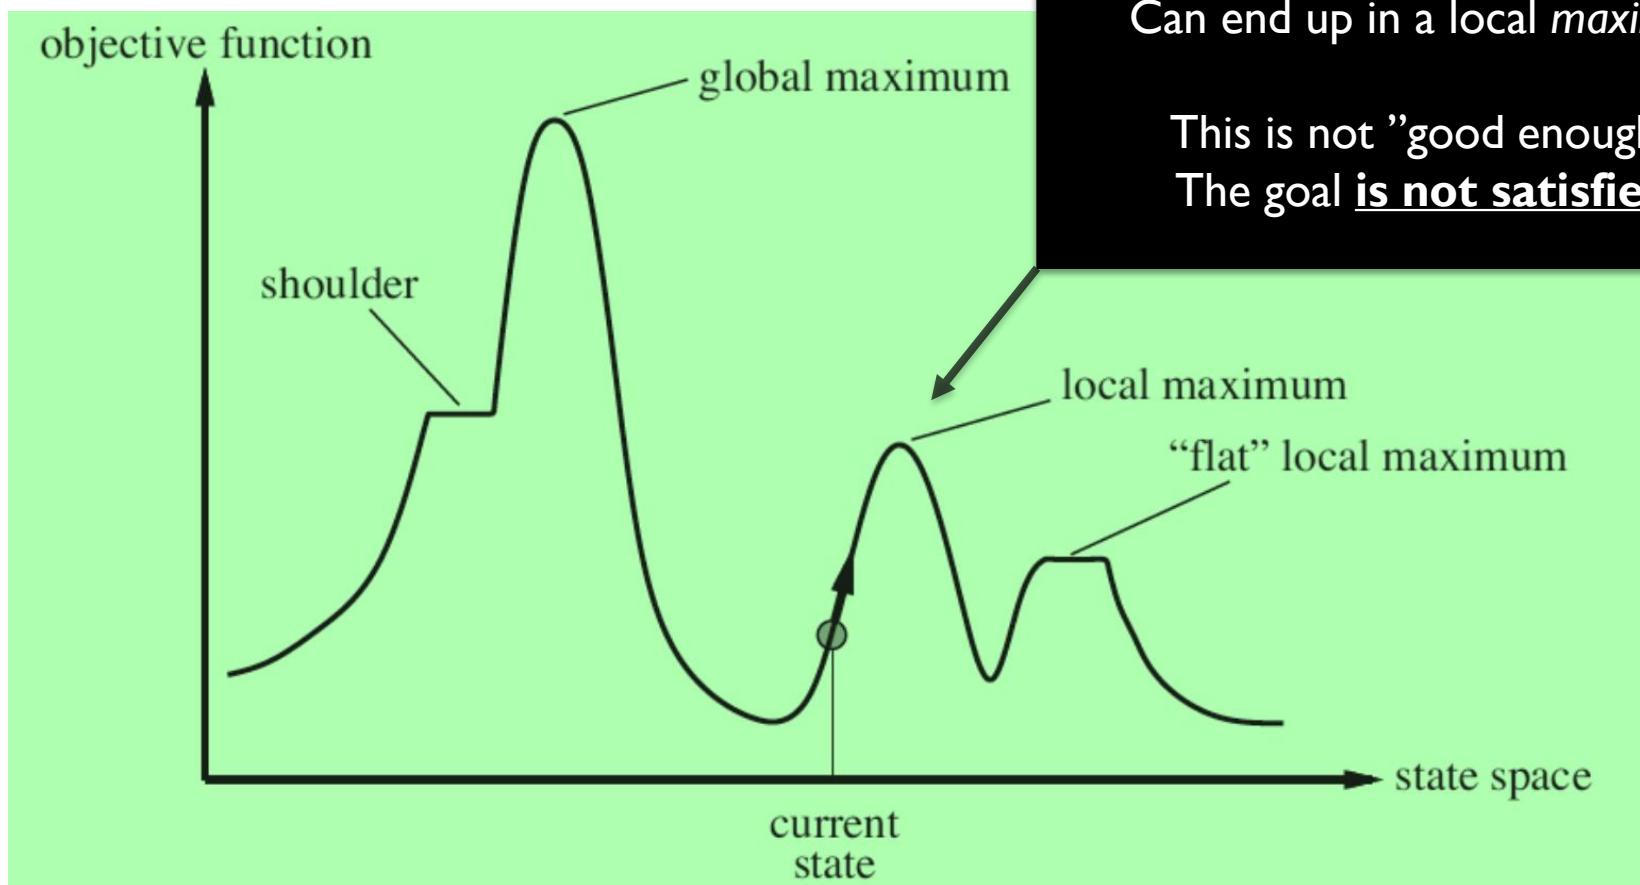
Value: Calculated as $-h(s)$
(maximize $-h(s)$ \rightarrow minimize $h(s)$)

Only considers children of the current node
→ tries to move quickly towards low heuristic values
→ doesn't examine "all nodes" with $\text{cost}(n) + f(n) \leq \text{some value}$

Local Optima in Planning

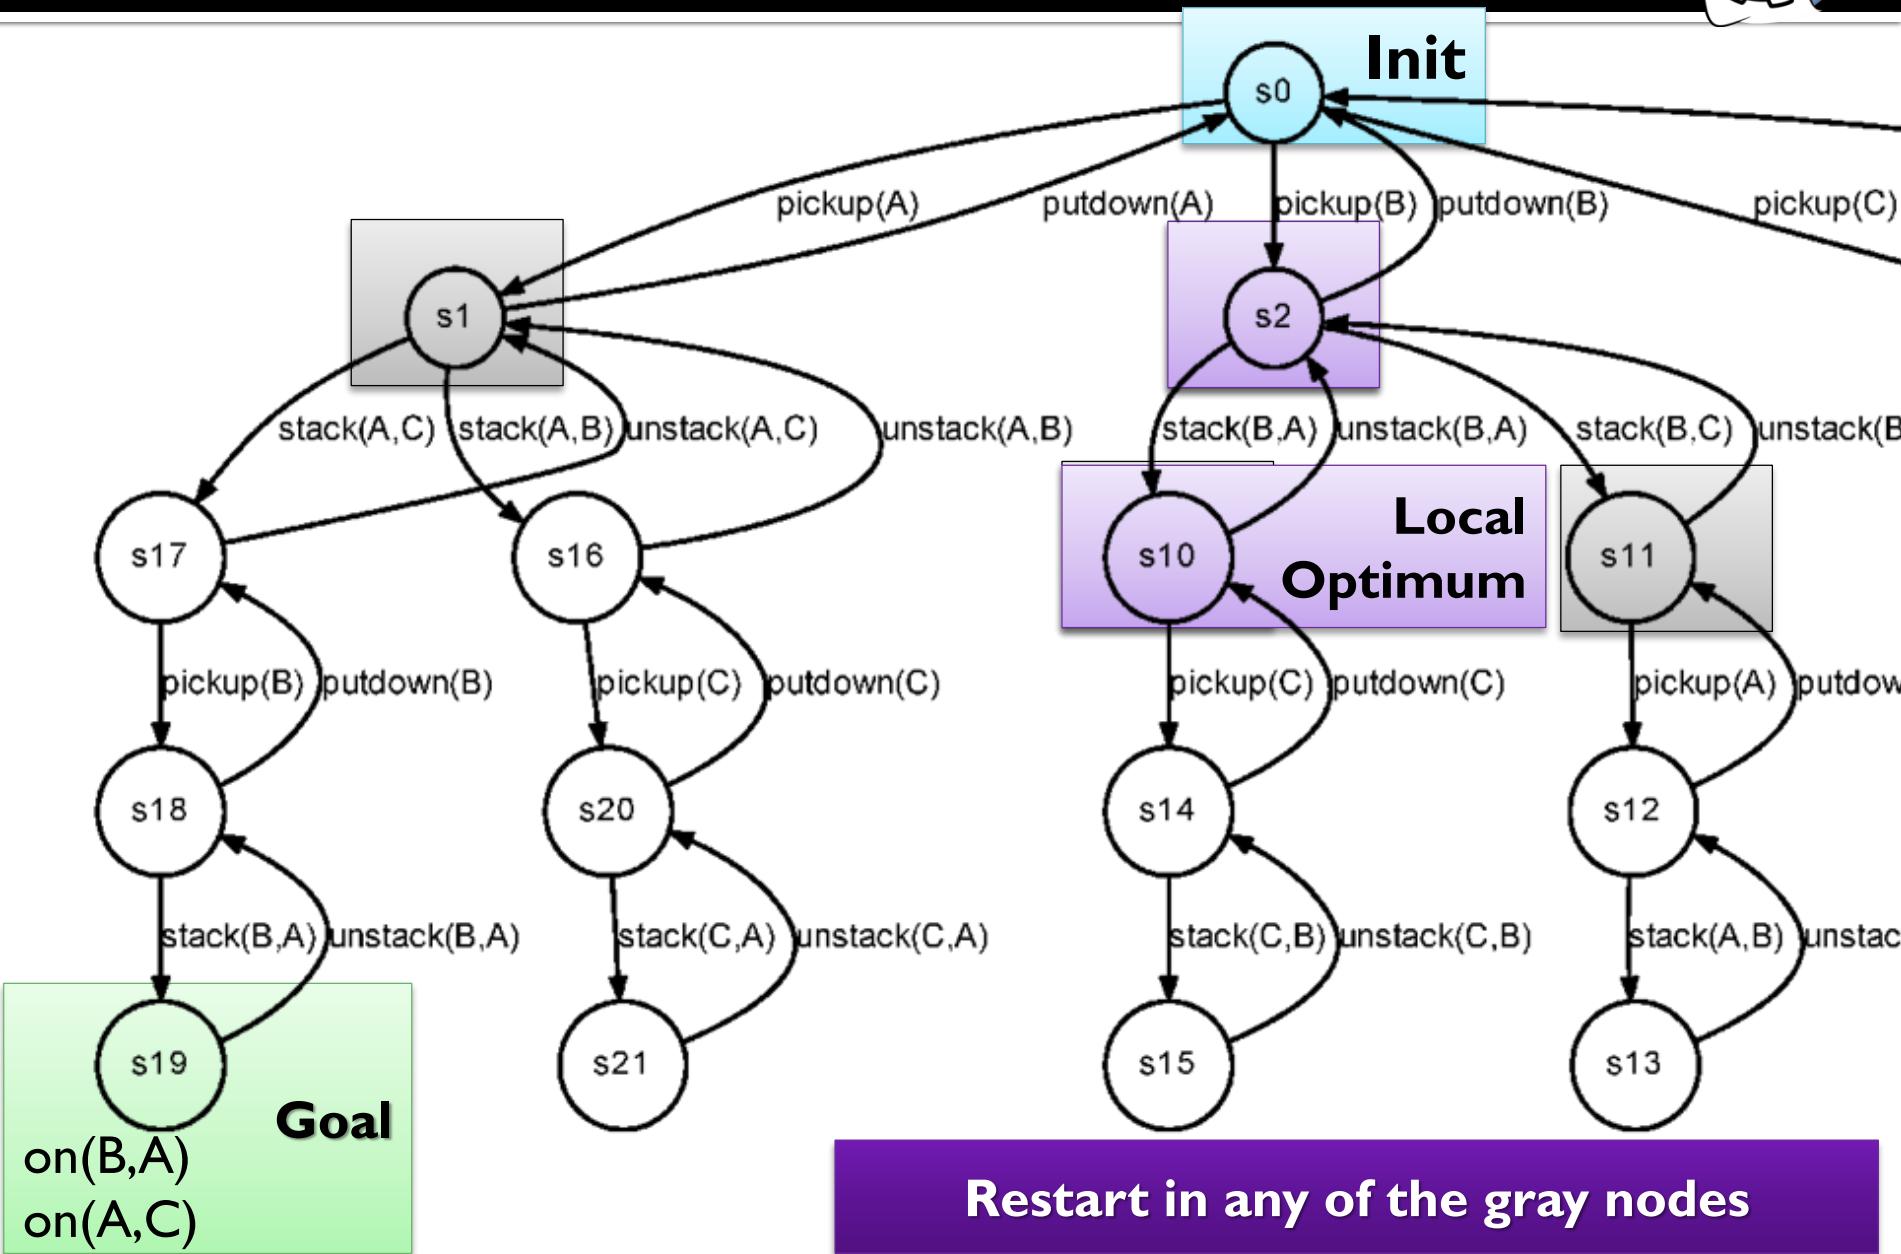


■ Problem:



Possible solution:
Restart search from another expanded state

Local Optima: Restarting



Enforced Hill Climbing



- FastForward uses **enforced** hill climbing – approximately:

- $s \leftarrow$ init-state
- repeat**

expand breadth-first until a better state s' is found
until a goal state is found

Step 1

$h(n) = 40!$

$h(n) = 72$

$h(n) = 50$

$h(n) = 44$

Not expanded

Step 2

$h(n) = 44$

$h(n) = 55$

$h(n) = 41$

$h(n) = 44$

$h(n) = 42$

$h(n) = 37!$

Wait longer to decide which branch to take
Don't restart – keep going

Summary of Forward Search



■ Forward State Space Search:

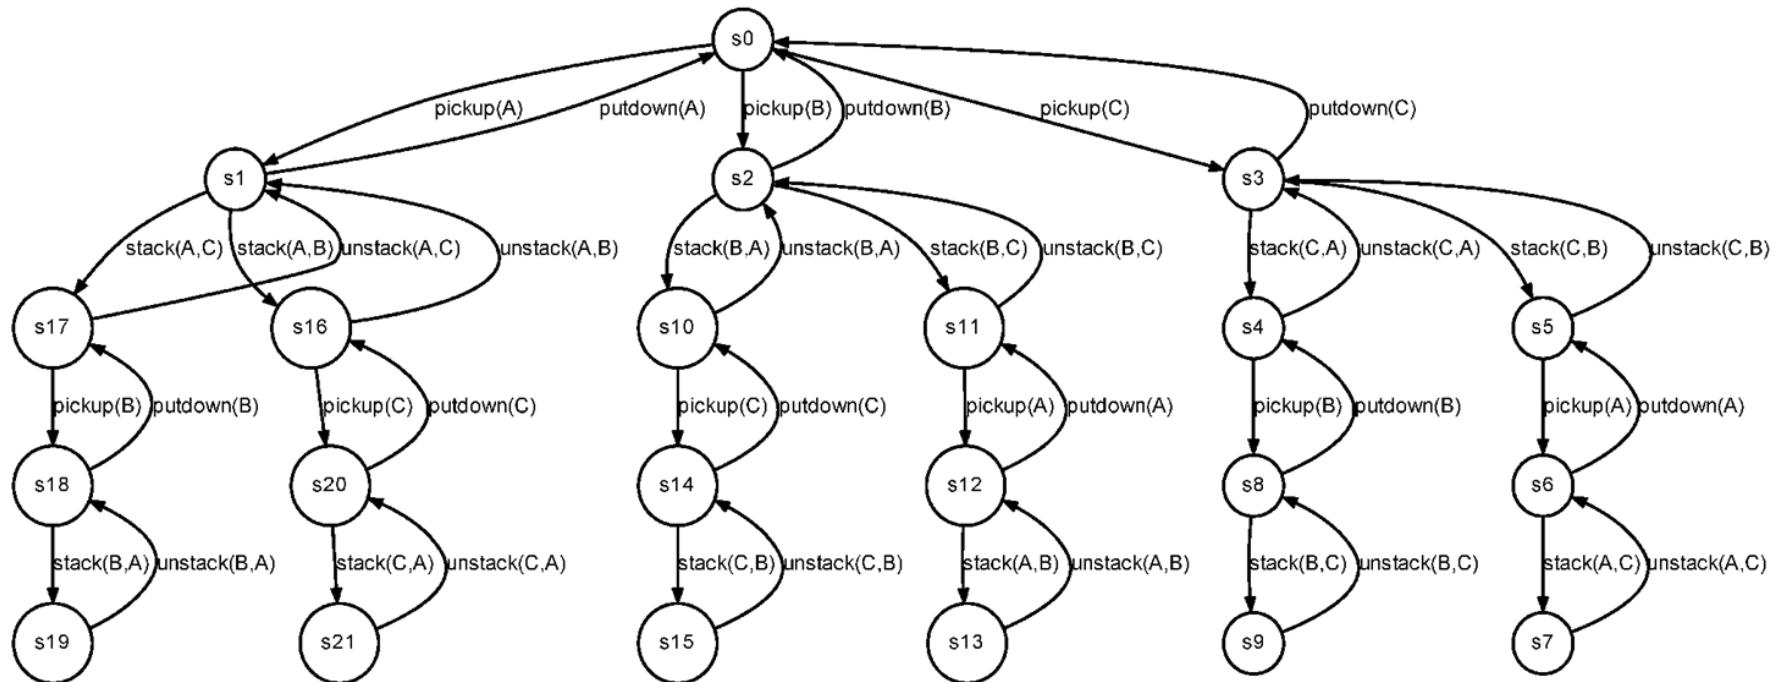
- Very common planning framework
- **Many** heuristic functions defined
 - TDDC17: Focus on understanding the basics, and general principles for defining general heuristics without knowing the problem to be solved (blocks, helicopters, ...)
 - TDDD48: Many additional heuristics (and principles for creating them)
- **Many** search strategies defined
 - TDDC17: Focus on applying standard strategies from general search
 - TDDD48: Additional planning-inspired strategies and their relations to specific heuristic functions

Plan-Space Search: An overview

Introduction



- Just because we are looking for a path in this graph...



- ...doesn't mean we have to use it as a search space!

Motivating Problem



- Simple planning problem:

- Two crates
 - At A
 - Should be at B



- One robot

- Can carry up to two crates
 - Can move between locations, which requires one unit of fuel
 - Has only two units of fuel



Let's see what a forward-chaining planner *might* do (depending on heuristics)...

Motivating Problem 2: Forward Search

We have:

- robotat(A)
- at(c1,A)
- at(c2,A)
- has-fuel(2)

pickup(c1,A)
effects:
 holding(c1),
 \neg at(c1,A)

robotat(A)
holding(c1)
 at(c2,A)
 has-fuel(2)

drive(A,B)
effects:
 \neg robotat(A),
 robotat(B)

robotat(B)
 holding(c1)
 at(c2,A)
has-fuel(1)

put(c1, B)
effects:
 at(c1,B),
 \neg holding(c1)

robotat(B)
at(c1,B)
 at(c2,A)
 has-fuel(1)

pickup(c1, B)
effects:
 \neg at(c1,B),
 holding(c1)

robotat(B)
holding(c1)
 at(c2,A)
 has-fuel(1)

Cycle, backtrack

drive(B,A)
effects:
 \neg robotat(B),
 robotat(A)

robotat(A)
 at(c1,B)
 at(c2,A)
has-fuel(0)

pickup(c2,A)
effects:
 holding(c2),
 \neg at(c2,A)

robotat(A)
 at(c1,B)
holding(c2)
 has-fuel(0)

Dead end,
backtrack

drive(B,A)
effects:
 robotat(A),
 \neg robotat(B)

robotat(A)
 holding(c1)
 at(c2,A)
has-fuel(0)

...dead end, backtrack

And so on...

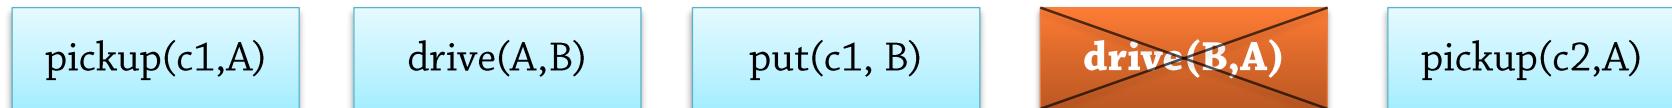
We want:
 at(c1,B)
 at(c2,B)

Motivating Problem 3

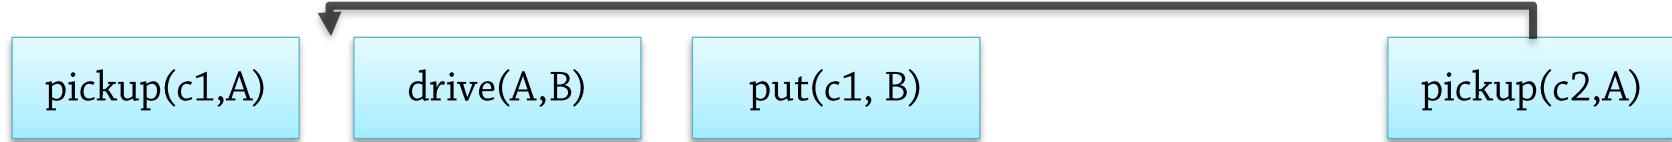


Observations:

- **Most** actions we added before backtracking were useful and necessary!



- At first, we added them in the wrong order



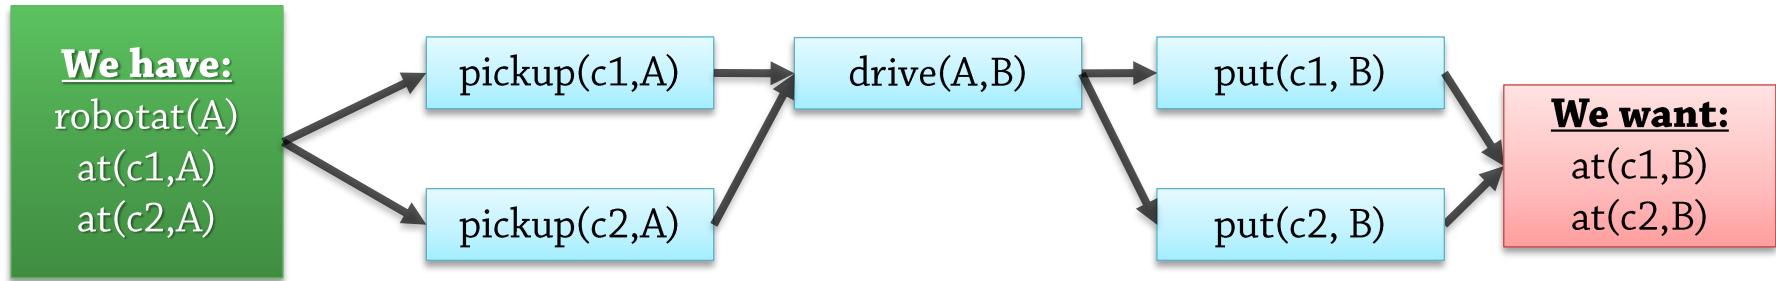
- **Forward** planning commits immediately to action order
 - Puts each action in its final place in the plan

- State space heuristics must be smart enough tell us:
 - Which actions are **useful**
 - **When** to add them to the plan

Partial Order Planning



- What if we could insert actions anywhere?
 - Requires different plan structures, book-keeping, heuristics, ...
- Turns out we might as well take it one step further
 - Don't use sequences; change the plan structure to a partial order
 - Keep track of necessary and possible orderings



Haven't decided if we handle c1 or c2 first

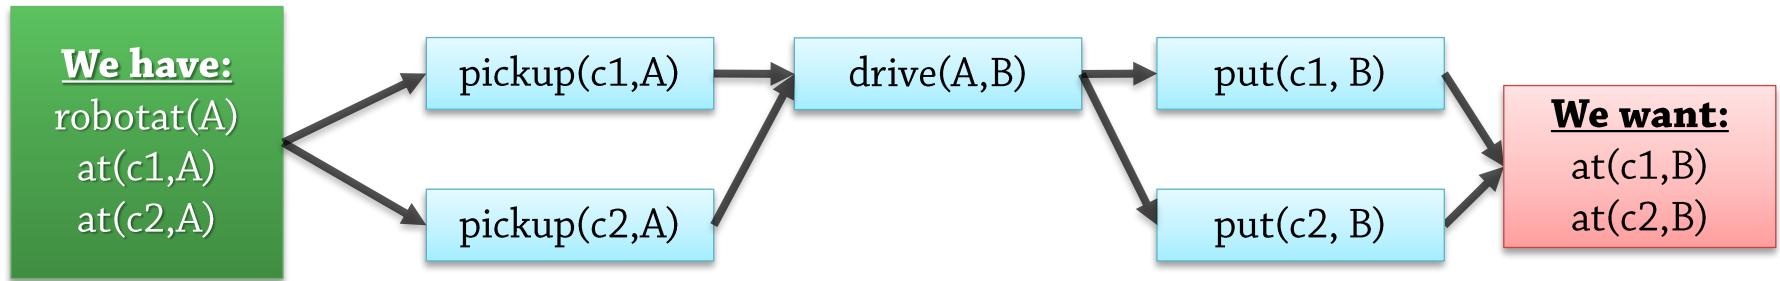
Least/late commitment to ordering – "don't decide until you have to"

Partial Order Planning: Consequences (1)

117

jonkv@ida

- Consequence I: If plans look like this...



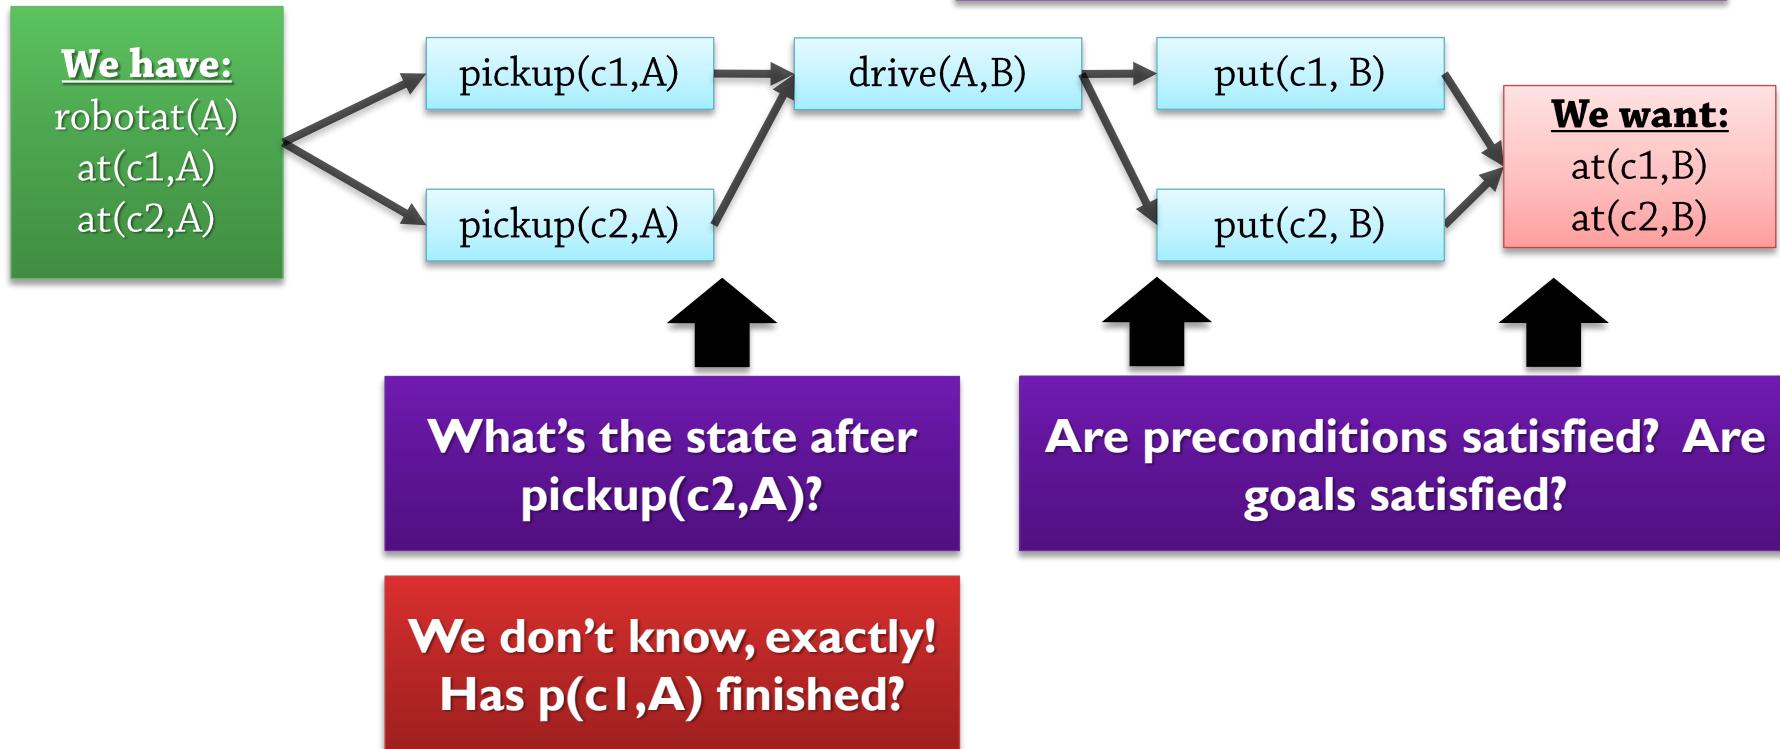
How do we define the set of successors?
Can we add any action anywhere?

Partial Order Planning: Consequences (2)



Consequence 2:

Can the put actions interfere with each other in some way?



More sophisticated "bookkeeping" required!

Partial-Order Planning



- How to keep track of initial state and goal?

A "fake" initial action,
whose effects are
the facts of the initial
state

A "fake" goal action,
whose preconditions are
the conjunctive goal facts



Streamlines the planner: *Initial state and action effects* are handled equivalently
goal and action preconditions are handled equivalently.

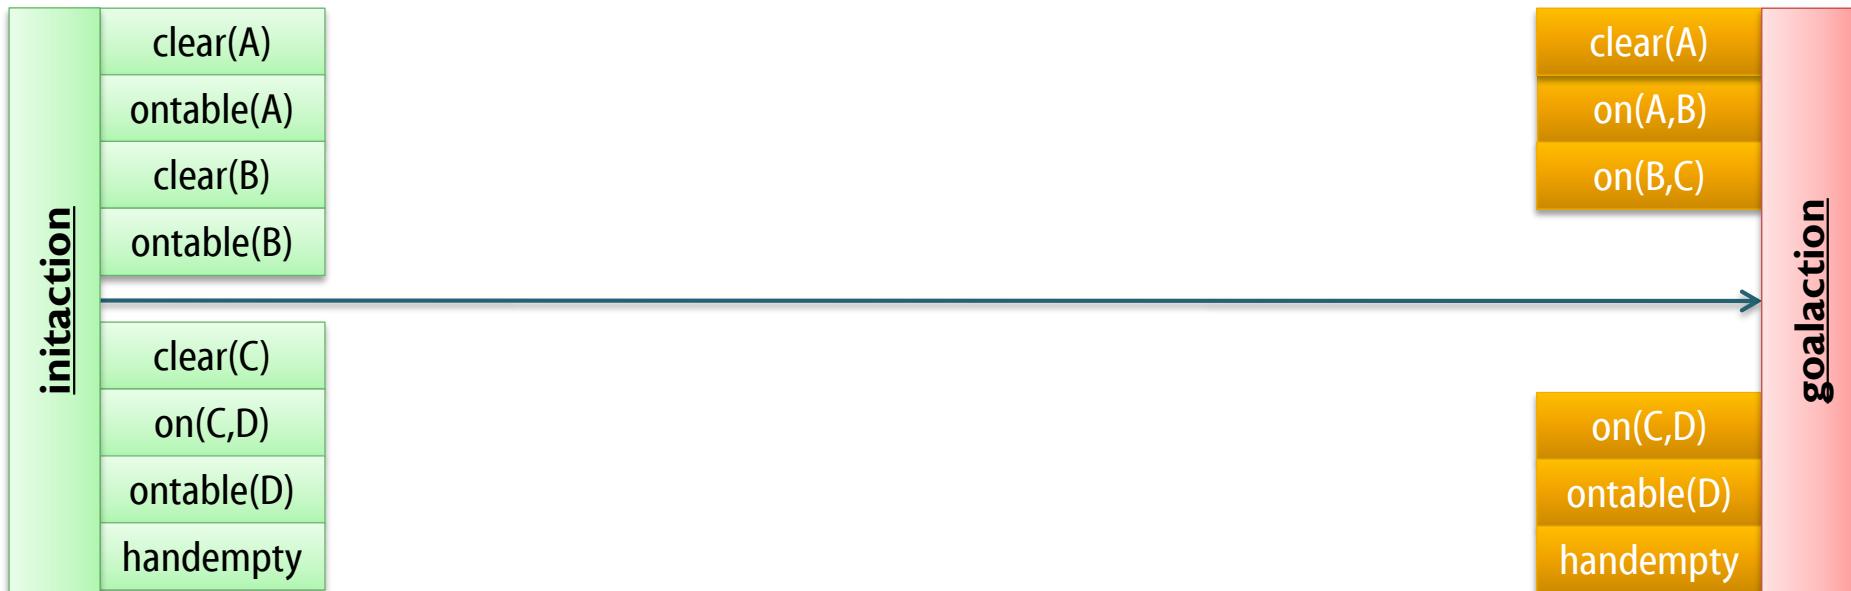
Partial-Order Planning



- No sequence → must have explicit **precedence constraints**

A "fake" initial action,
whose effects are
the facts of the initial
state

A "fake" goal action,
whose preconditions are
the conjunctive goal facts



initaction happens *before* goalaction

Partial-Order Planning

- This is the initial node in the search space – not a state!

Initial node containing initial plan

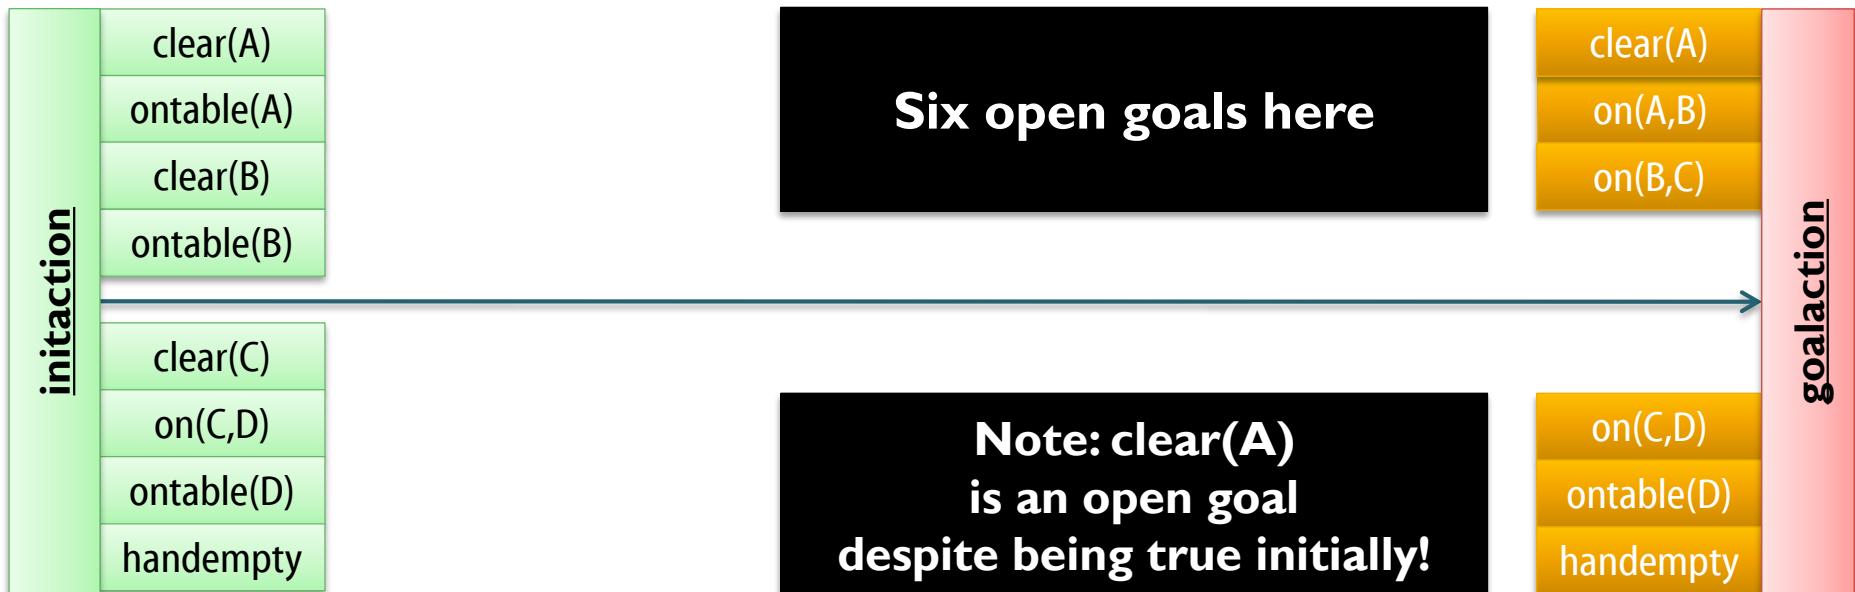


Partial-Order Planning

122

jonkv@ida

- No current state – what are the successors?
 - Don't need to "add any action anywhere"
 - We can identify a set of **issues to take care of** ("todos") called **flaws**
- First type of flaw: **Open goal**
 - When an **action** has a **precondition** that is not **explicitly marked** as handled



Resolving Open Goals 1

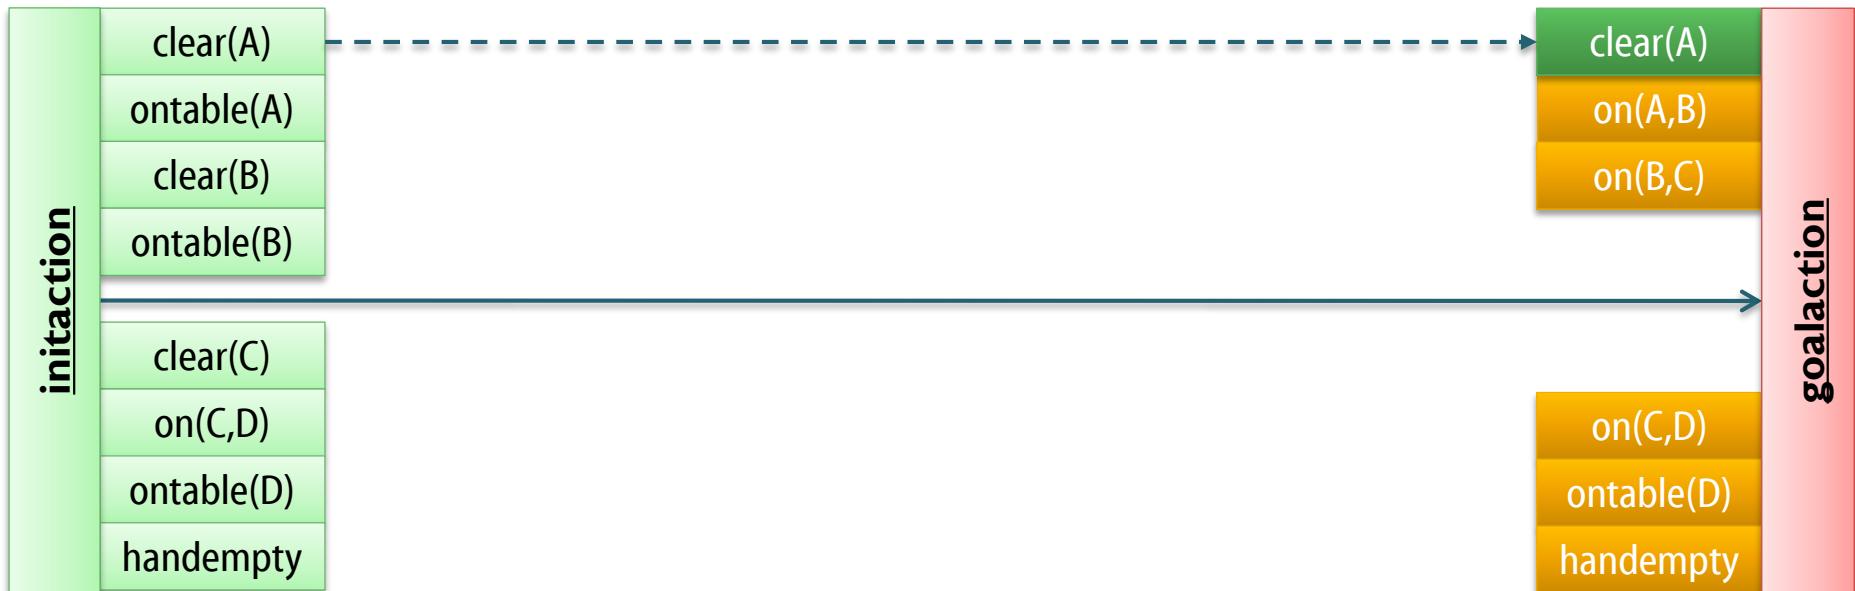


- To resolve an open goal:

Alternative I:
Add a causal link
from an existing action
that achieves it

Causal link:

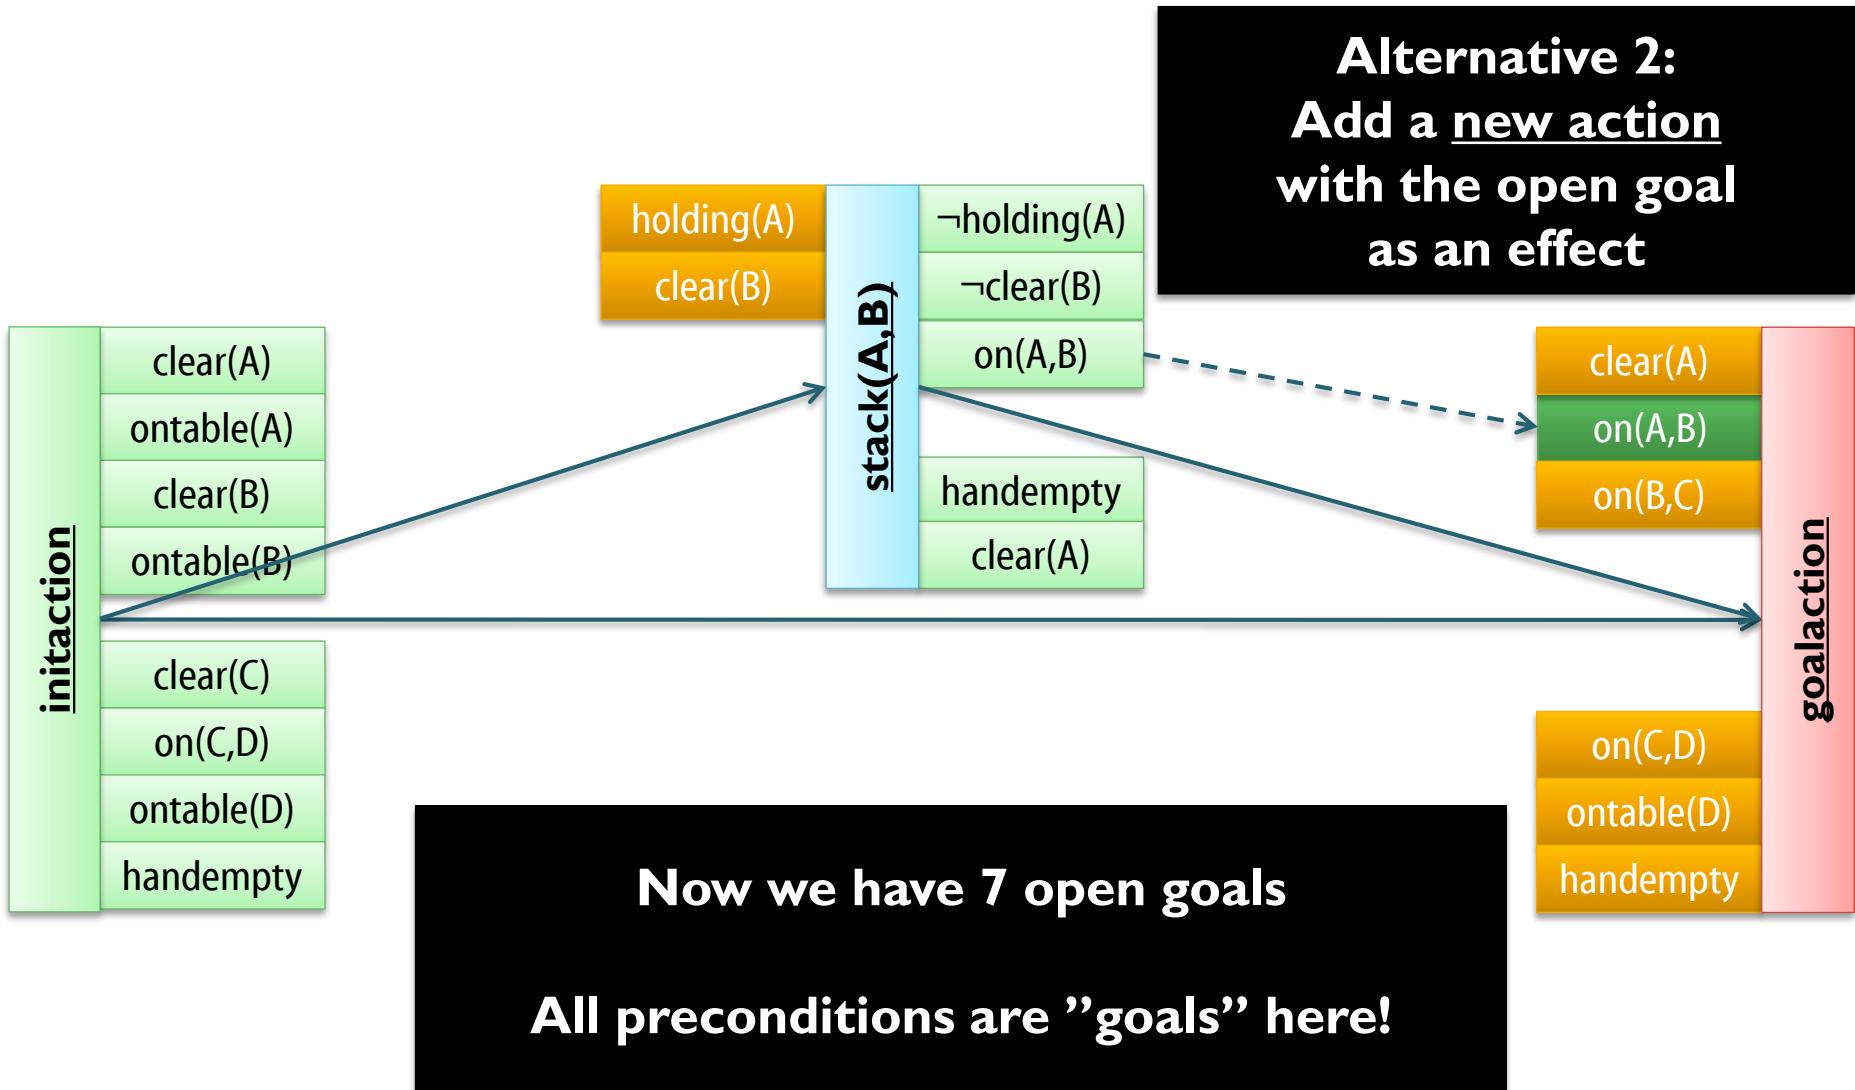
- 1) "initaction causes (achieves) clear(A) for goalaction"
- 2) "From the start of the link until its end, the property [clear(A)] will be true!"



Resolving Open Goals 2



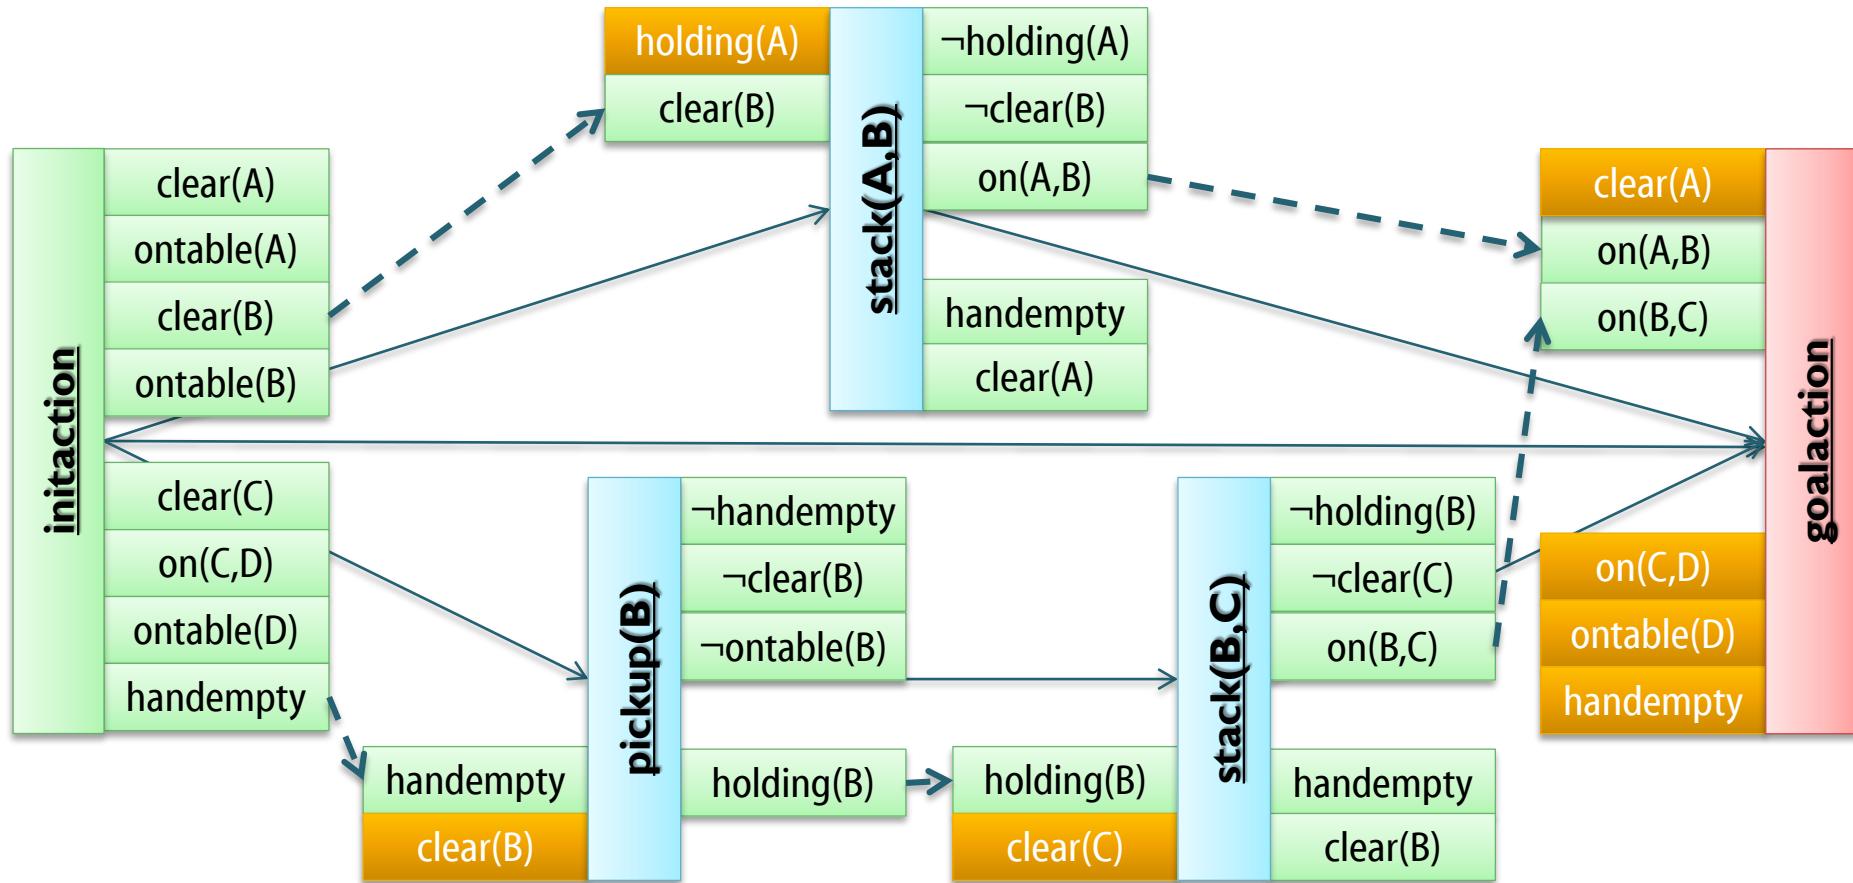
- To resolve an open goal:



Partial Orders



- Some pairs of actions can remain unordered



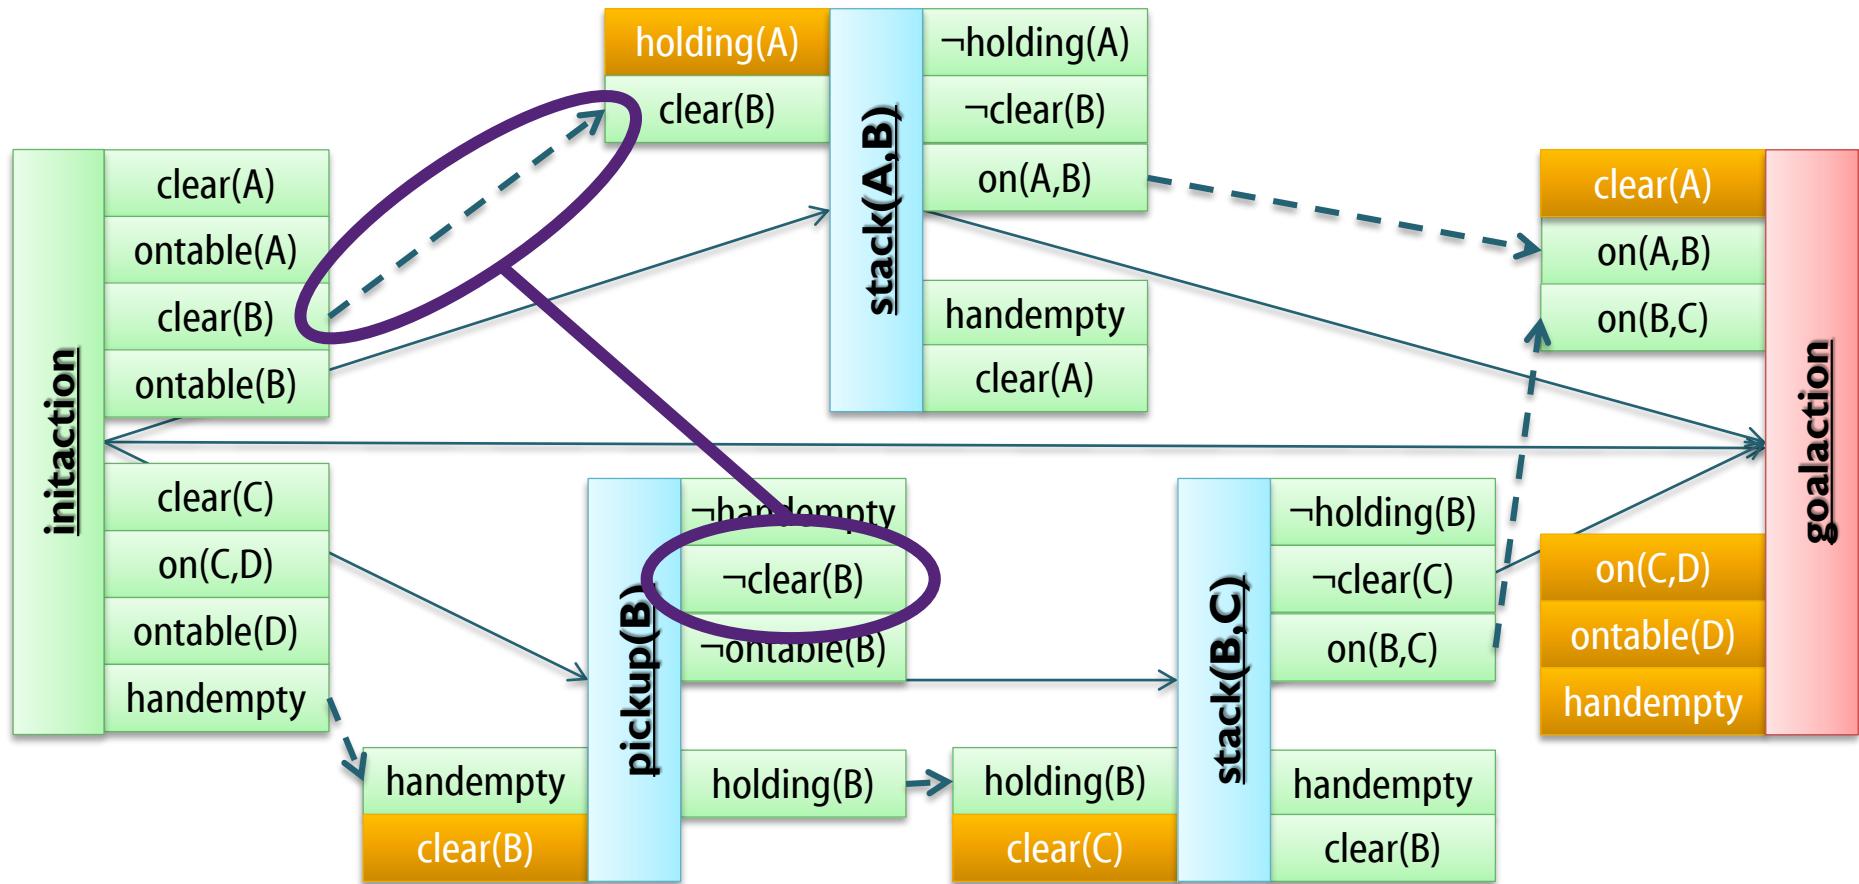
Threats

126

jonkv@ida

■ Second flaw type: A threat to be resolved

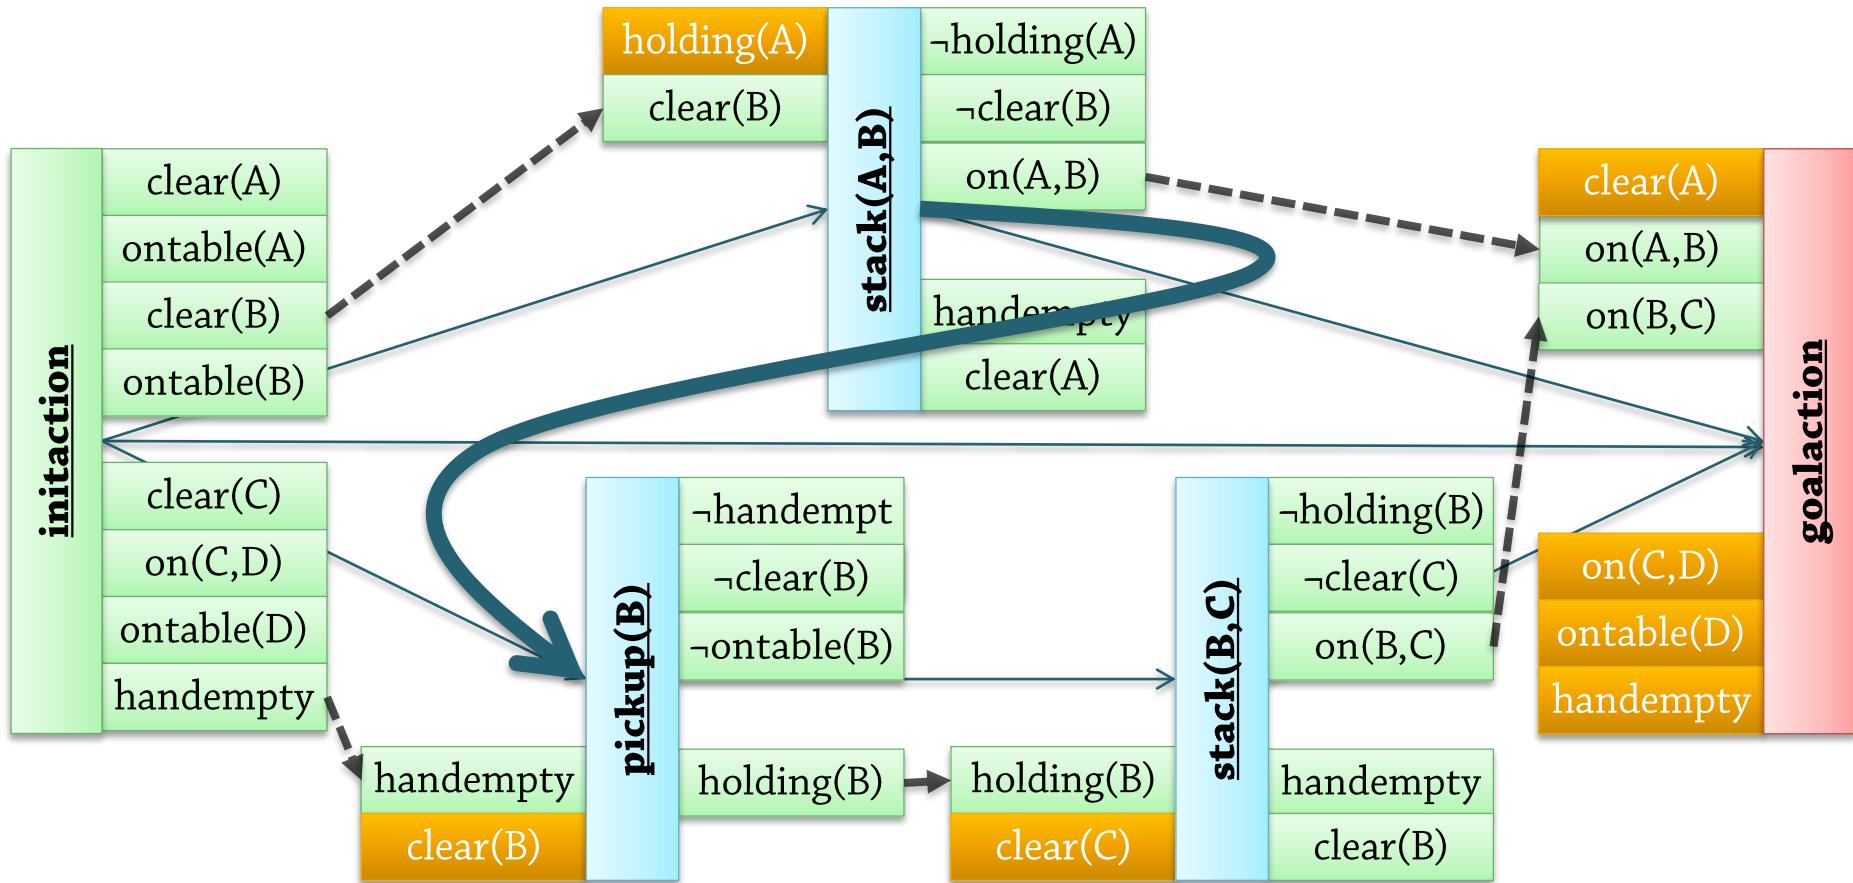
- initiation supports clear(B) for stack(A,B) – there's a causal link
- pickup(B) deletes clear(B), and may occur "during" the link
- We can't be certain that clear(B) still holds when stack(A,B) starts!



Resolving Threats 1



- How to make sure that $\text{clear}(B)$ holds when $\text{stack}(A,B)$ starts?
 - Alternative 1: The action that disturbs the precondition is placed after the action that has the precondition
 - Only possible if the resulting partial order is consistent (acyclic)!

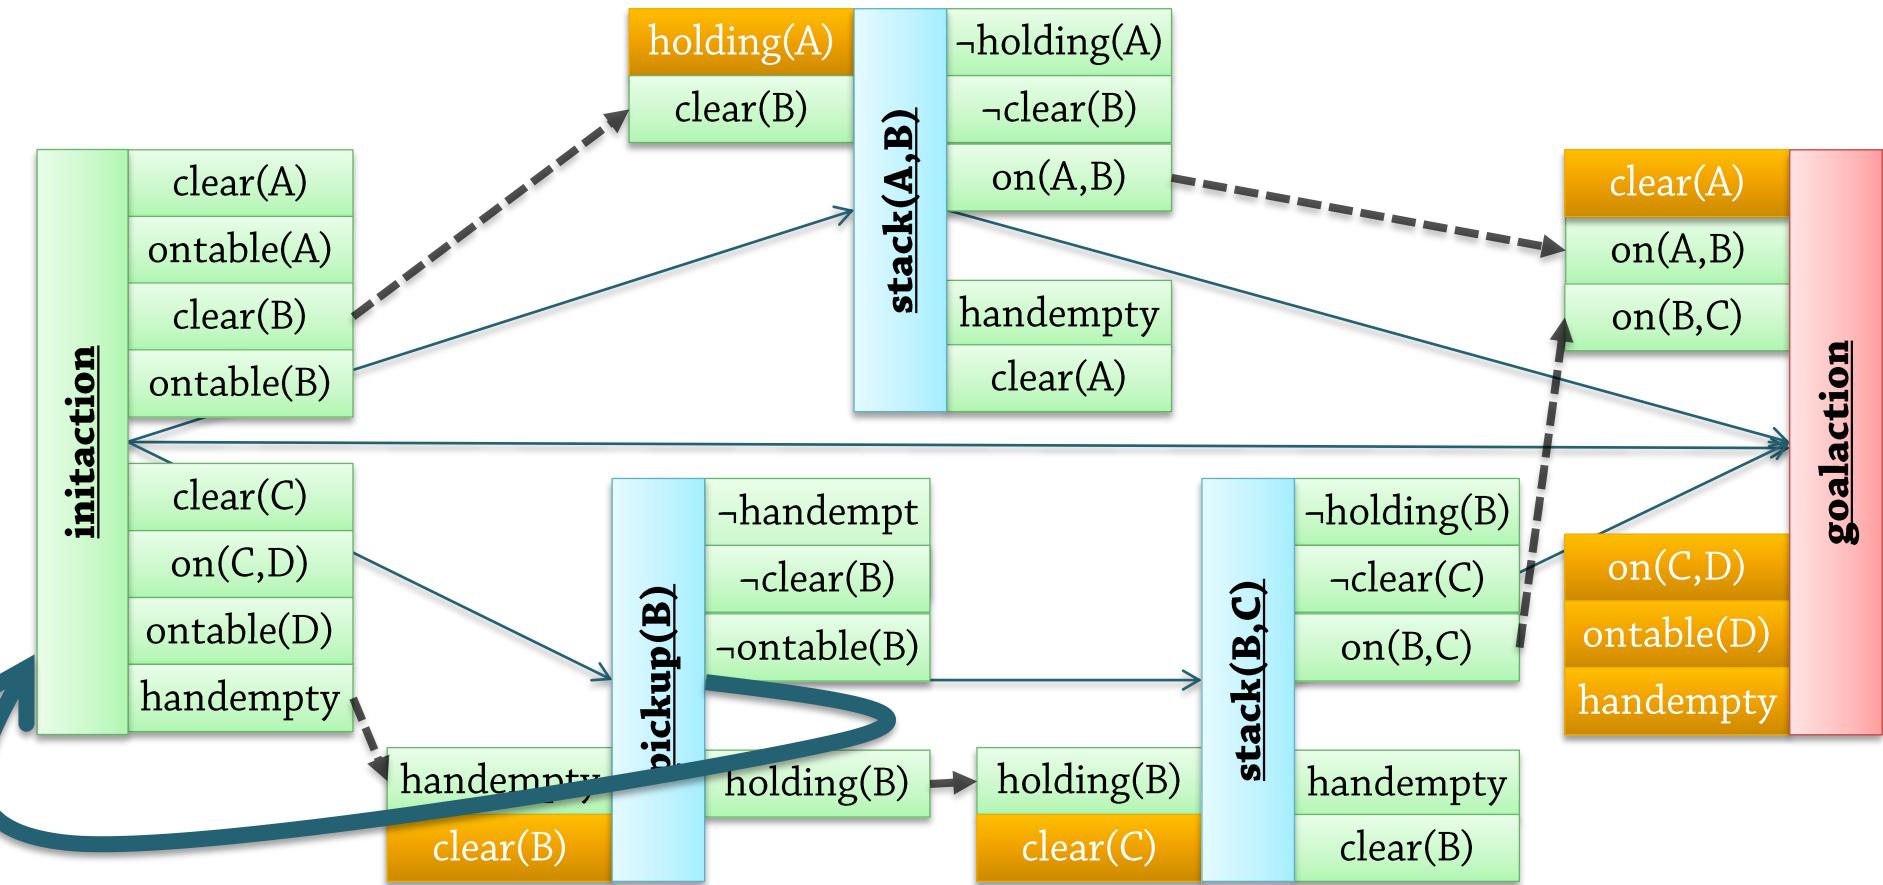


Resolving Threats 2

128

jonkv@ida

- Alternative 2:
 - The action that disturbs the precondition is placed before the action that supports the precondition
 - Only possible if the resulting partial order is consistent – not in this case!



- As usual, uninformed search is not sufficient
 - Heuristics are required

Remember what a flaw is!

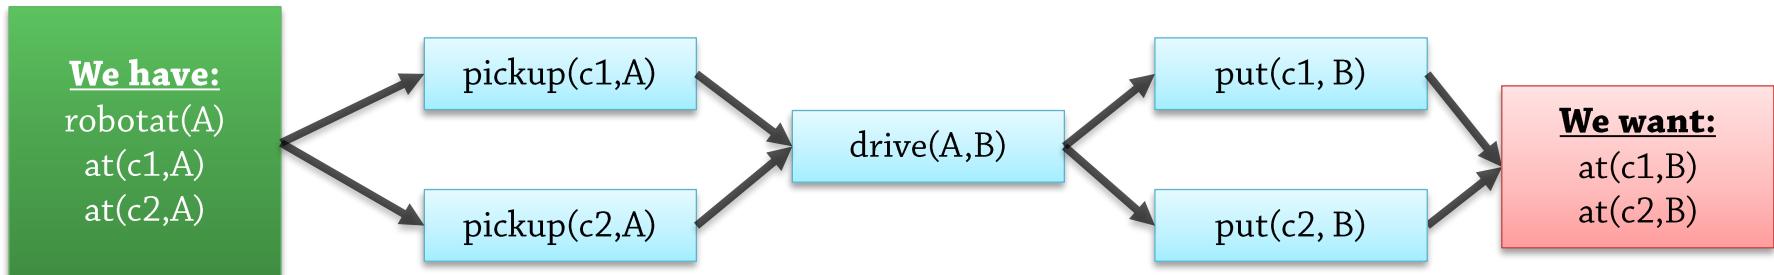
- Not something "wrong" in a final solution
- Not a mistake
- Simply bookkeeping for "something we haven't taken care of yet"
 - Open goal: Must decide how to achieve a precondition
 - Threat: Must decide how to order actions

Partial-Order Solutions



■ Original motivation: **performance**

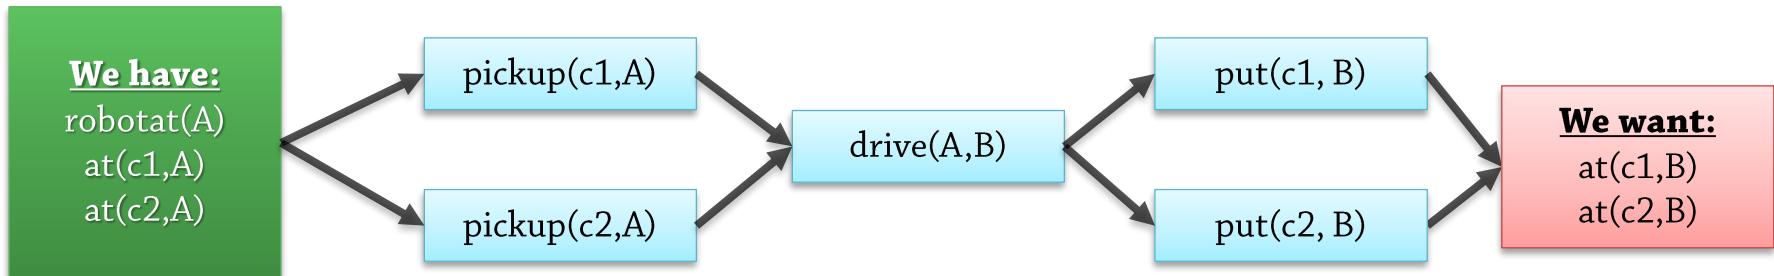
- Therefore, a partial-order plan is a **solution** iff **all sequential** plans satisfying the ordering are solutions
 - Similarly, **executable** iff corresponding sequential plans are executable
 - <pickup(c1,A), pickup(c2,A), drive(A,B), put(c1,B), put(c2,B)>
 - <pickup(c2,A), pickup(c1,A), drive(A,B), put(c1,B), put(c2,B)>
 - <pickup(c1,A), pickup(c2,A), drive(A,B), put(c2,B), put(c1,B)>
 - <pickup(c2,A), pickup(c1,A), drive(A,B), put(c2,B), put(c1,B)>



Partial Orders and Concurrency



- Can be extended to allow concurrent execution
 - Requires an extended action model!
 - Our model *does not define* what happens if c_1 and c_2 are picked up simultaneously (or even if we are allowed to do this)!



Beyond State Space Planning: An Overview

Many alternative search spaces exist;
some examples here!

SAT Planning

- SAT: Boolean satisfiability, first problem proven NP-complete
 - Translate planning to a SAT problem – propositions + complex formula
 - Apply existing SAT solver
 - Each **SAT solution** corresponds to a **solution plan**

"Action propositions": Which actions are executed, and when?

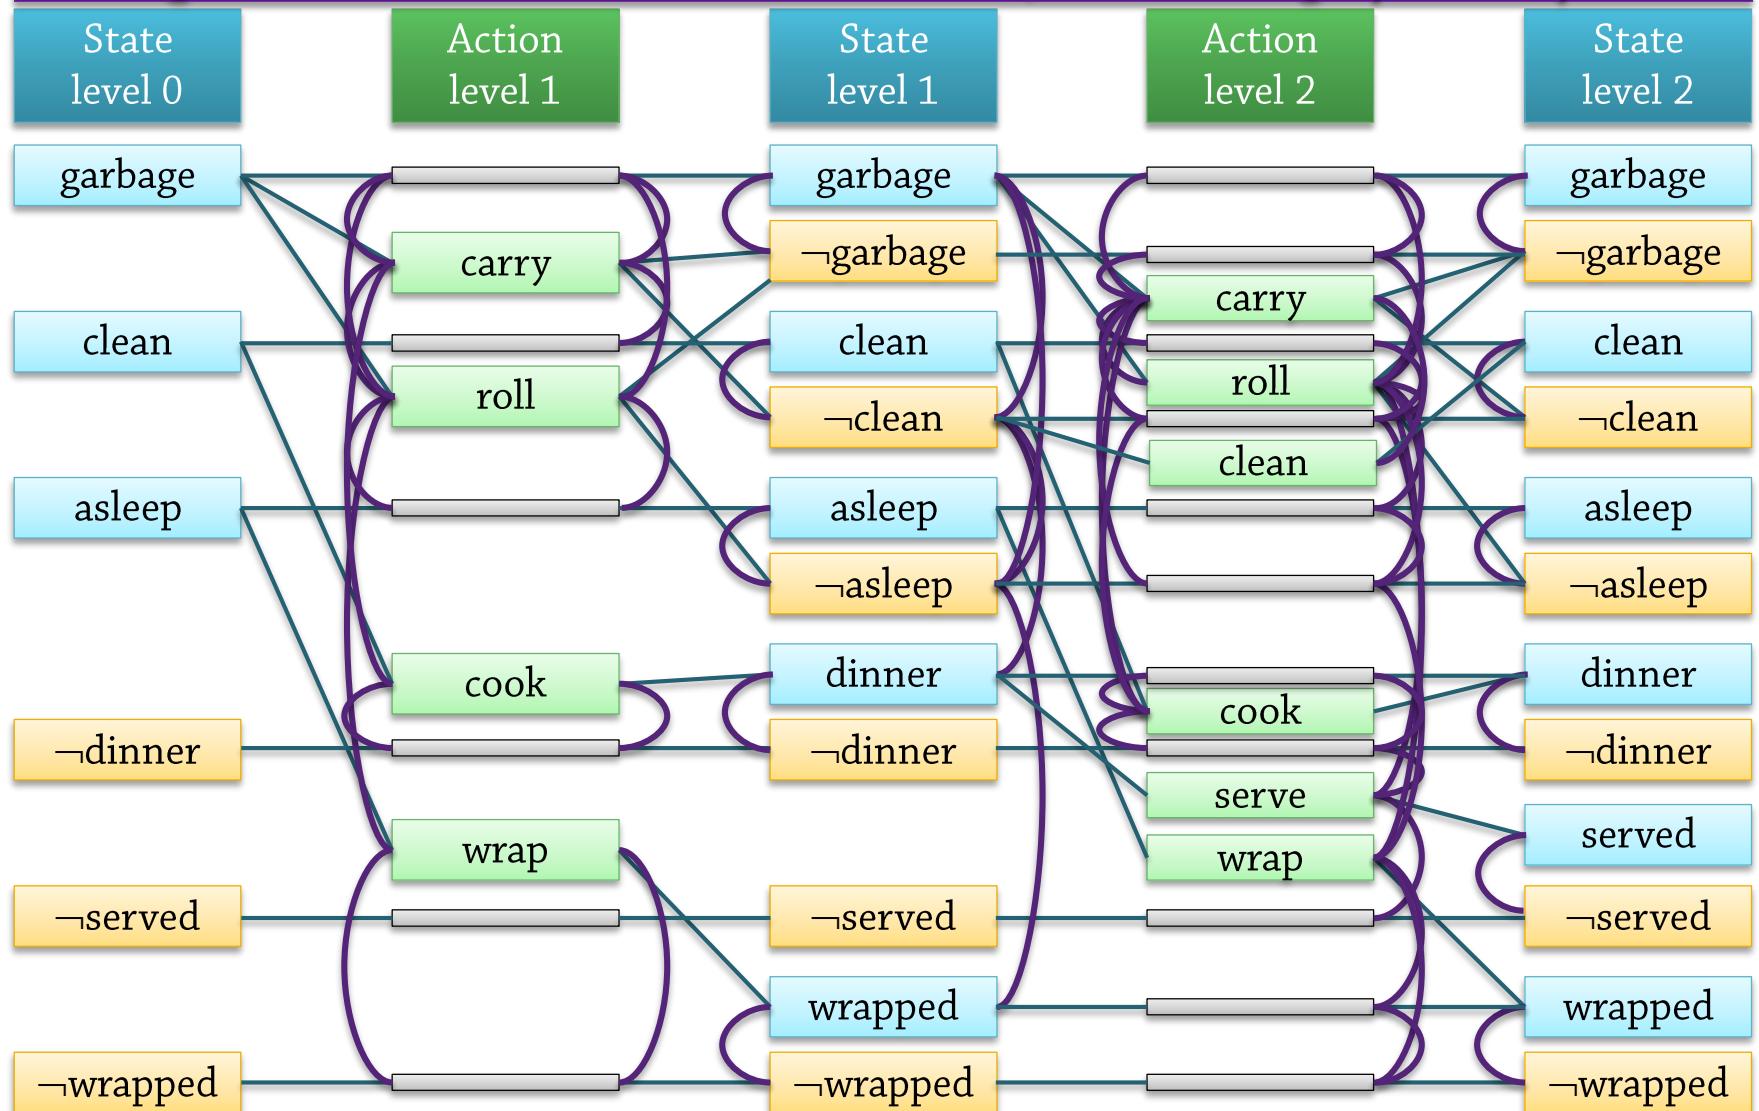
"Fact propositions": Which facts are true, and when?

Planning Graphs

134



Compute a graph representing what might be achievable per time step, together with mutual exclusion relations; search this graph for a plan



Beyond Classical Planning: Time, Concurrency, Resources

Time, Concurrency and Resources



- Many planners support **time** and **concurrency**
 - Actions have **durations**
 - Actions can be **scheduled** in parallel, suitable for multiple agents
 - Example: Duration determined by current state**

(:durative-action heat-water

:parameters (?p - pan)

:duration (= ?duration (/ (- 100 (temperature ?p)) (heat-rate)))

:condition (and (at start (**full** ?p))

(at start (**onHeatSource** ?p))

(at start (**byPan**))

(over all (**full** ?p))

(over all (**onHeatSource** ?p))

(over all (**heating** ?p))

(at end (**byPan**)))

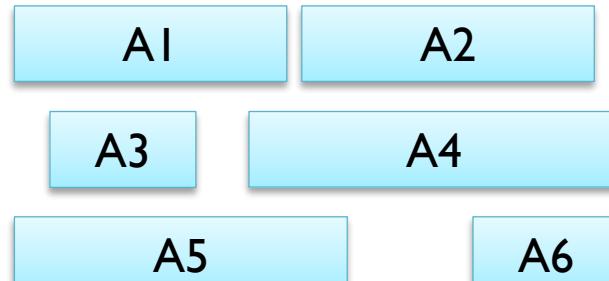
Conditions / effects
at the *start* or *end*
of the action

:effect (and (at start (**heating** ?p))

(at end (not (**heating** ?p))))

(at end (**assign (temperature** ?p) 100))))

)



Implemented by quite a few planners –
affects algorithms and heuristics!

Time, Concurrency and Resources (2)



- Example: Duration selected by planner

(:durative-action heat-water

:parameters (?p - pan)

:duration (at end (<= ?duration (/ (- 100 (temperature ?p)) (heat-rate))))

:condition (and (at start (full ?p))

(at start (onHeatSource ?p))

(at start (byPan))

(over all (full ?p))

(over all (onHeatSource ?p))

(over all (heating ?p))

(at end (byPan)))

:effect (and (at start (heating ?p))

(at end (not (heating ?p))))

(at end (increase (temperature ?p) (* ?duration (heat-rate)))))

General increase and decrease effects

for numeric values

(handles concurrent production and consumption of resources)

Time, Concurrency and Resources (3)



■ Some planners support continuous linear effects

- Occurring throughout the duration of the action

- (:durative-action fly

:parameters (?p - airplane ?a ?b - airport)

:duration (= ?duration (flight-time ?a ?b))

:condition (and (at start (at ?p ?a))

(over all (inflight ?p))

(over all (>= (fuel-level ?p) 0)))

:effect (and (at start (not (at ?p ?a)))

(at start (inflight ?p))

(at end (not (inflight ?p)))

(at end (at ?p ?b))

(decrease (fuel-level ?p) (* #t (fuel-consumption-rate ?p))))

- (:action midair-refuel

:parameters (?p)

:precondition (inflight ?p)

:effect (assign (fuel-level ?p) (fuel-capacity ?p)))

Time, Concurrency and Resources (4)



■ Hybrid discrete/continuous planning

- (:durative-action navigate

:control-variables ((velX) (velY))

:duration (and (<= ?duration 5000))

:condition (and

(over all (>= (velX) -4)) (over all (<= (velX) 4))

(over all (>= (velY) -4)) (over all (<= (velY) 4))

(over all (<= (x) 700)) (over all (>= (x) 0))

(over all (<= (y) 700)) (over all (>= (y) 0))

(at start (AUV-ready)))

:effect (and

(at start (not (AUV-ready)))

(at end (AUV-ready))

(increase (x) (* 1.0 (velX) #t))

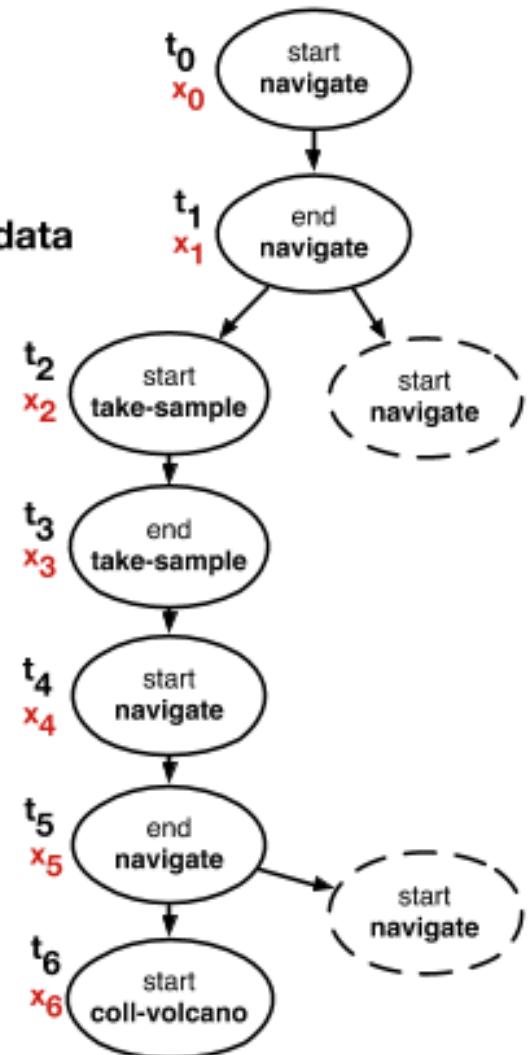
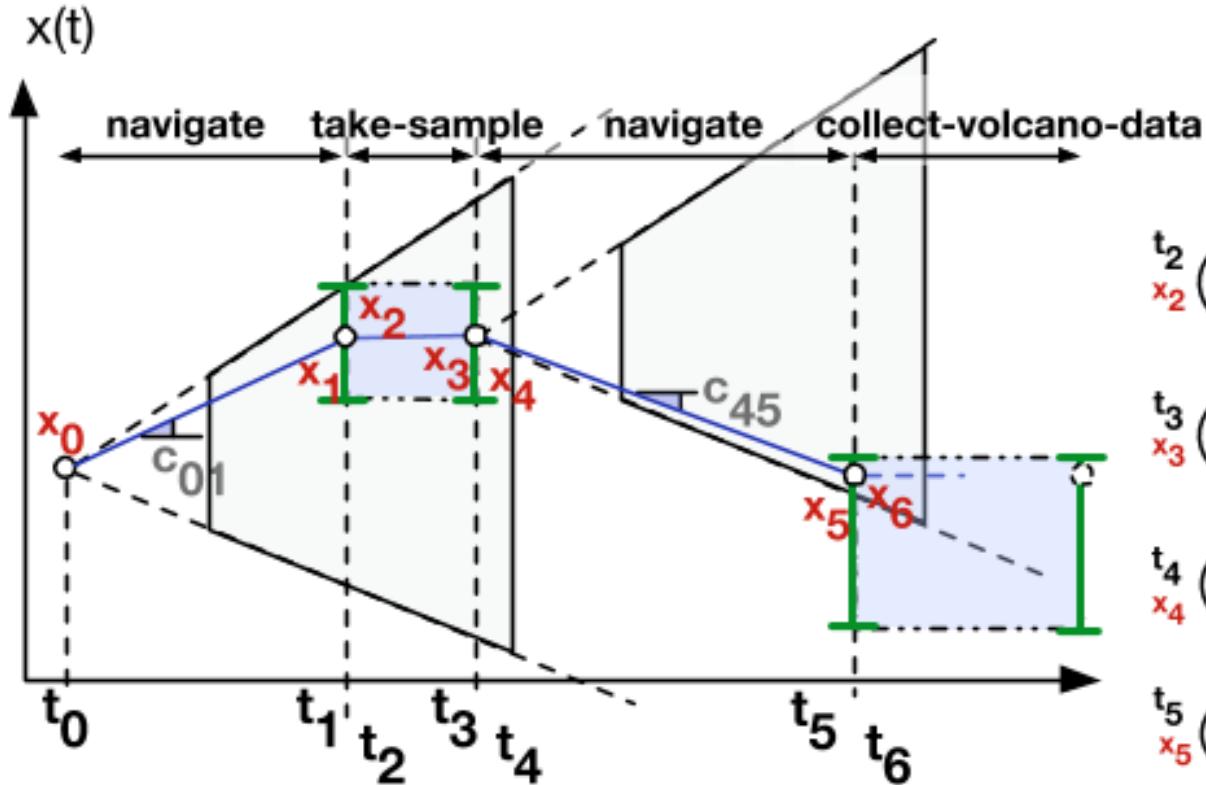
(increase (y) (* 1.0 (velY) #t))))

Can vary continuously
during the action

Continuous effects
depend on the
instantaneous values of
control variables

Time, Concurrency and Resources (5)

- Flow tubes result in reachable regions (shaded)
 - Rectangles are conditions for **take-sample**, **collect**



Beyond Classical Planning: Incomplete Information

Multiple Outcomes



- In reality, actions may have **multiple outcomes**
 - Some outcomes can indicate **faulty / imperfect execution**
 - **pick-up(object)**
Intended outcome: carrying(object) is true
Unintended outcome: carrying(object) is false
 - **move(100,100)**
Intended outcome: xpos(robot)=100
Unintended outcome: xpos(robot) != 100
 - **jump-with-parachute**
Intended outcome: alive is true
Unintended outcome: alive is false
 - Some outcomes are more **random**, but clearly **desirable / undesirable**
 - Pick a present at random – do I get the one I longed for?
 - Toss a coin – do I win?
 - Sometimes we have **no clear idea** what is desirable
 - Outcome will affect how we can continue, but in less predictable ways

To a **planner**,
there is generally
no difference
between these
cases!

FOND Planning

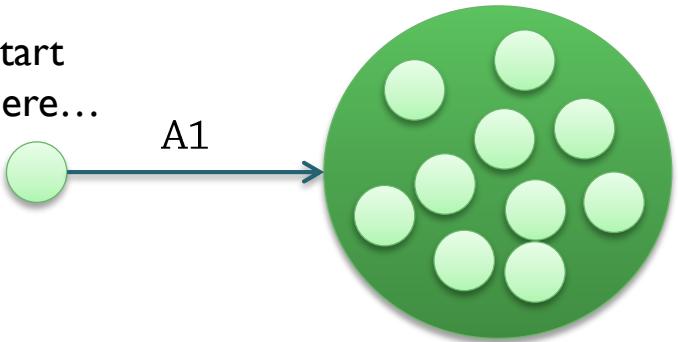
- **Nondeterministic** planning:
 - Many possible outcomes for each action
- FOND: **Fully Observable** Non-Deterministic
 - After executing an action, sensors determine exactly which state we are in

Planning

Model says: we end up in one of these states

Start here...

A1

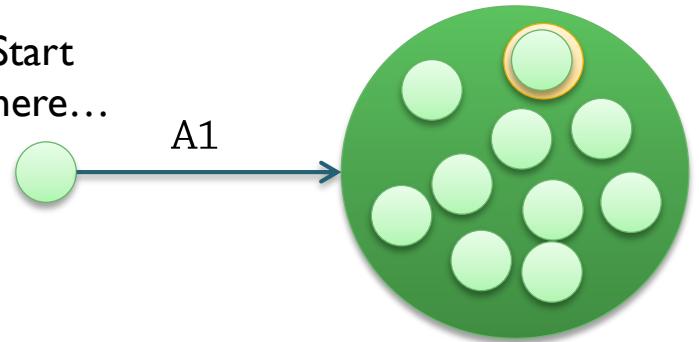


Execution

Sensors say: we are in this state!

Start here...

A1

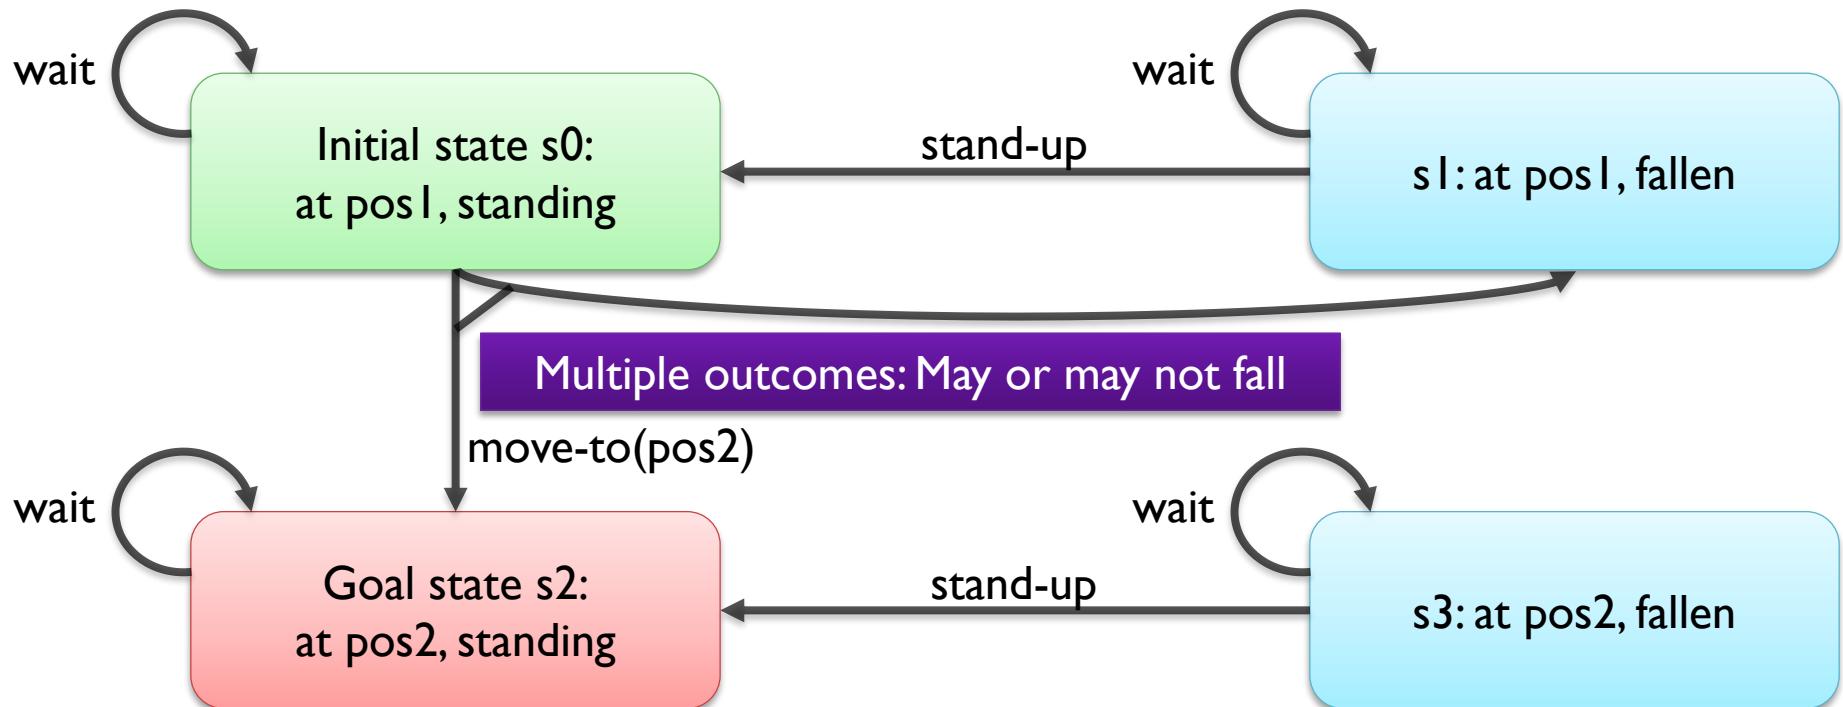


FOND Planning: Plan Structure (1)

144

jonkv@ida

- Example state transition system:



- Intuitive strategy:**

- while** (not in s2) {
move-to(pos2);
if (fallen) stand-up;
}

In FOND planning, action choice
should depend on actual outcomes
(through the current state)

There may be no upper bound on how
many actions we may have to execute!

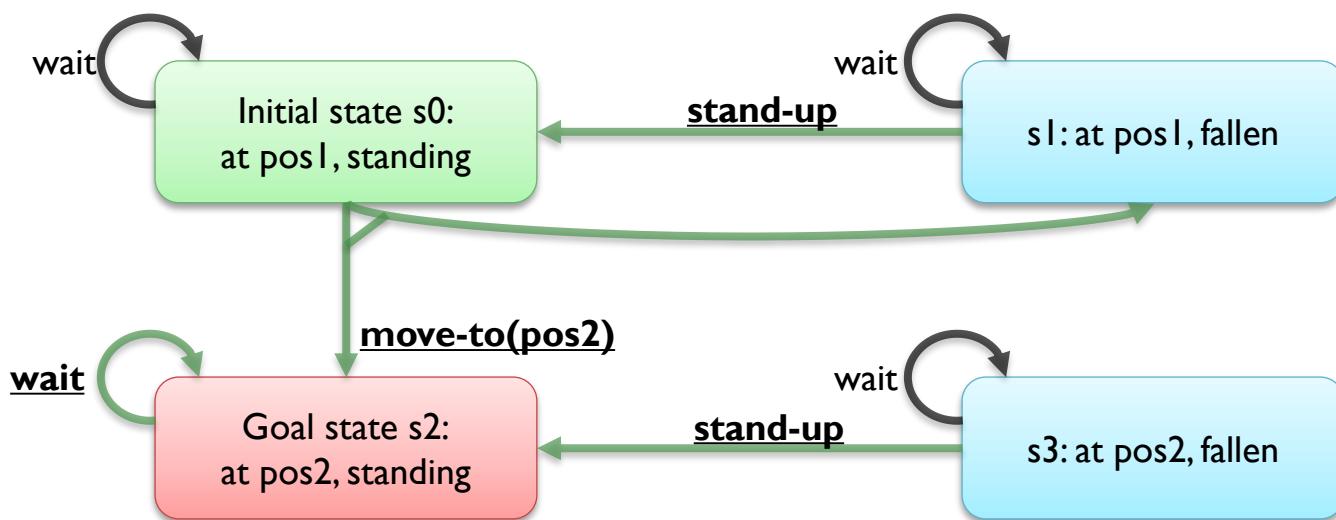
FOND Planning: Plan Structure (2)



■ Examples of formal plan structures:

- **Conditional plans** (with if/then/else statements)
- **Policies** $\pi : S \rightarrow A$
 - Defining, for each state, which action to execute whenever we end up there
 - $\pi(s_0) = \text{move-to(pos2)}$
 - $\pi(s_1) = \text{stand-up}$
 - $\pi(s_2) = \text{wait}$
 - $\pi(s_3) = \text{stand-up}$

Or at least, for every state
that is *reachable* from the given initial state
(A policy can be a *partial function*)

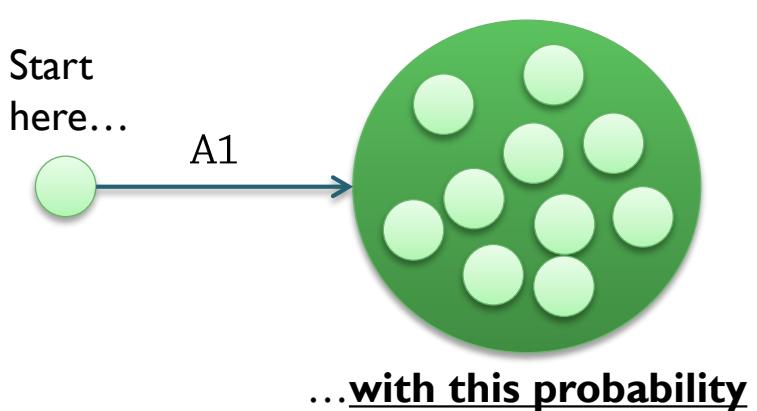


■ Probabilistic planning → a stochastic system

- $P(s, a, s')$: Given that we are in s and execute a , the probability of ending up in s'
- For every state s and action a , we have $\sum_{s' \in S} P(s, a, s') = 1$: The world gives us 100% probability of ending up in some state

Planning

Model says: we end up in one of these states



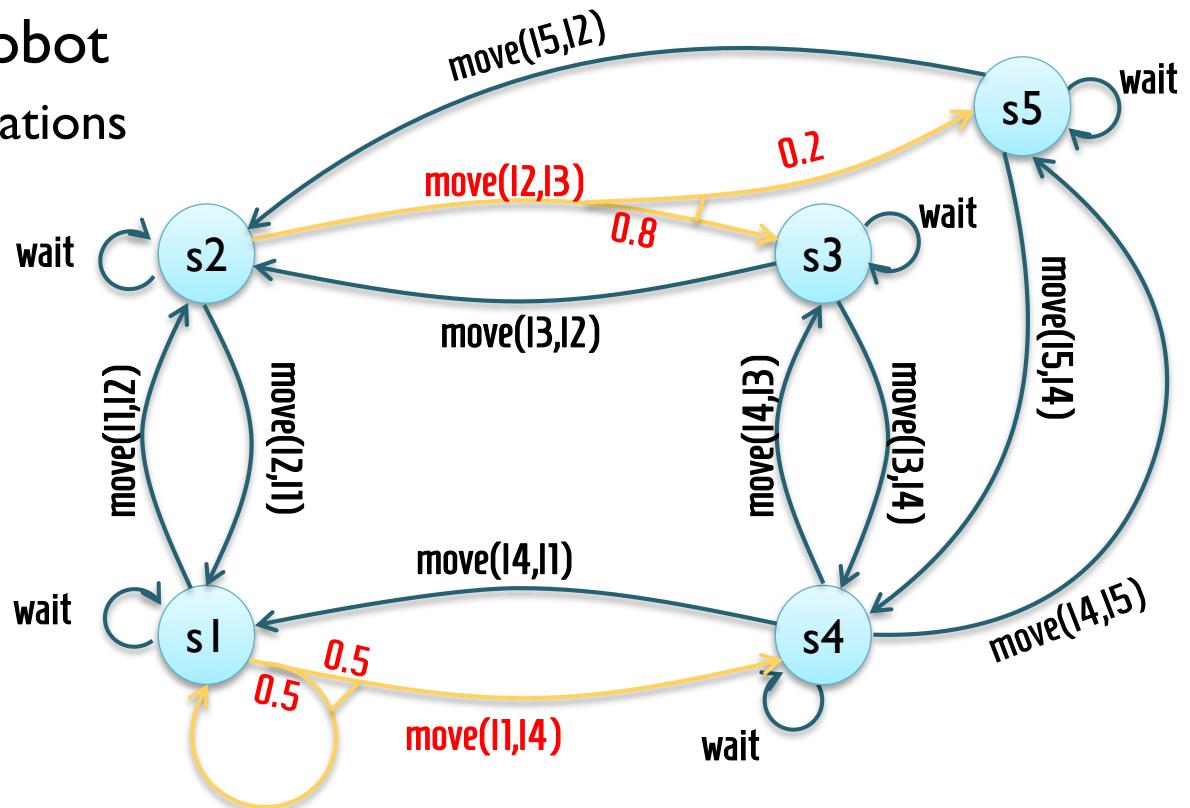
Stochastic System Example

147

jonkv@ida

Example: A single robot

- Moving between 5 locations
- For simplicity, states correspond directly to locations
 - s1: at(r1, l1)
 - s2: at(r1, l2)
 - s3: at(r1, l3)
 - s4: at(r1, l4)
 - s5: at(r1, l5)



- Some transitions are deterministic, some are stochastic
 - Trying to move from l2 to l3: You may end up at l5 instead (20% risk)
 - Trying to move from l1 to l4: You may stay where you are instead (50% risk)
- (Can't always move in both directions, e.g. due to terrain gradient)

Overview

148

jonkv@ida

	Non-Observable: No information gained after action	Fully Observable: Exact outcome known after action
Deterministic: Exact outcome known in advance	Classical planning (possibly with extensions) (Information dimension is irrelevant; we know everything <i>before</i> execution)	
Non-deterministic: Multiple outcomes, no probabilities	NOND: Conformant Planning	FOND: Conditional / Contingent Planning
Probabilistic: Multiple outcomes with probabilities	Probabilistic Conformant Planning (Special case of POMDPs, Partially Observable MDPs)	Probabilistic Conditional Planning Stochastic Shortest Path Problems Markov Decision Processes (MDPs)

Conclusions

Example Questions

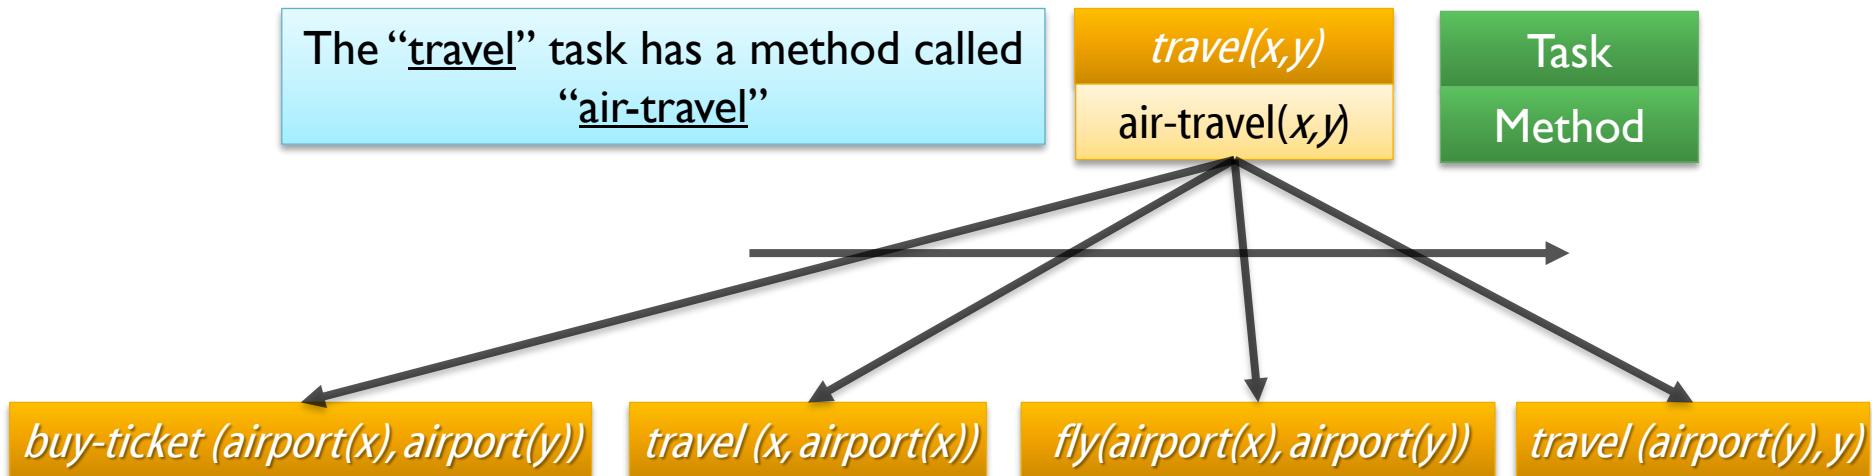


- Example questions:
 - Describe and explain:
 - Relaxation, and how it can be applied to create an admissible heuristic
 - PDB heuristics
 - Landmarks
 - The main advantage of partial-order planning over forward state space planning
 - The different kinds of flaws that can occur in a partially ordered plan
 - The ways in which you can resolve an open goal
 - Why planners don't use Dijkstra's algorithm to search the state space
 - Is the number of open goals an *admissible* heuristic? Why / why not?
 - Given a domain, problem and initial state:
 - Expand the forward state space by two levels
 - How can you apply hill climbing to planning?
 - How do you have to modify standard hill climbing, and why?

- **Deeper discussions** about all of these topics
- **Formal basis** for planning
 - Alternative representations of planning problems
 - Simple and complex state transition systems
- Completely different **principles for heuristics**
 - Landmarks
 - Pattern databases
- Alternative **search spaces**
 - Planning by conversion to boolean satisfiability problems
- Extended **expressivity**
 - Planning with time and concurrency
 - Planning with probabilistic actions (Markov Decision Processes)
 - Planning with non-classical goals

TDDD48 Automated Planning (2)

- Planning with **domain knowledge**
 - Using what you know: Temporal control rules
 - Breaking down a task into smaller parts: Hierarchical Task Networks



- Alternative **types** of planning
 - Path planning
- And so on...

